

# New hardware support in OSIRIS 4.0



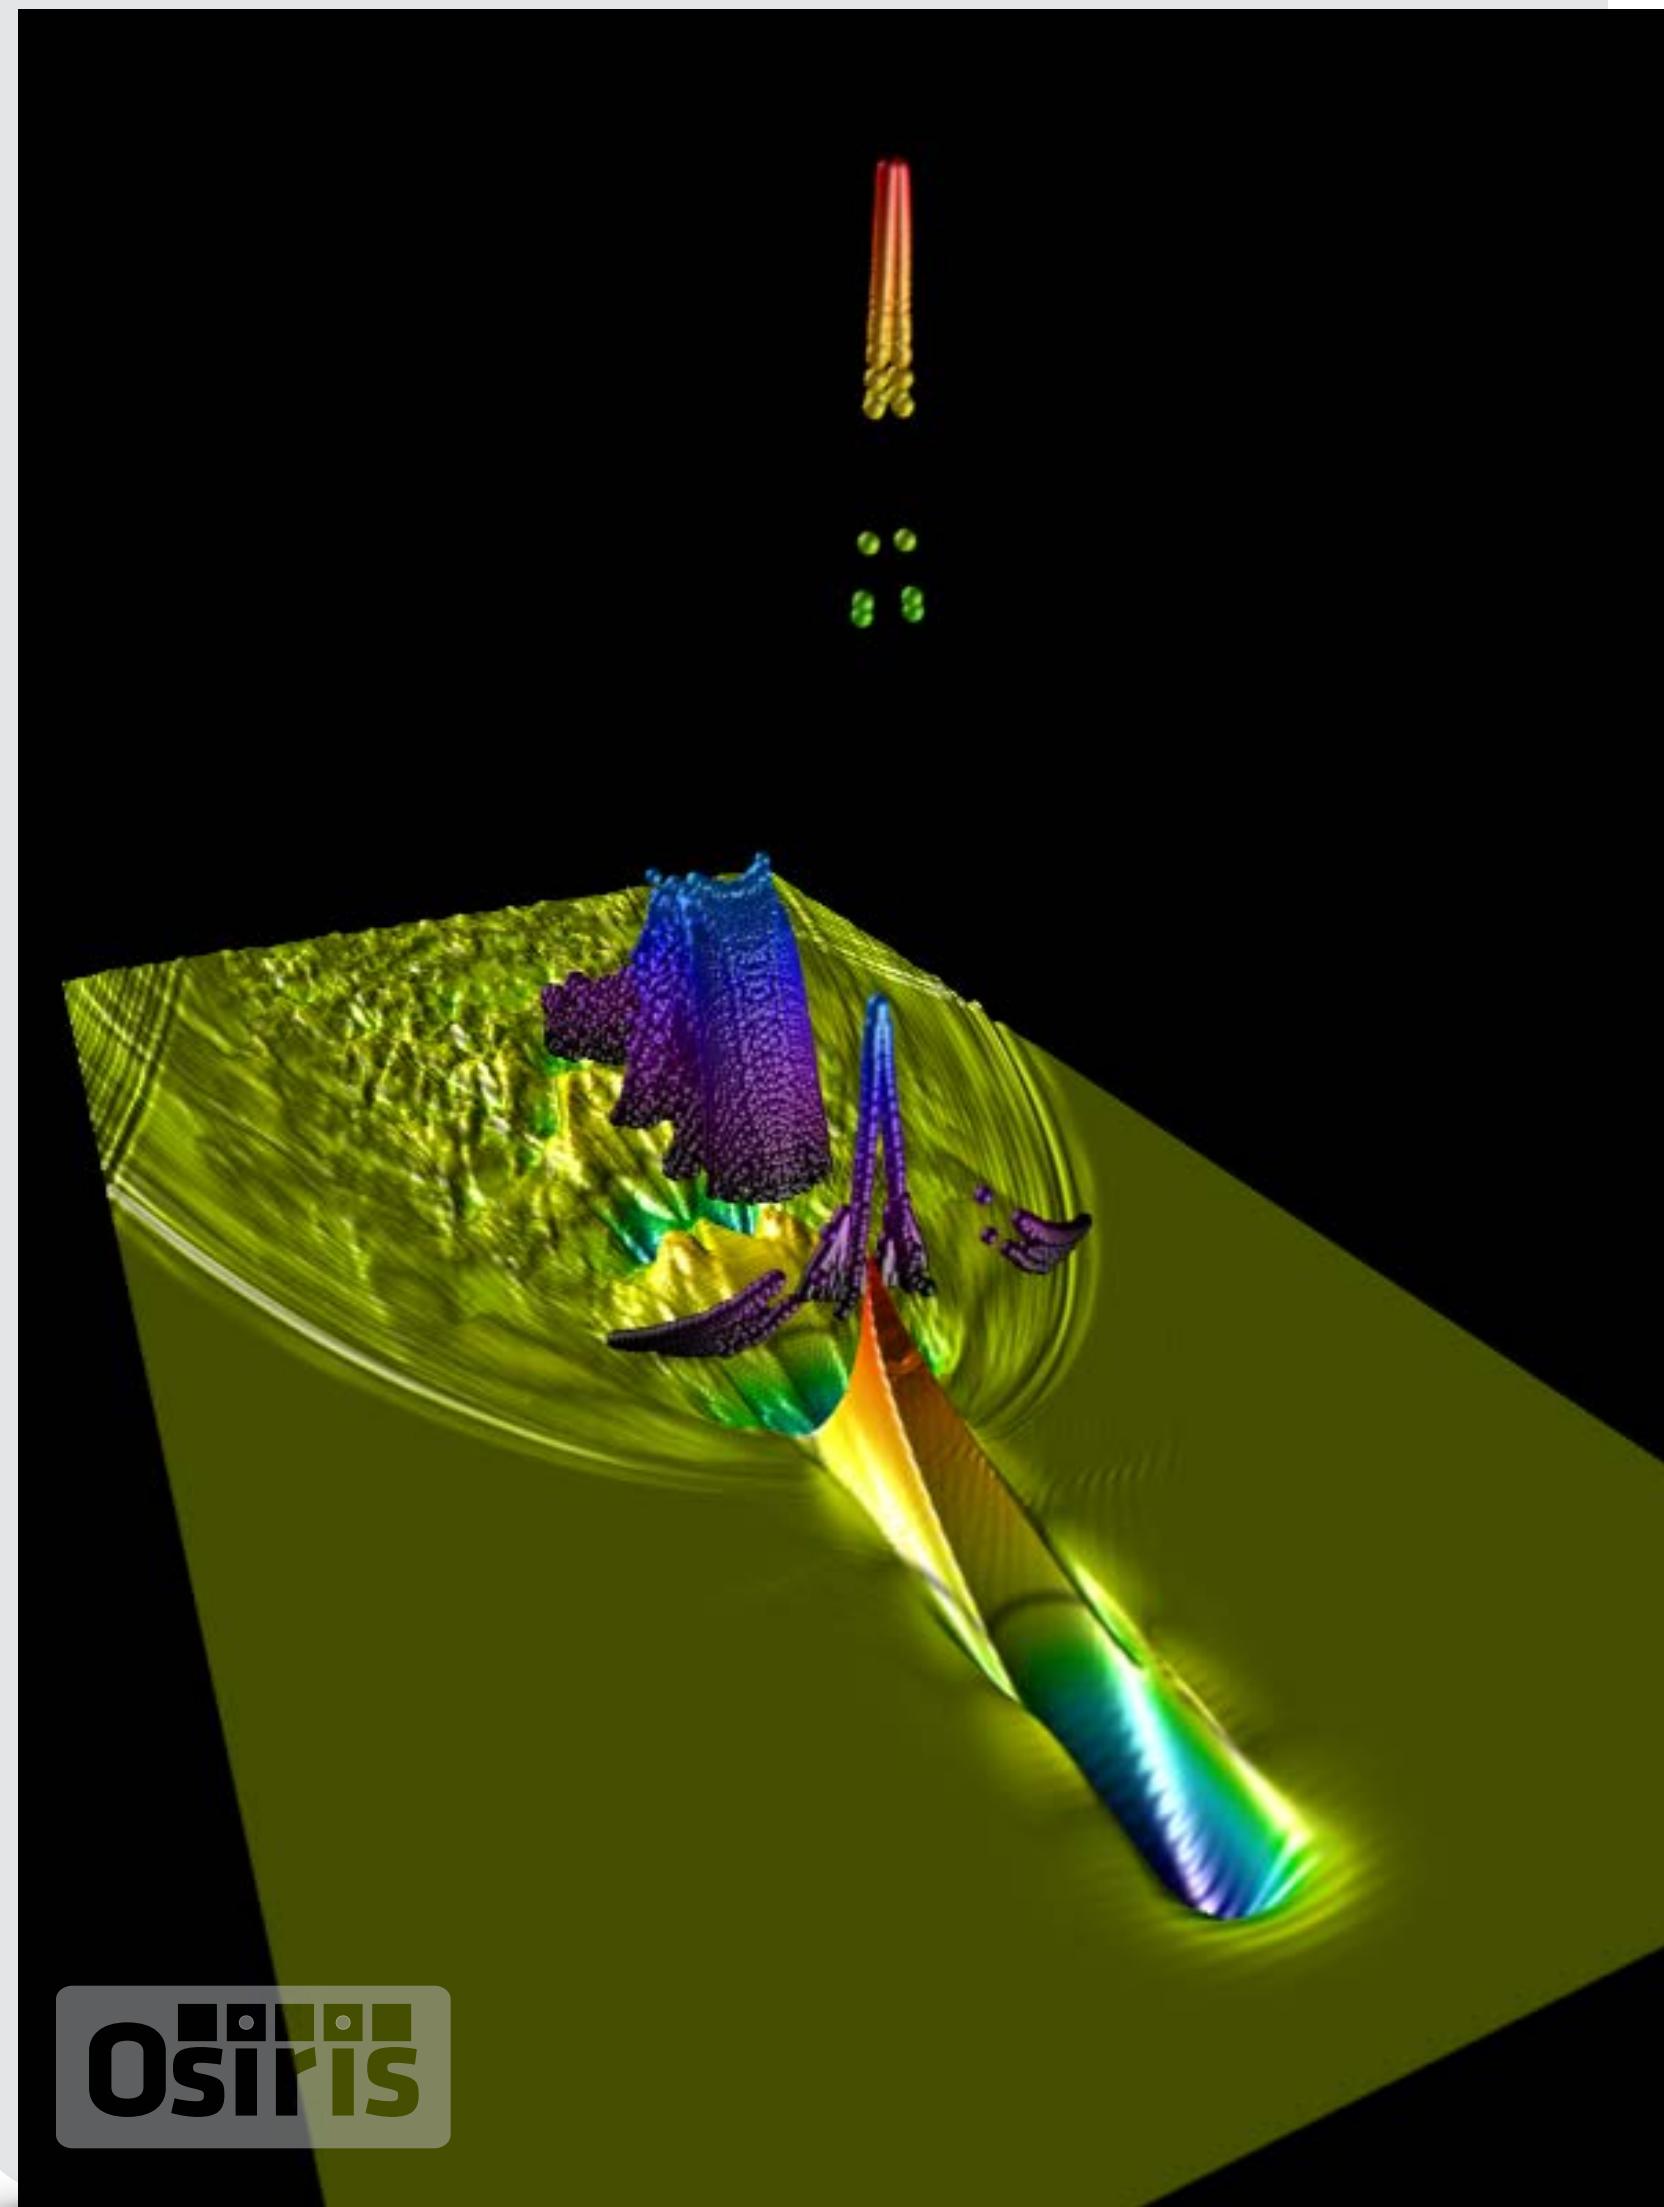
R. A. Fonseca<sup>1,2</sup>

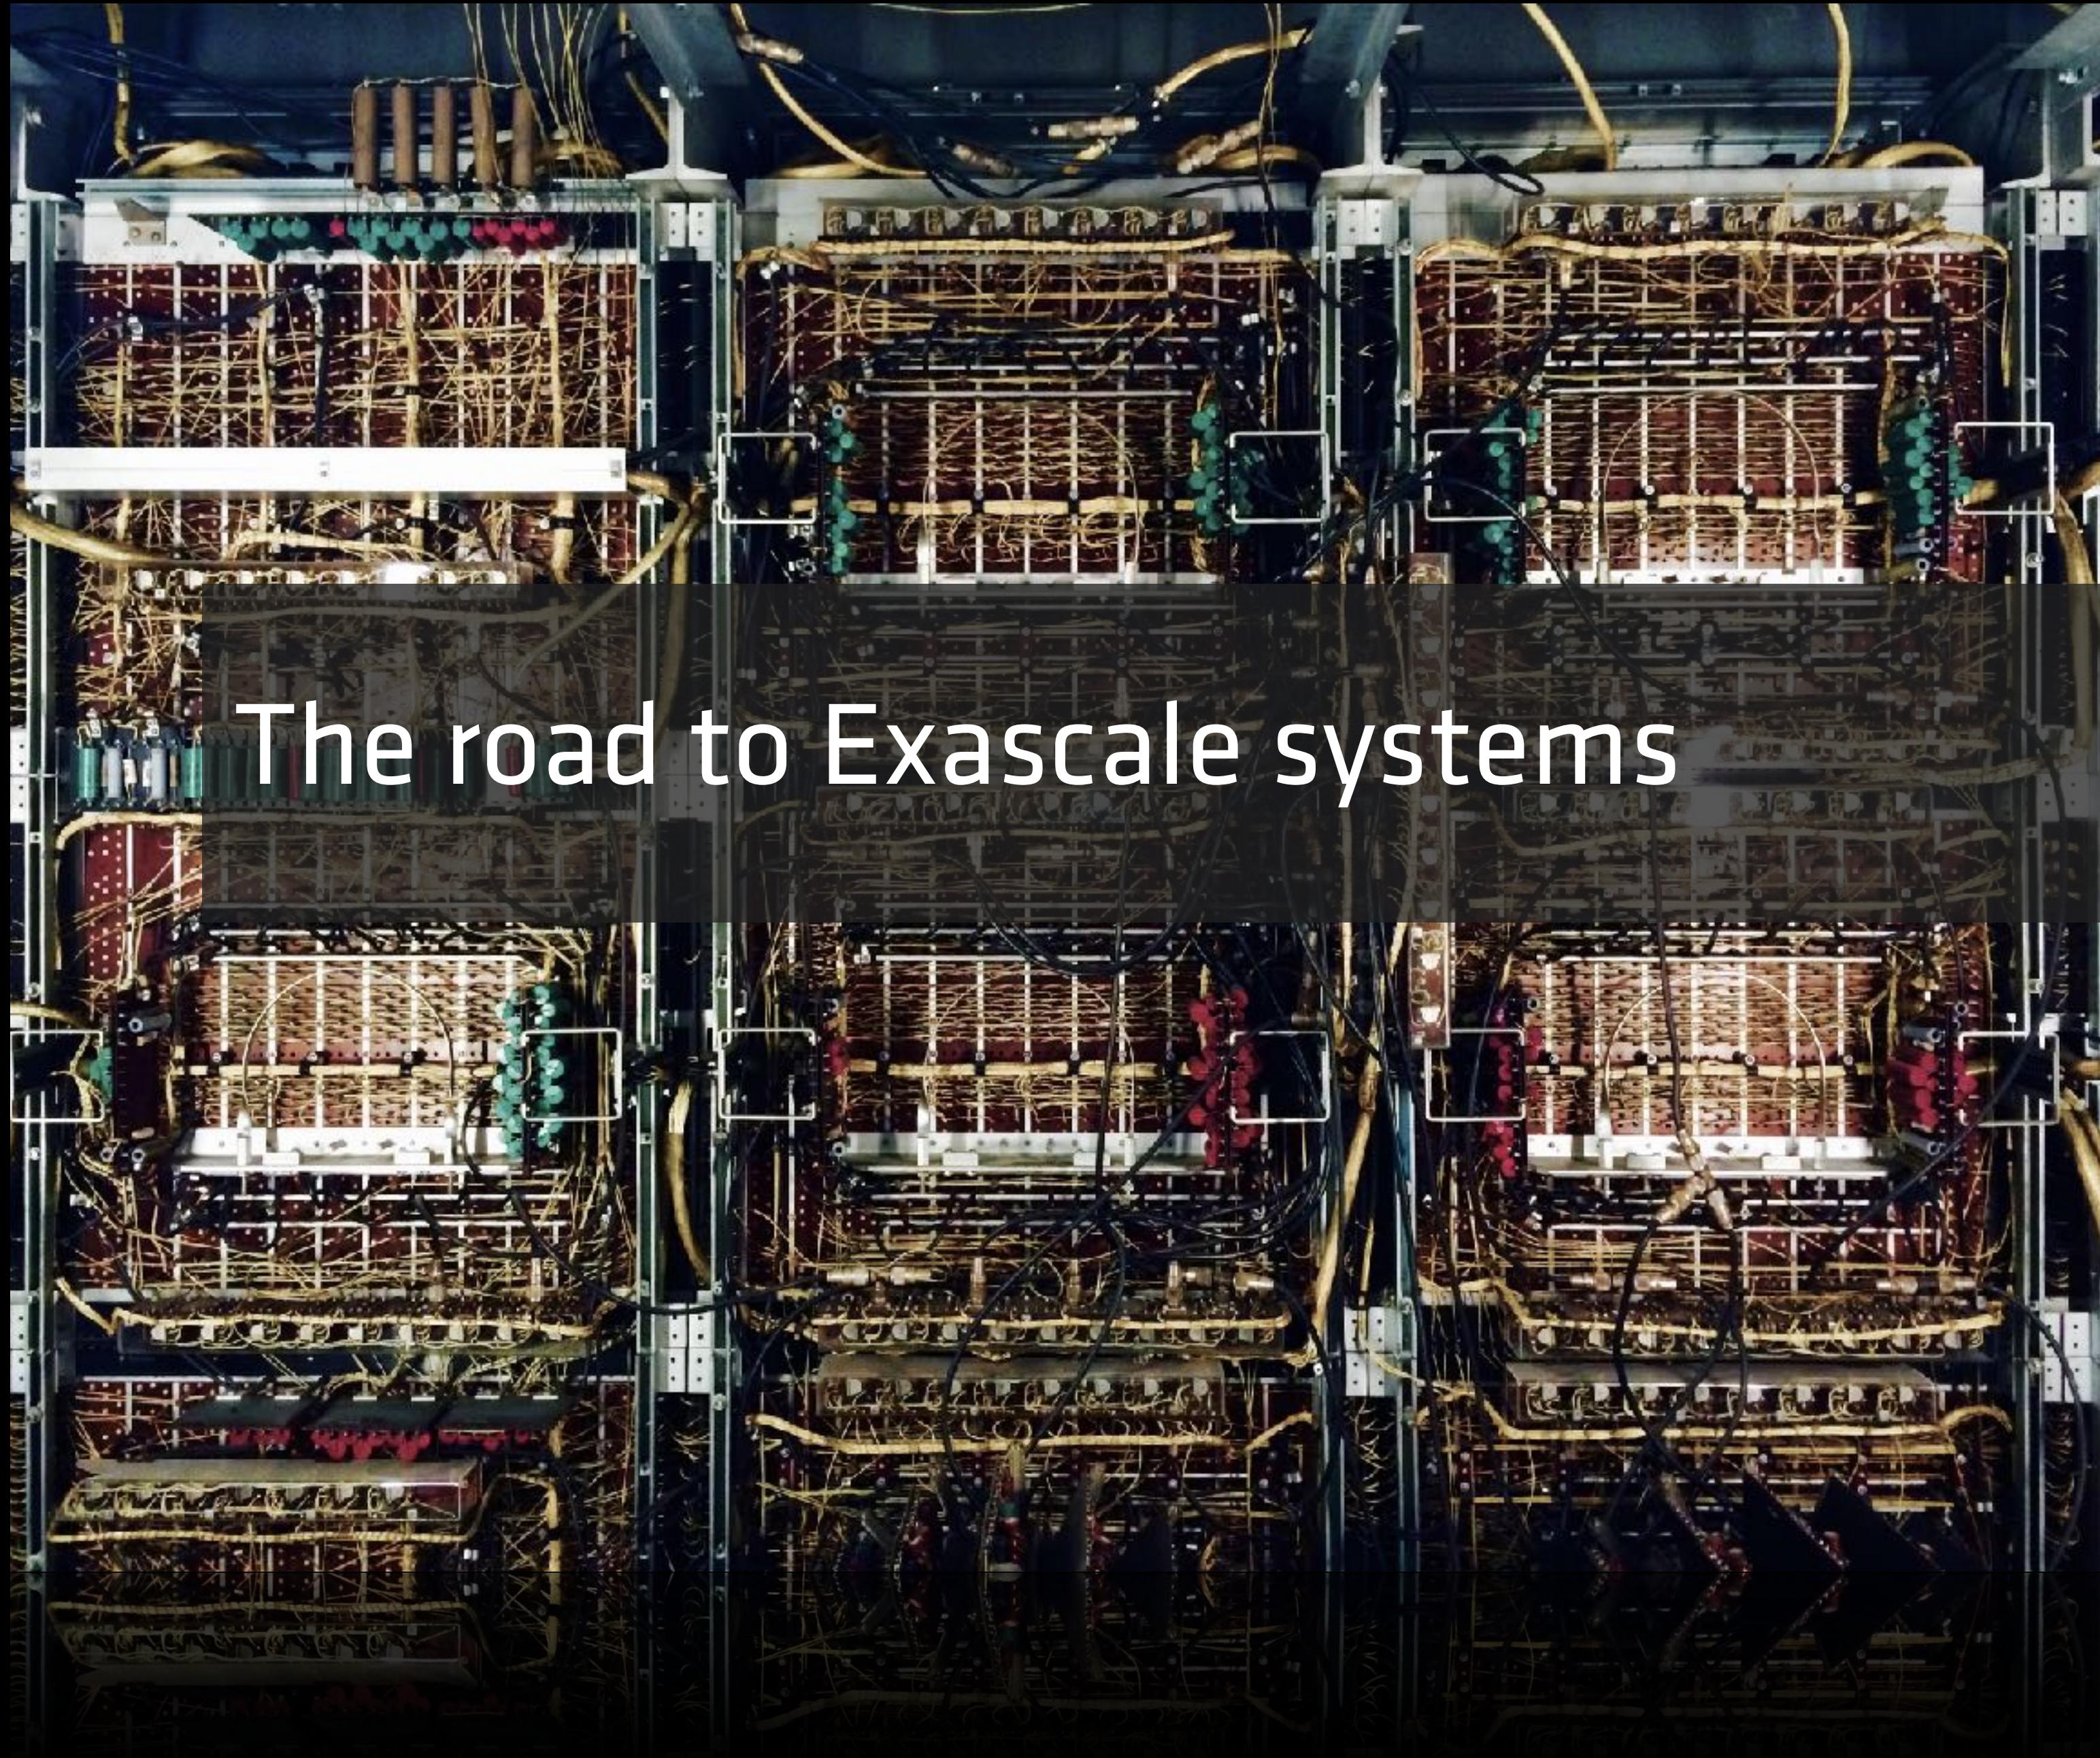
<sup>1</sup> GoLP/IPFN, Instituto Superior Técnico, Lisboa, Portugal

<sup>2</sup> DCTI, ISCTE-Instituto Universitário de Lisboa, Portugal

# Outline

- The road to Exascale Systems
  - HPC system evolution
  - Current trends
  - OSIRIS support for top10 systems
- The new Intel Knights Landing architecture
  - Architecture details
  - OSIRIS on the KNL
- OSIRIS on modern CPUs
  - Overview of modern CPUs
  - Compiling and running for these
- OSIRIS on GPUs





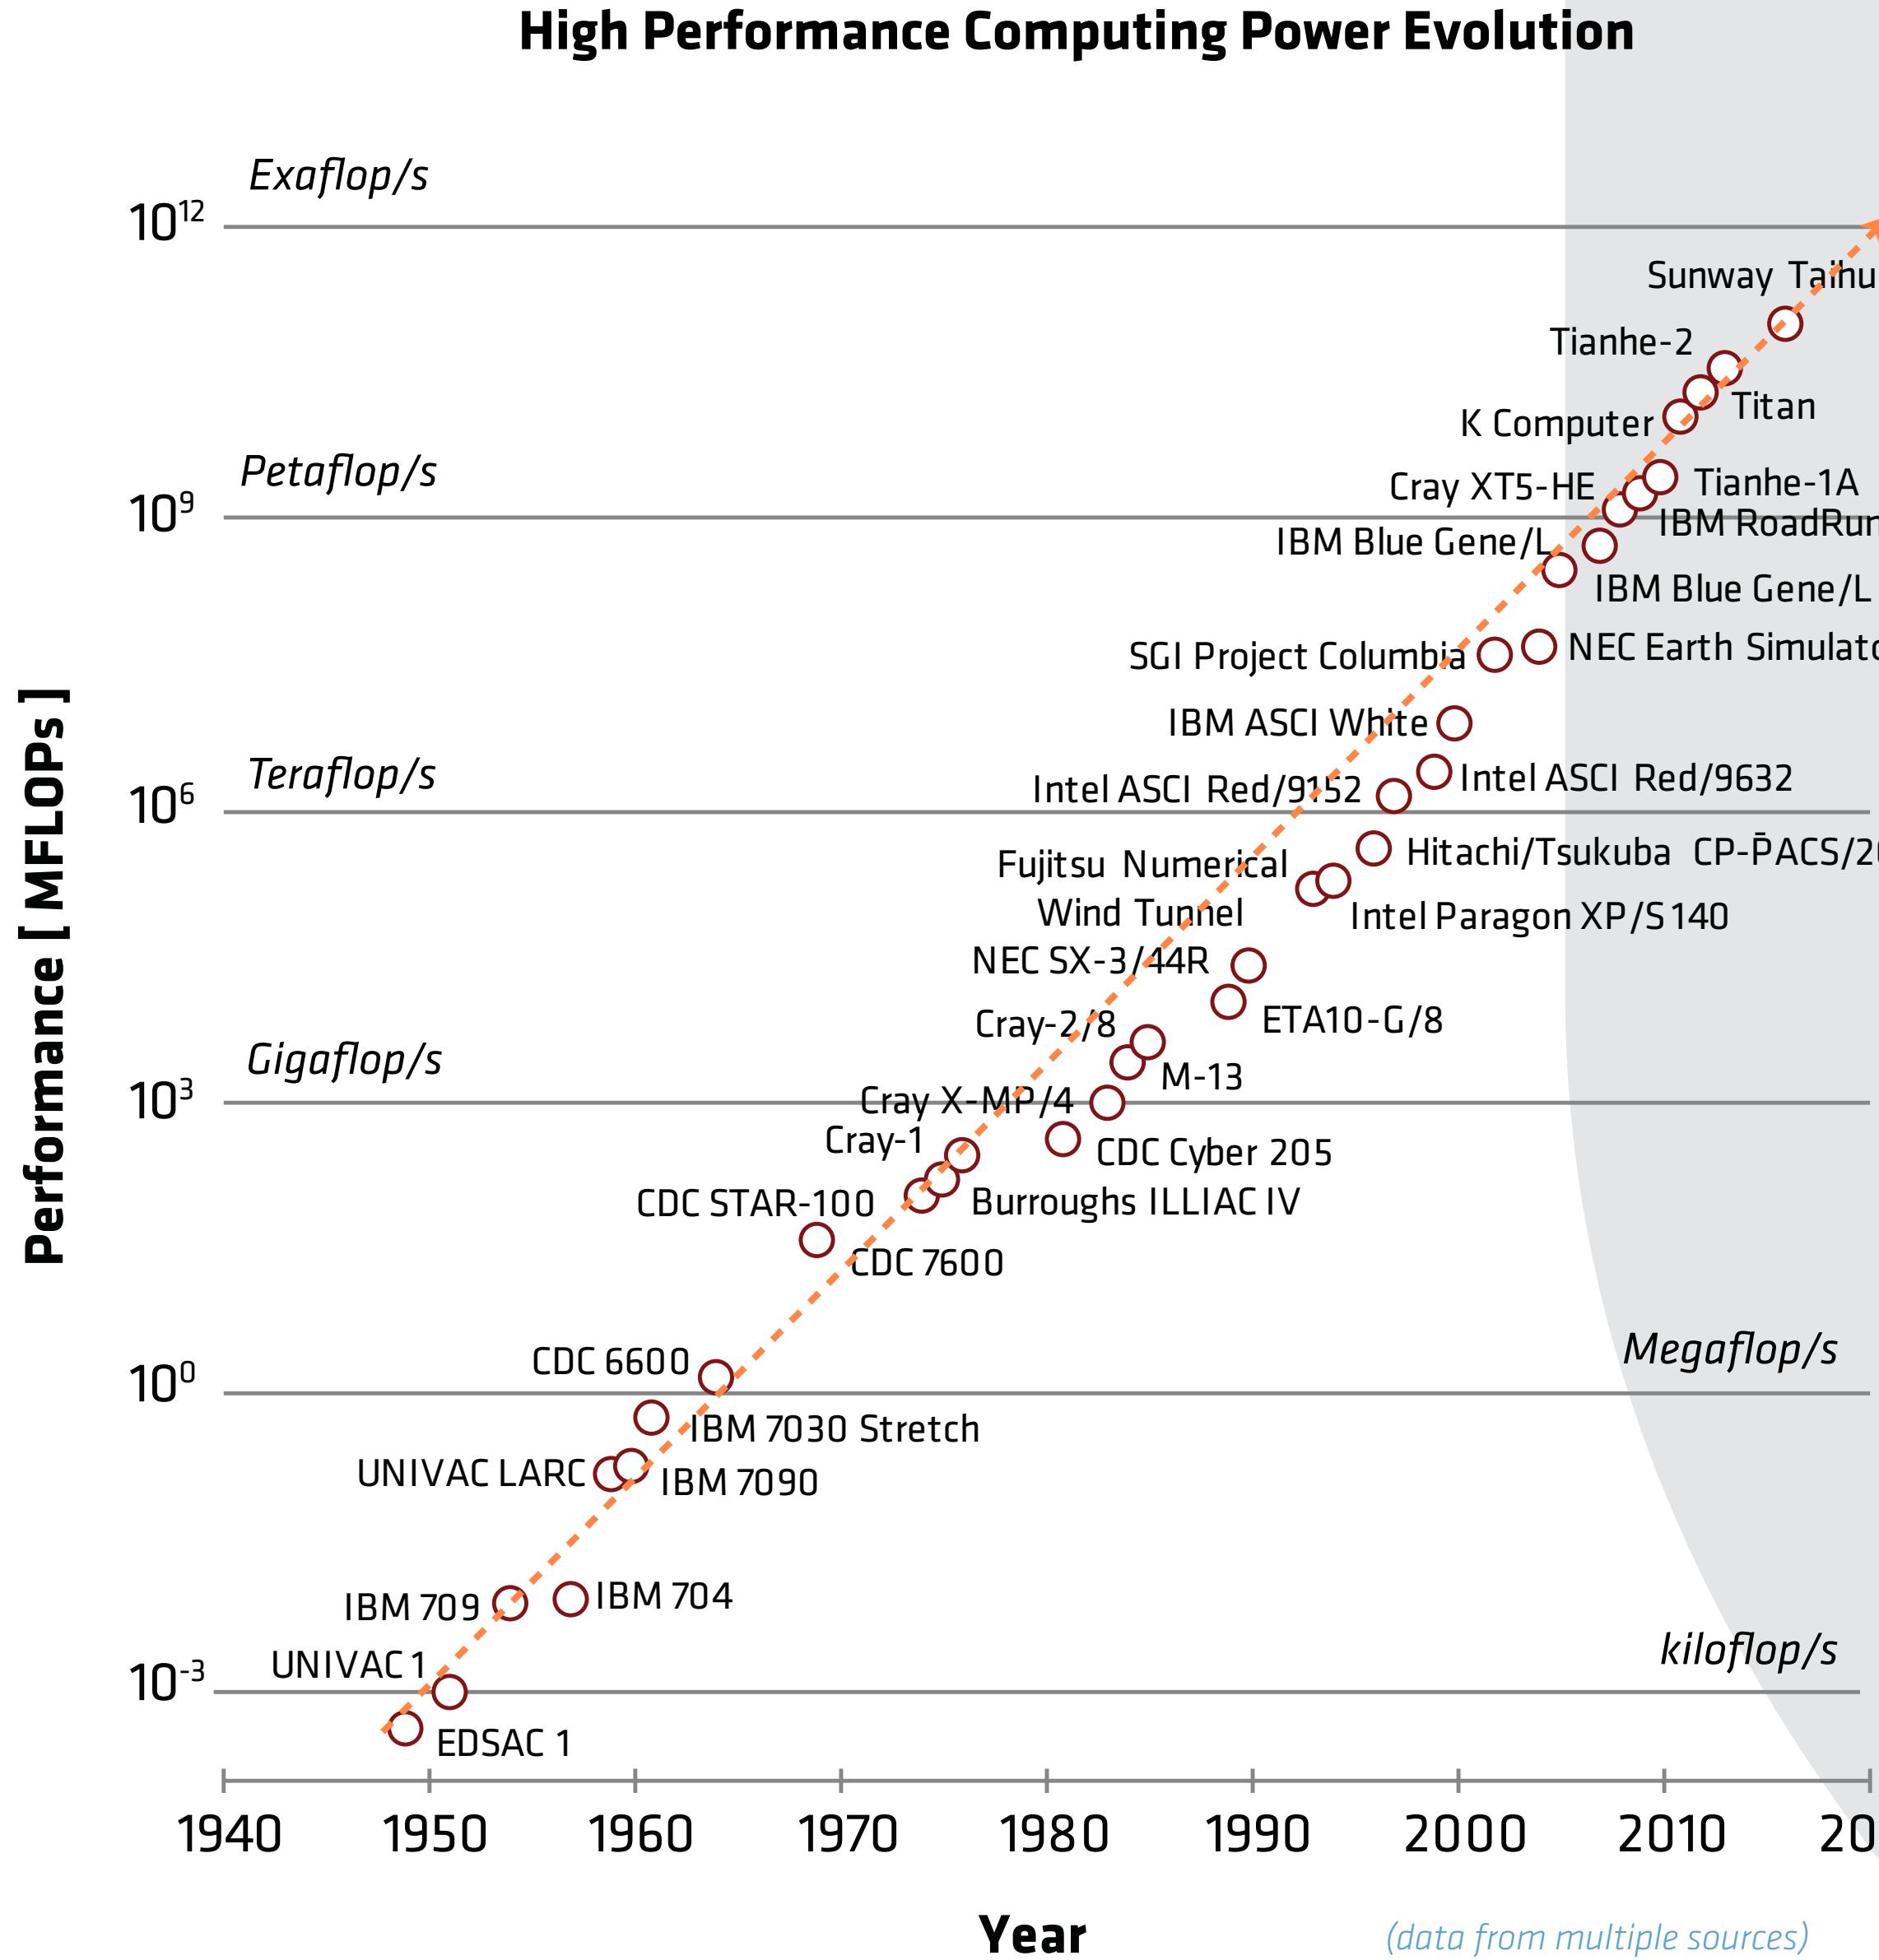
# The road to Exascale systems

**UNIVAC 1 - 1951**

Internal view



# The Road to Exascale Computing



## Sunway Taihulight

- 40 960 compute nodes

## Node Configuration

- 1x SW26010 manycore processor
- 4x(64+1) cores @ 1.45 GHz
- 4x 8 GB DDR3

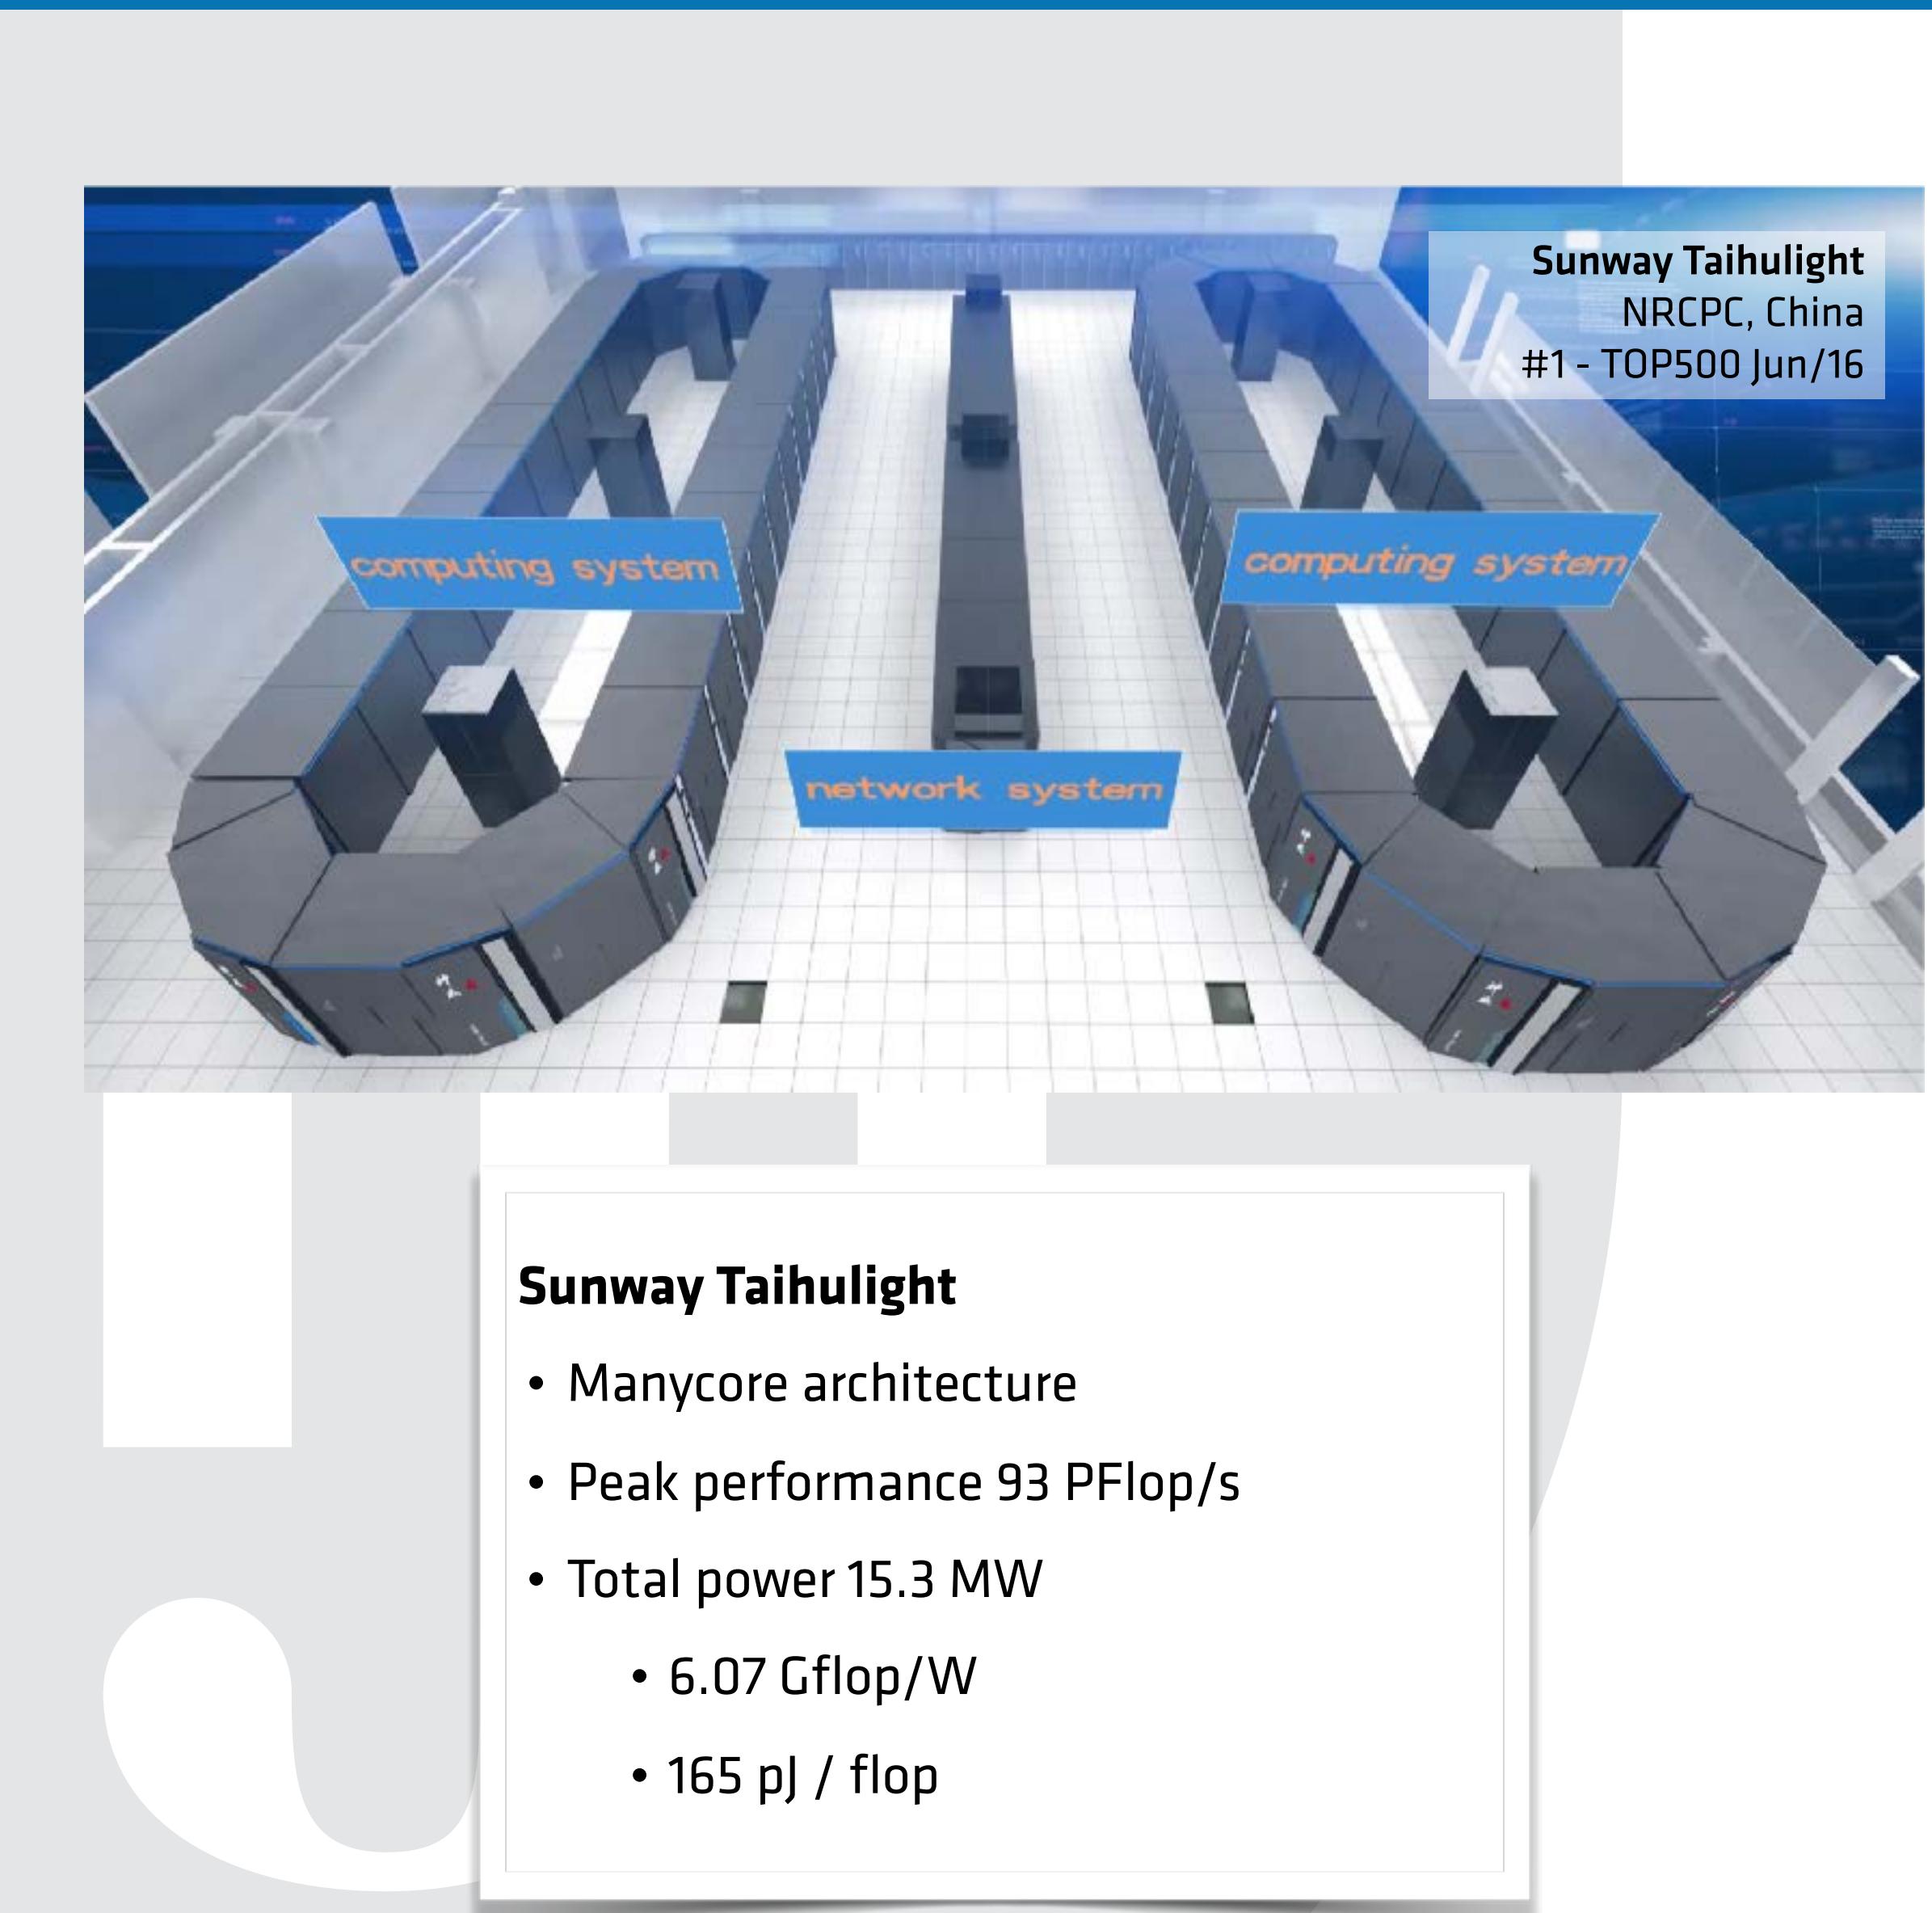
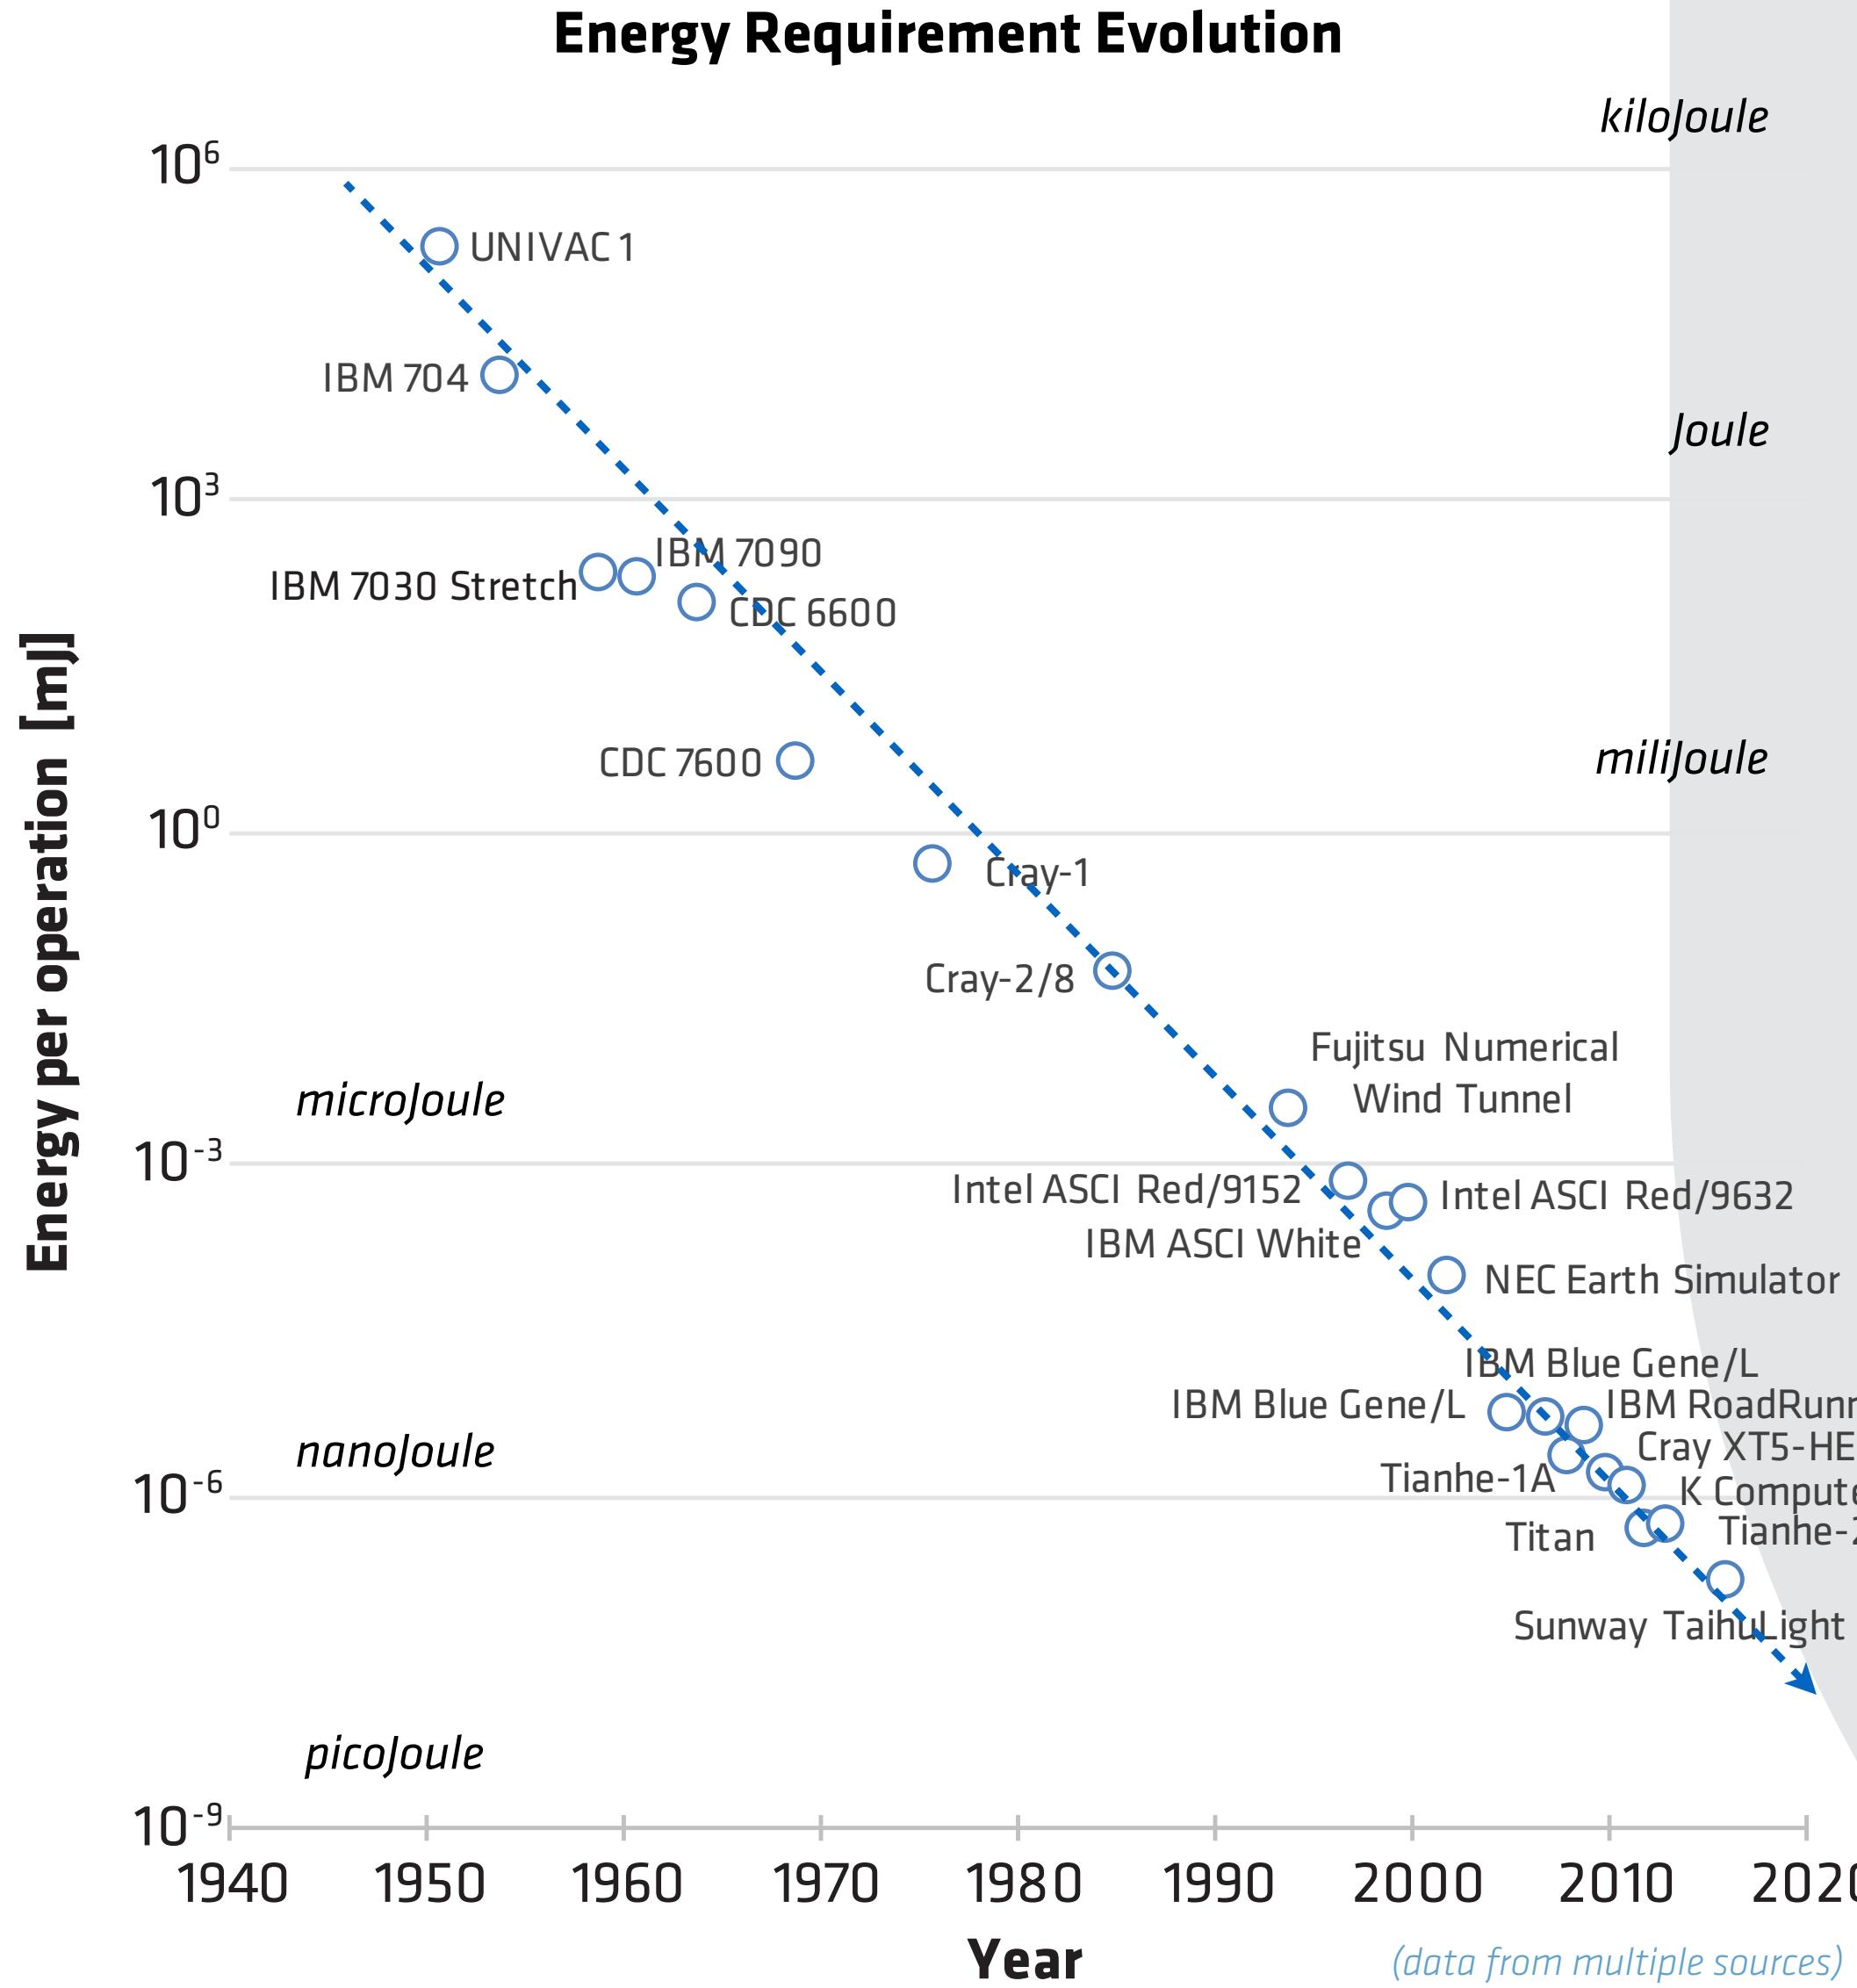
## Total system

- 10 649 600 cores
- 1.31 PB RAM

## Performance

- $R_{\text{peak}}$  125.4 Pflop/s
- $R_{\text{max}}$  93.0 Pflop/s

# The Road to Power Efficient Computing



# Petaflop systems firmly established

## The drive towards Exaflop

- Steady progress for over 60 years
  - 138 systems above 1 PFlop/s
- Supported by many computing paradigm evolutions
- Trend indicates Exaflop systems by next decade
- Electric power is one of the limiting factors
  - Target < 20 MW
  - Top system achieves ~ 6 Gflop/W
    - ~ 0.2 GW for 1 Exaflop
    - Factor of 10× improvement still required
  - Best energy efficiency
    - 14 Gflop/W
    - NVIDIA Tesla P100

### Multicore systems

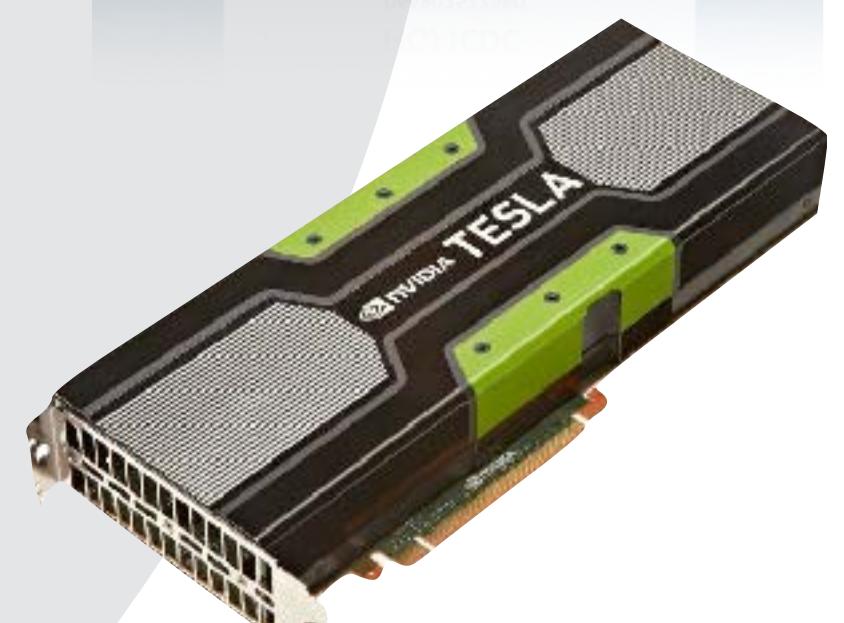
- Maintain complex cores and replicate
- 2 systems in the top 10 are based on multicore CPUs
  - 1x Fujitsu SPARK #8 (K-Computer)
  - 1x Intel Xeon E5 #10 (Trinity)

### Manycore

- Use many lower power cores
- 5 systems in the top 10
  - #1 (Sunway Taihulight)
  - 2x IBM BlueGene/Q in the top
    - #5 (Sequoia) and #9 (Mira)
    - *Seem to be the last of their kind*
  - 2x Intel Knights Landing systems
    - #6 (Cori) and #7 (Oakforest-PACS)

### Accelerator/co-processor technology

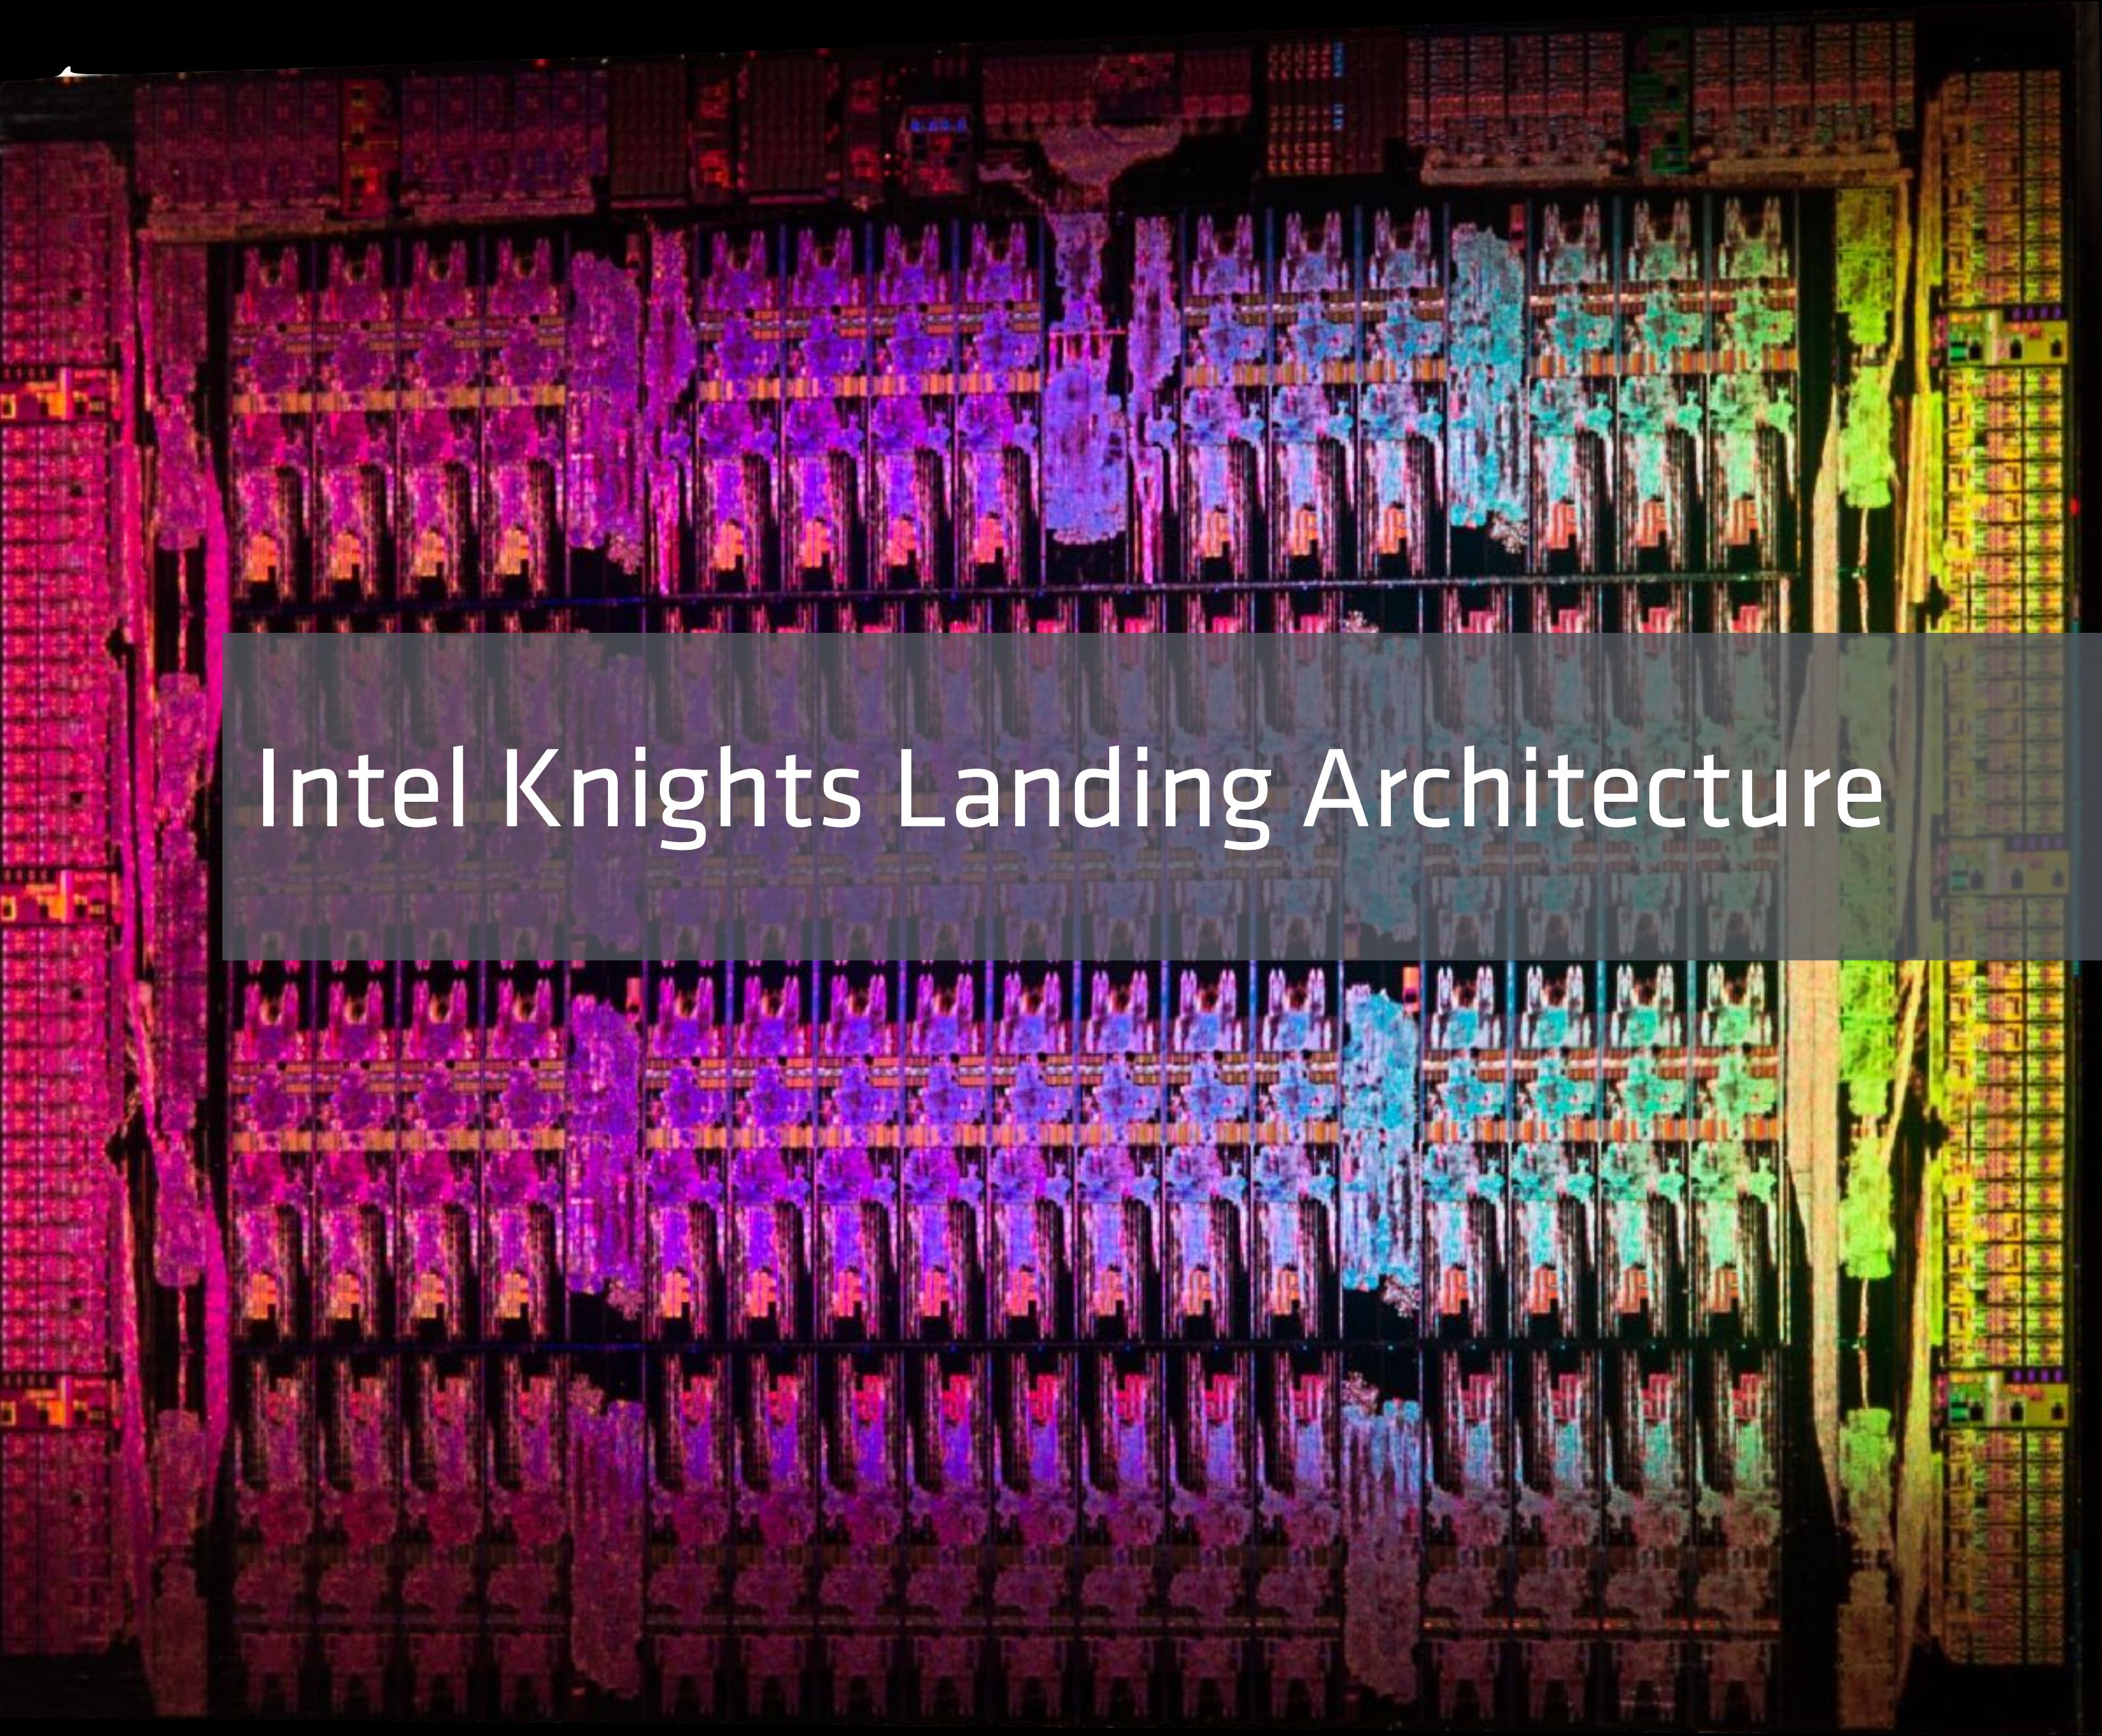
- 3 systems in top 10
  - #3 (Piz Daint) and #4 (Titan) use NVIDIA GPUs
  - #2 (Tianhe-2) uses Intel Knights Corner



# Which top 10 systems are currently supported by OSIRIS?



Rank	System	Cores	R <sub>max</sub> [TFlop/s]	R <sub>peak</sub> [TFlop/s]	Power [kW]	OSIRIS support
1	Sunway TaihuLight, China	10649600	93014.6	125435.9	15371	No
2	Tianhe-2 (MilkyWay-2), China	3120000	33862.7	54902.4	17808	Full (KNC)
3	Piz Daint, Switzerland	361760	19590.0	25326.3	2272	Full (CUDA)
4	Titan, United States	560640	17590.0	27112.5	8209	Full (CUDA)
5	Sequoia, United States	1572864	17173.2	20132.7	7890	Full (QPX)
6	Cori, United States	622336	14014.7	27880.7	3939	Full (KNL)
7	Oakforest-PACS, Japan	556104	13554.6	24913.5	2719	Full (KNL)
8	K computer, Japan	705024	10510.0	11280.4	12660	Standard (Fortran)
9	Mira, United States	786432	8586.6	10066.3	3945	Full (QPX)
10	Trinity, United States	301056	8100.9	11078.9	4233	Full (AVX2)



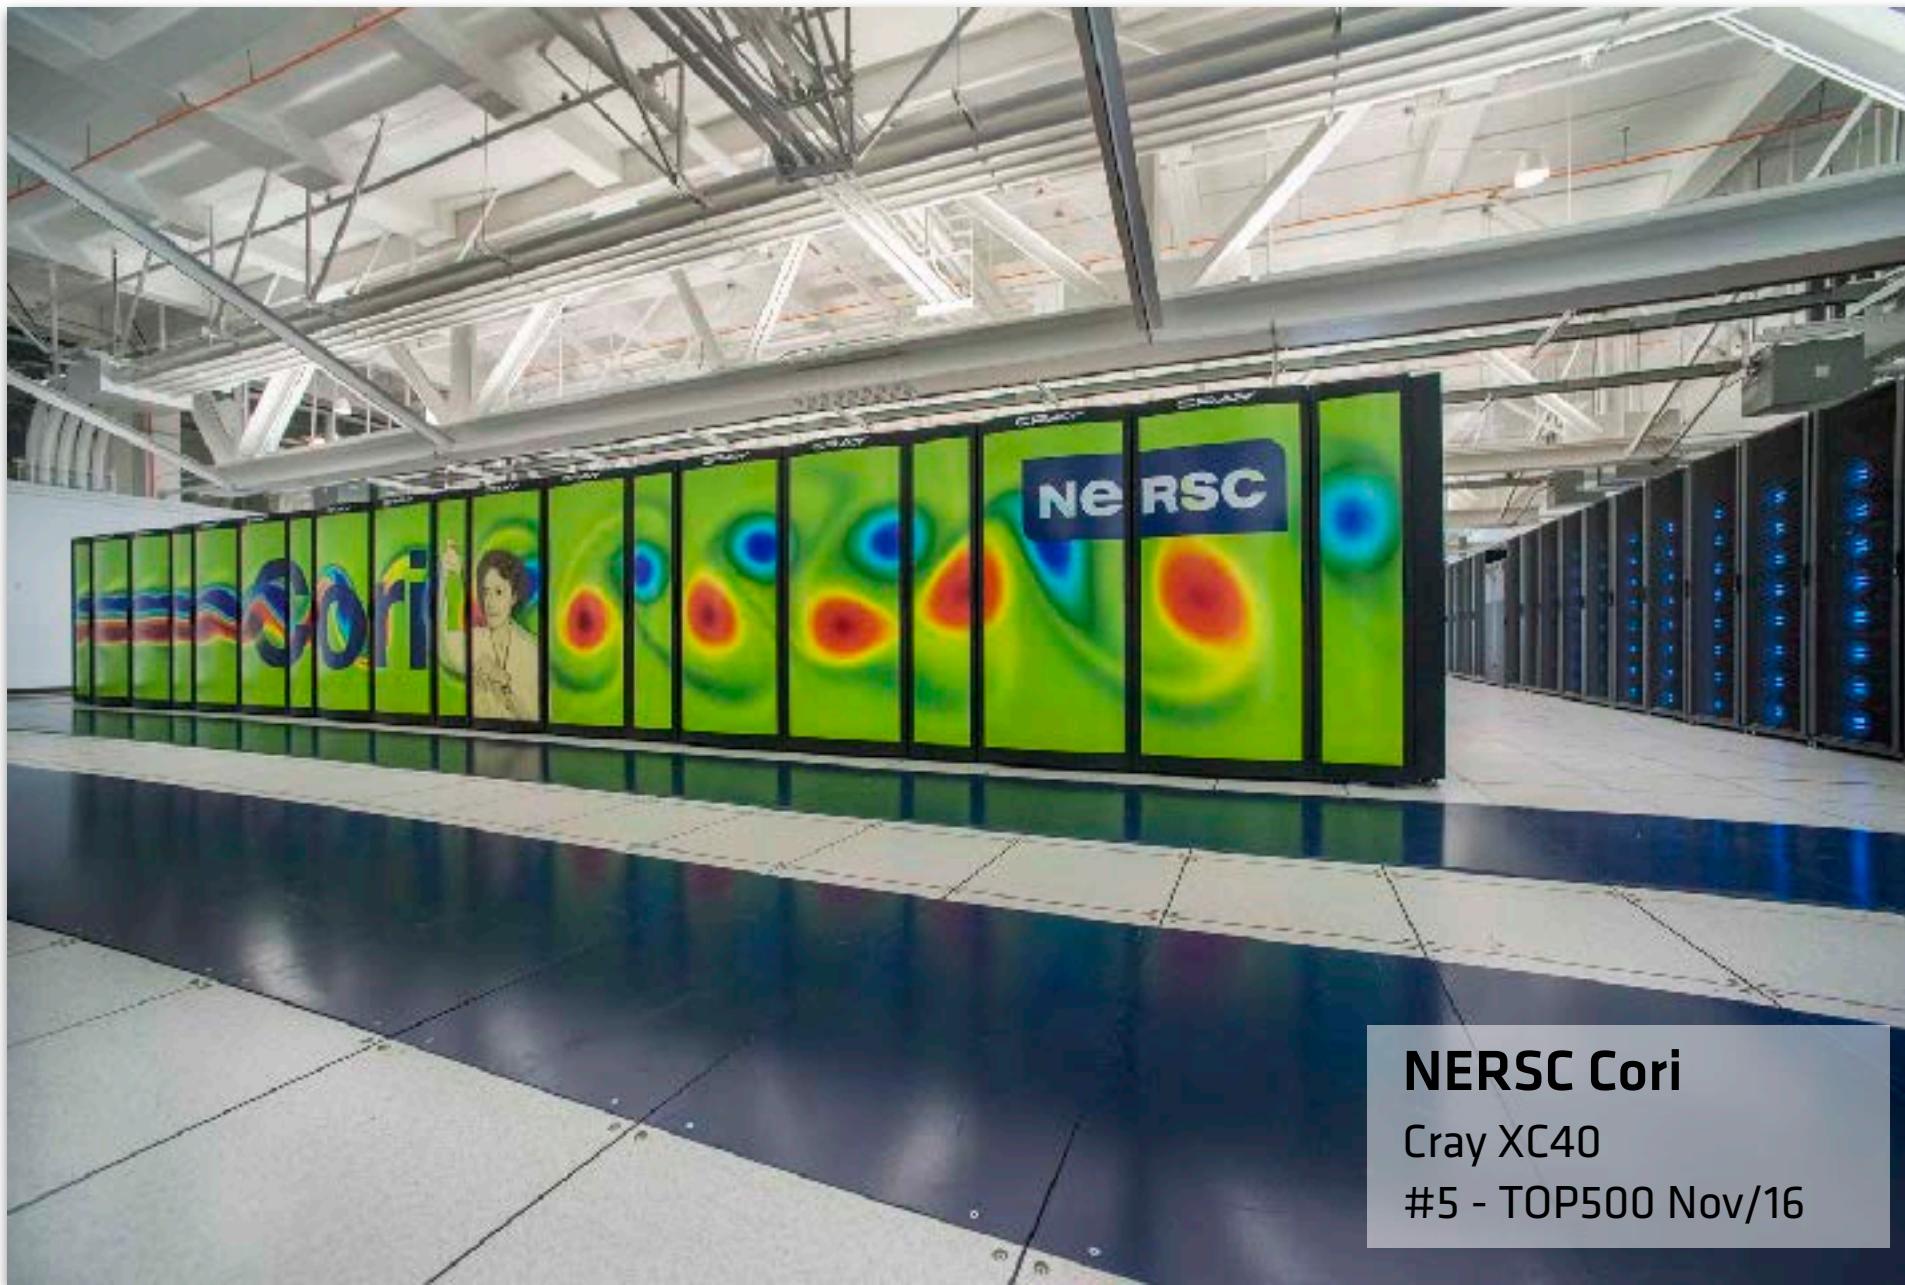
# Intel Knights Landing Architecture



Intel Xeon Phi 5110p die

# Intel Knights Landing Architecture

NERSC Cori



- **Cray XC40**
  - 9152 Compute Nodes
  - 3.9 MW
- **Interconnect**
  - Aries Interconnect
- **Node configuration**
  - 1x Intel® Xeon® Phi 7250 @ 1.4GHz (68 cores)
  - 16 GB (MCDRAM) + 96 GB (RAM)
  -
- **Total system**
  - 622 336 cores
  - 0.15 PB (MCDRAM) + 0.88 PB (RAM)
- **Performance**
  - $R_{MAX} = 14.0 \text{ PFlop/s}$
  - $R_{PEAK} = 27.9 \text{ PFlop/s}$

Intel Knights Landing (KNL)

## • Processing

- 68 x86\_64 cores @ 1.4 GHz
  - 4 threads / core
  - Scalar unit + 2× 512 bit vector unit
  - 32KB L1-I + 32KB L1-D cache
- Organized in 36 tiles
  - Connected by 2D Mesh
  - 1 MB L2 cache
  - 2 cores/tile
- Fully Intel Xeon ISA-compatible
  - Binary compatible with other x86 code
- No L3 cache (MCDRAM can be configured as an L3 cache)

## • Memory

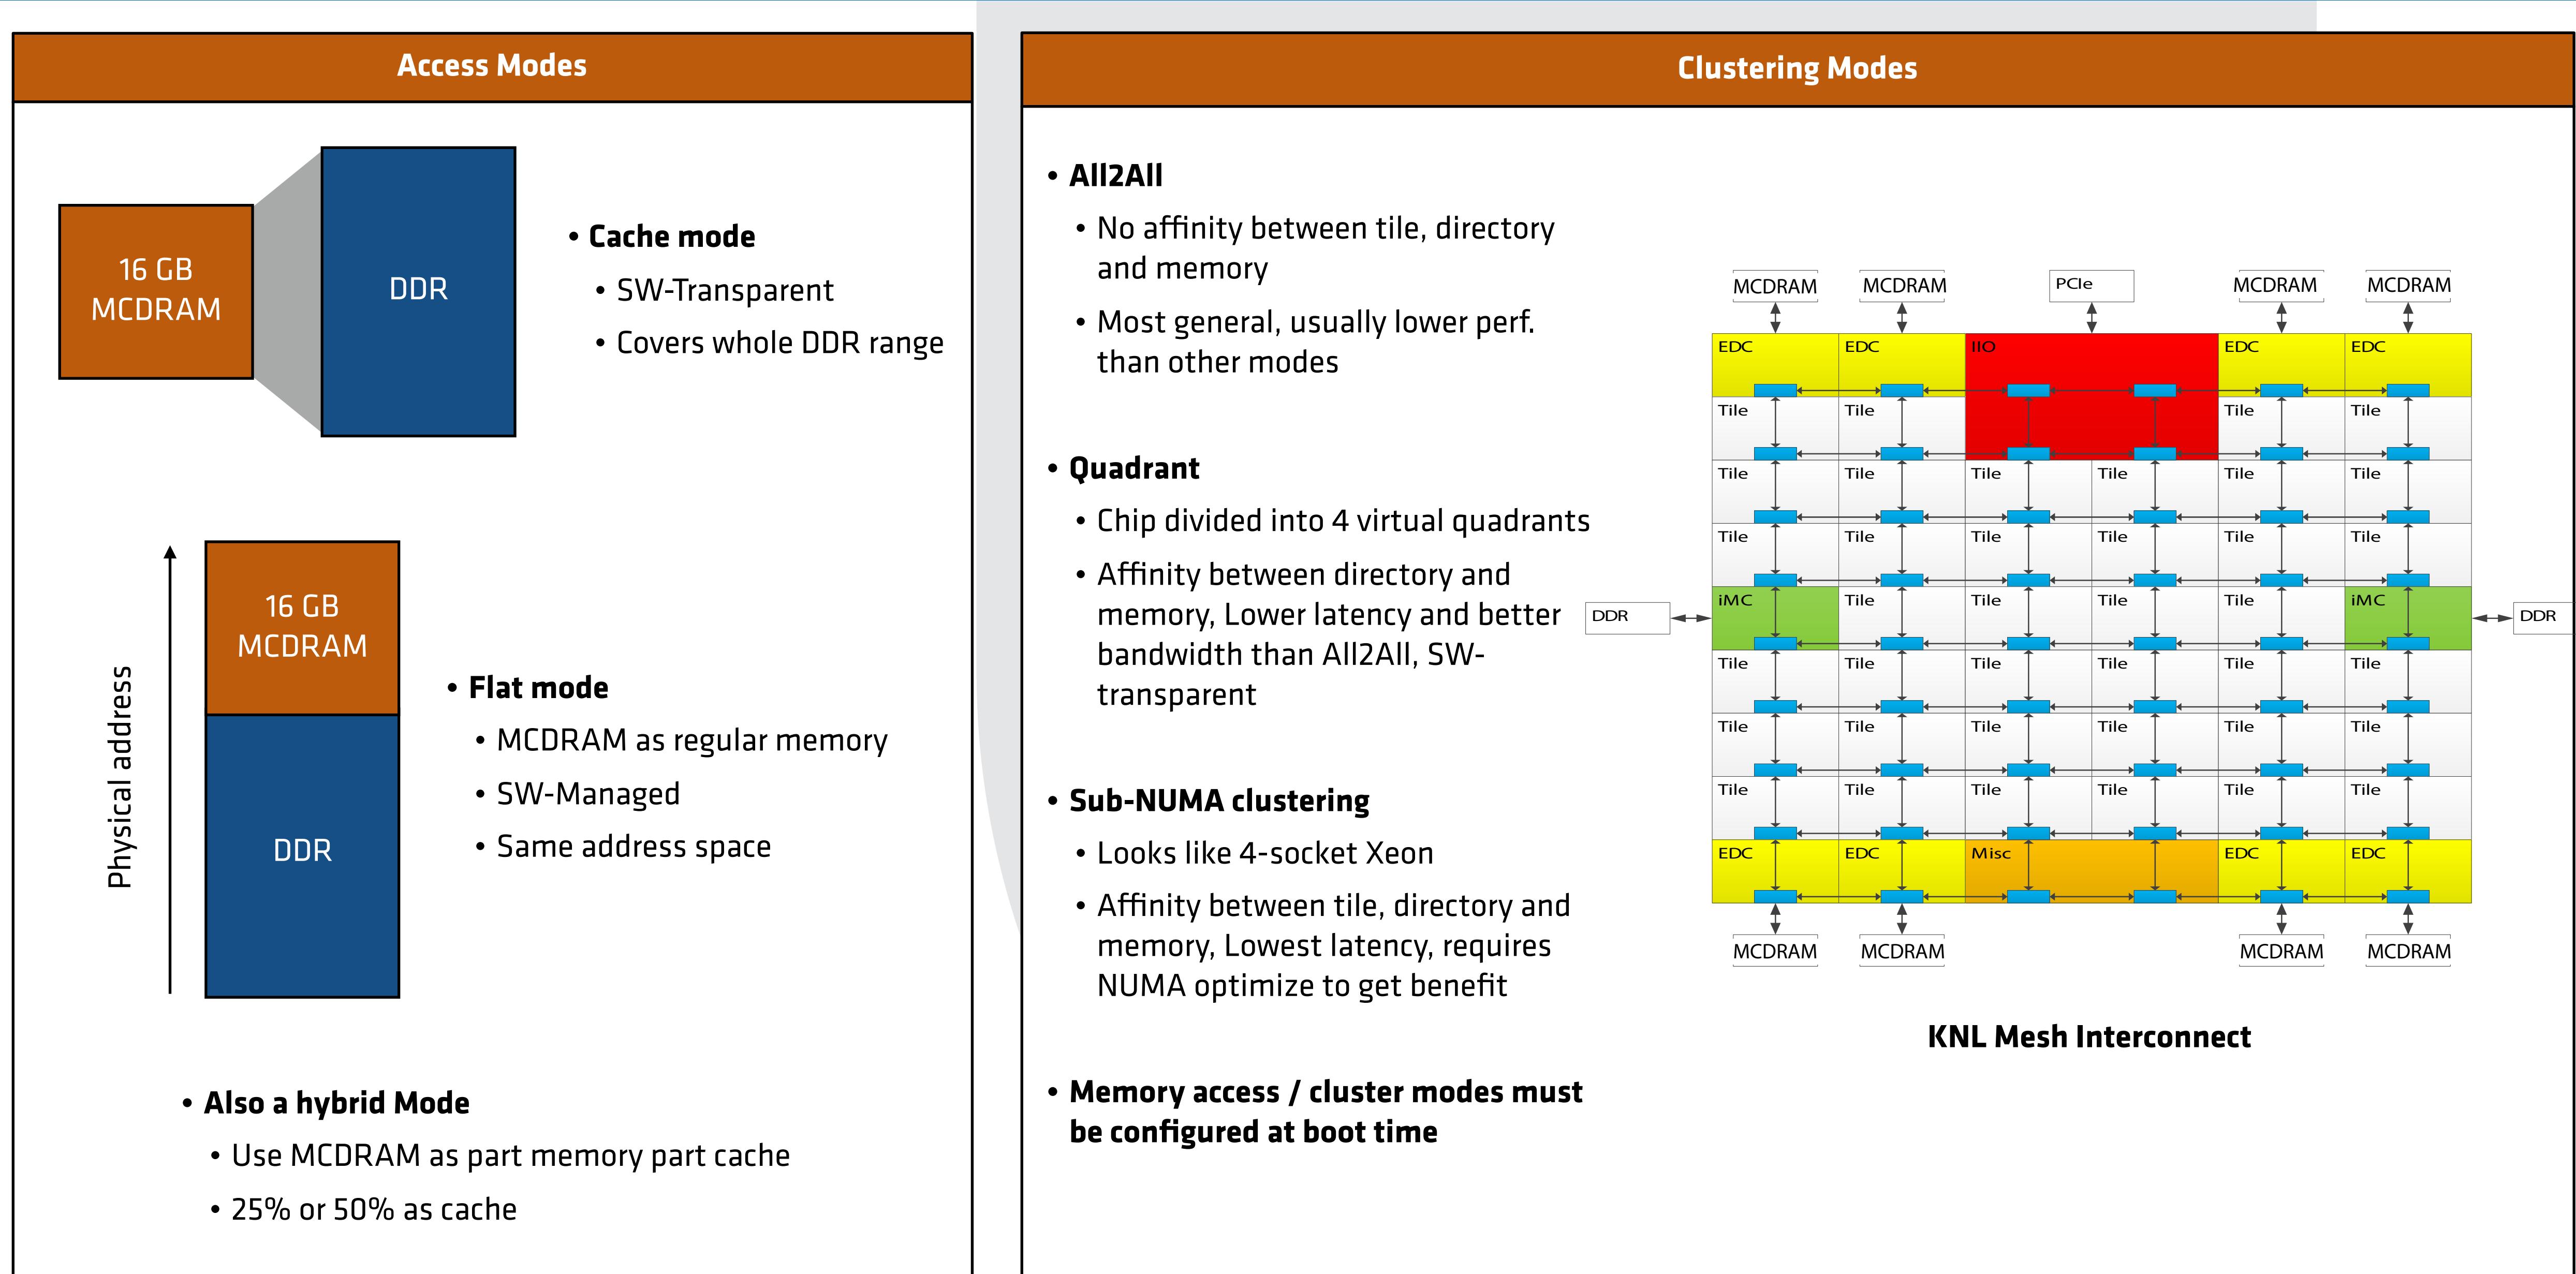
- 16 GB 8-channel 3D MCDRAM
  - 400+ GB/s
- (Up to) 384 GB 6-channel DDR4
  - 115.2 GB/s
- 8×2 memory controllers
- Up to 350 GB/s

## • System Interface

- Standalone host processor
- Max power consumption 215W



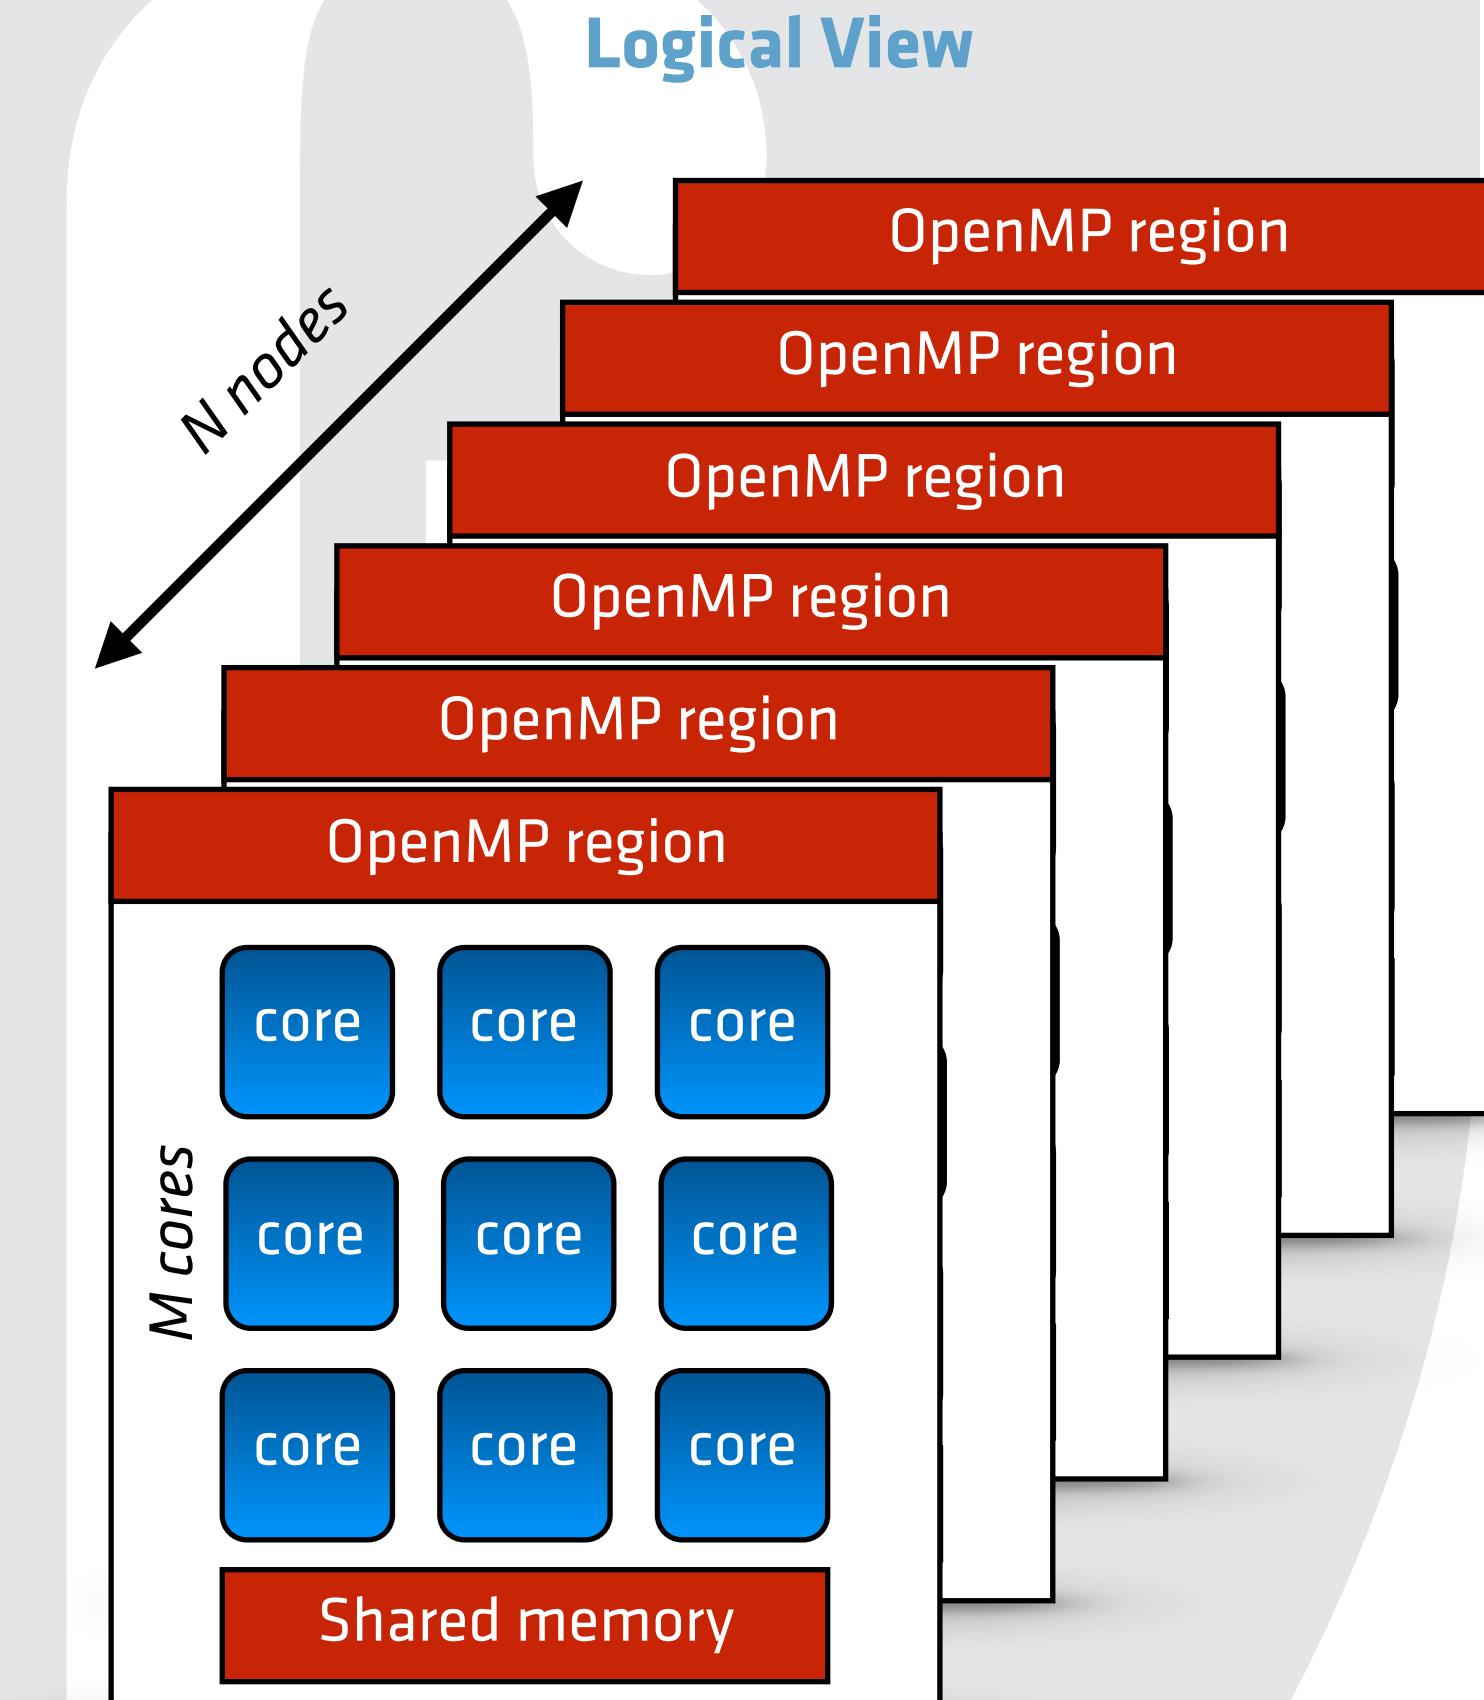
# KNL Memory Architecture



# Logical view of a KNL node

## Intel KNL systems

- The KNL node is simply a 68 core (4 threads per core) shared memory computer
  - Each core has dual 512 bit VPUs
- This can be viewed as small computer cluster with **N SMP nodes** with **M cores** per node
  - Use a distributed memory algorithm (**MPI**) for parallelizing across nodes
  - Use a shared memory algorithm (**OpenMP**) for parallelizing inside each node
- The exact same code used in "standard" HPC systems can run on the KNL nodes
- The vector units play a key role in system performance
  - Vectorization for this architecture must be considered

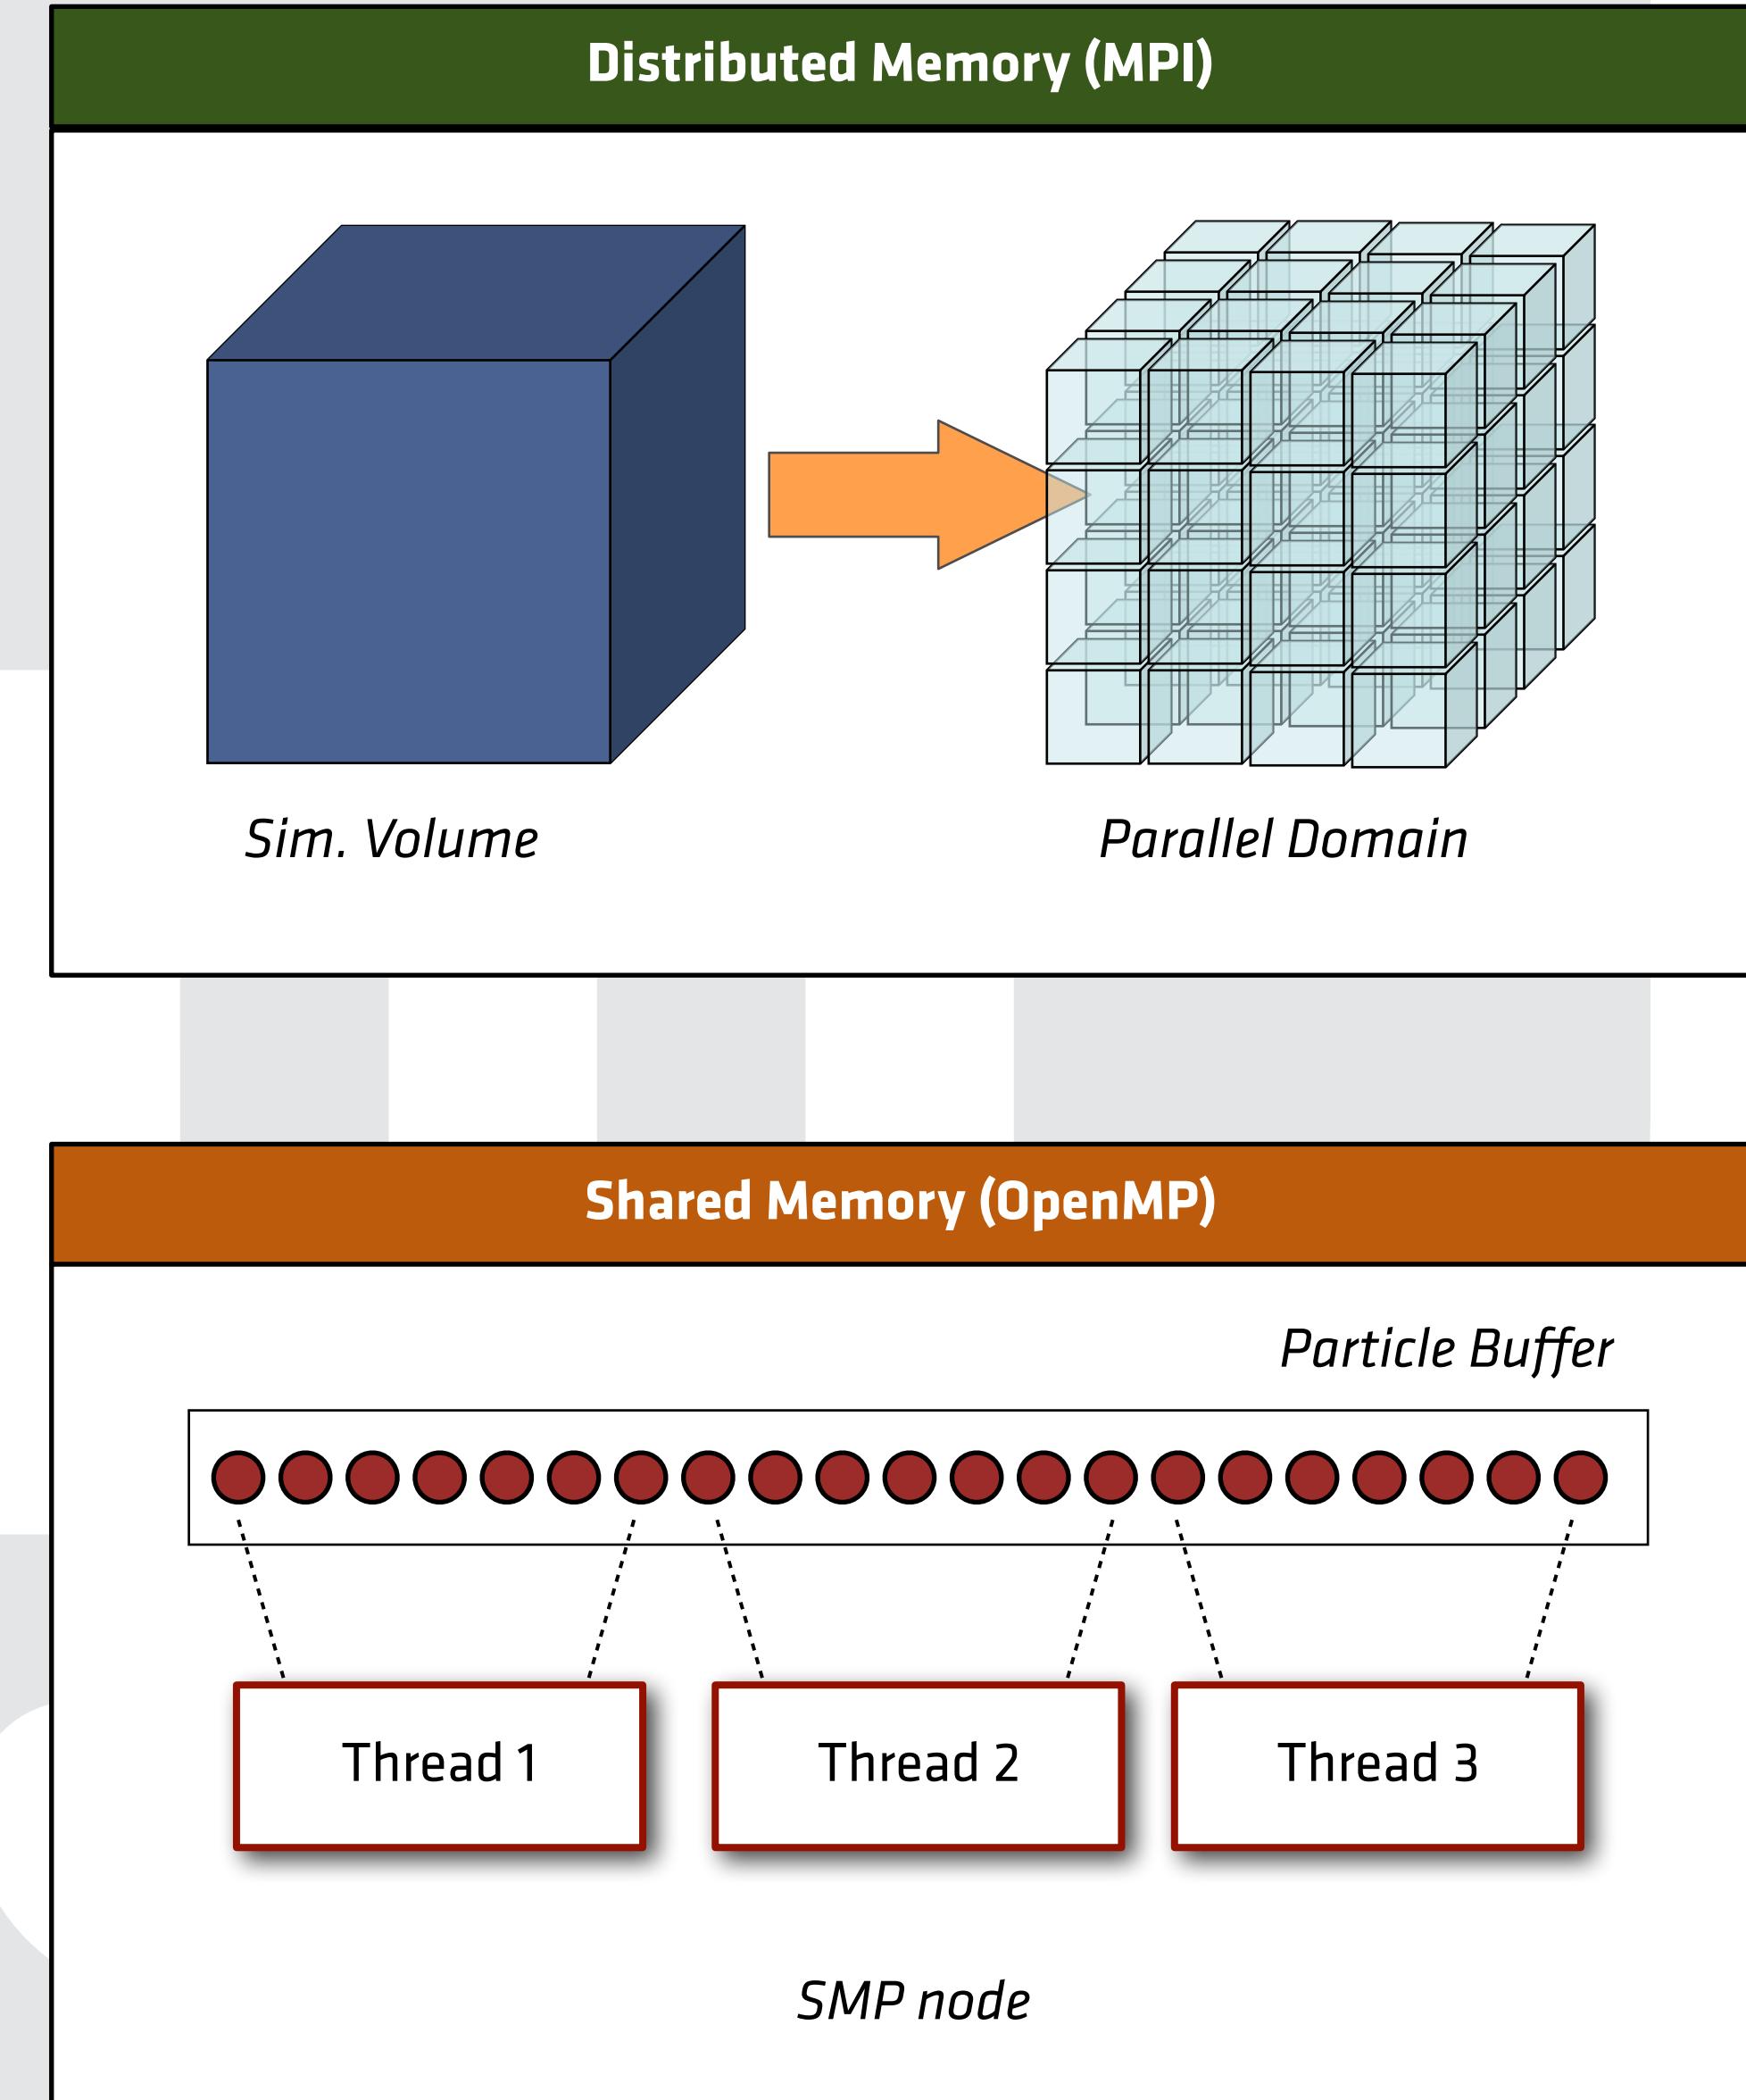


Intel KNL node

# Parallelization for a KNL node

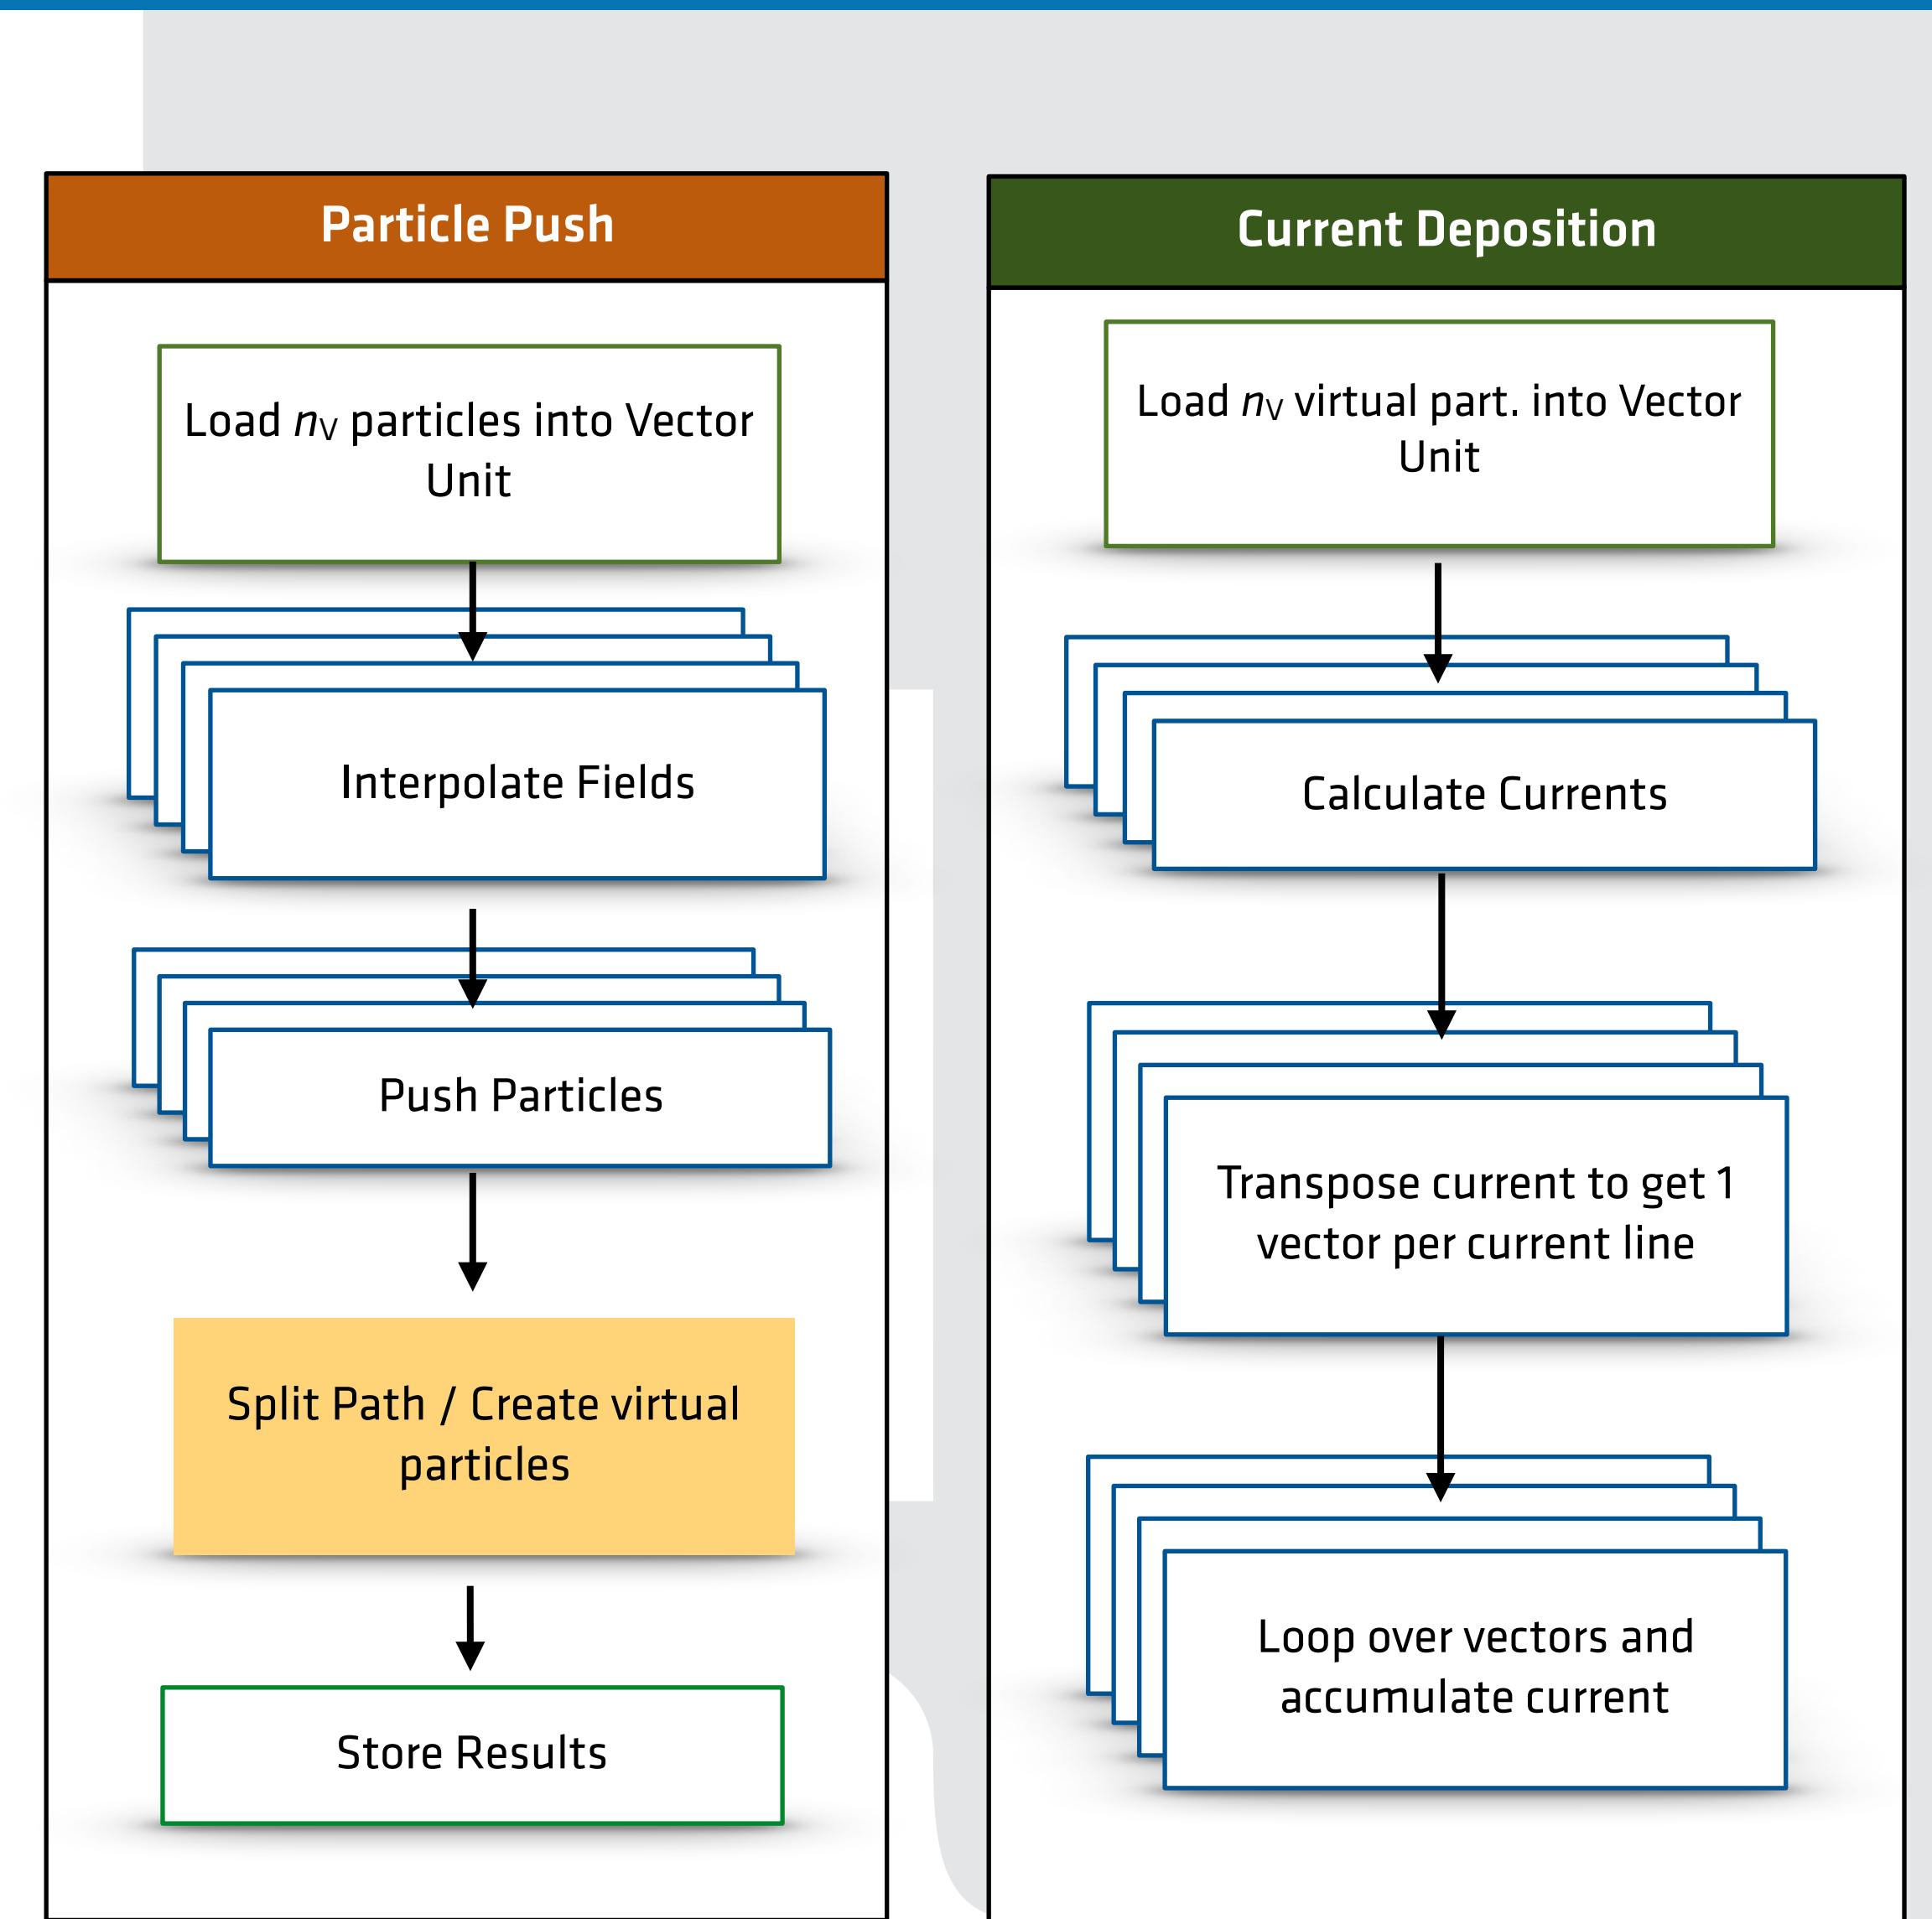


- HPC systems present a hierarchy of distributed / shared memory parallelism
- Top level is a network of computing nodes
  - Nodes exchange information through network messages
- Each computing node is a set of CPUs / cores
  - Memory can be shared inside the node
- Parallelization of PIC algorithm
  - Spatial decomposition across nodes: each node handles a specific region of simulation space
  - Split particles over cores inside node
  - *Same as for standard CPU systems*



# Vectorization of the PIC algorithm for KNL

- The code spends ~80-90% time advancing particles and depositing current
  - Focus vectorization effort here
- Previous strategy developed for vectorization works well
  - Process  $n_v$  (16) particles at a time for most of the algorithm
- For current deposition vectorize by current line to avoid serialization
  - Critical because of large vector width
  - Use a vector transpose operation



## Vectorized by particle

- Vectors contain same quantity for all particles

## Transpose vectors

- Done through vector operations

## Vectorized by current line

- Vectors contain 1 line of current for 1 particle

# Performance on the KNL

## Warm plasma tests

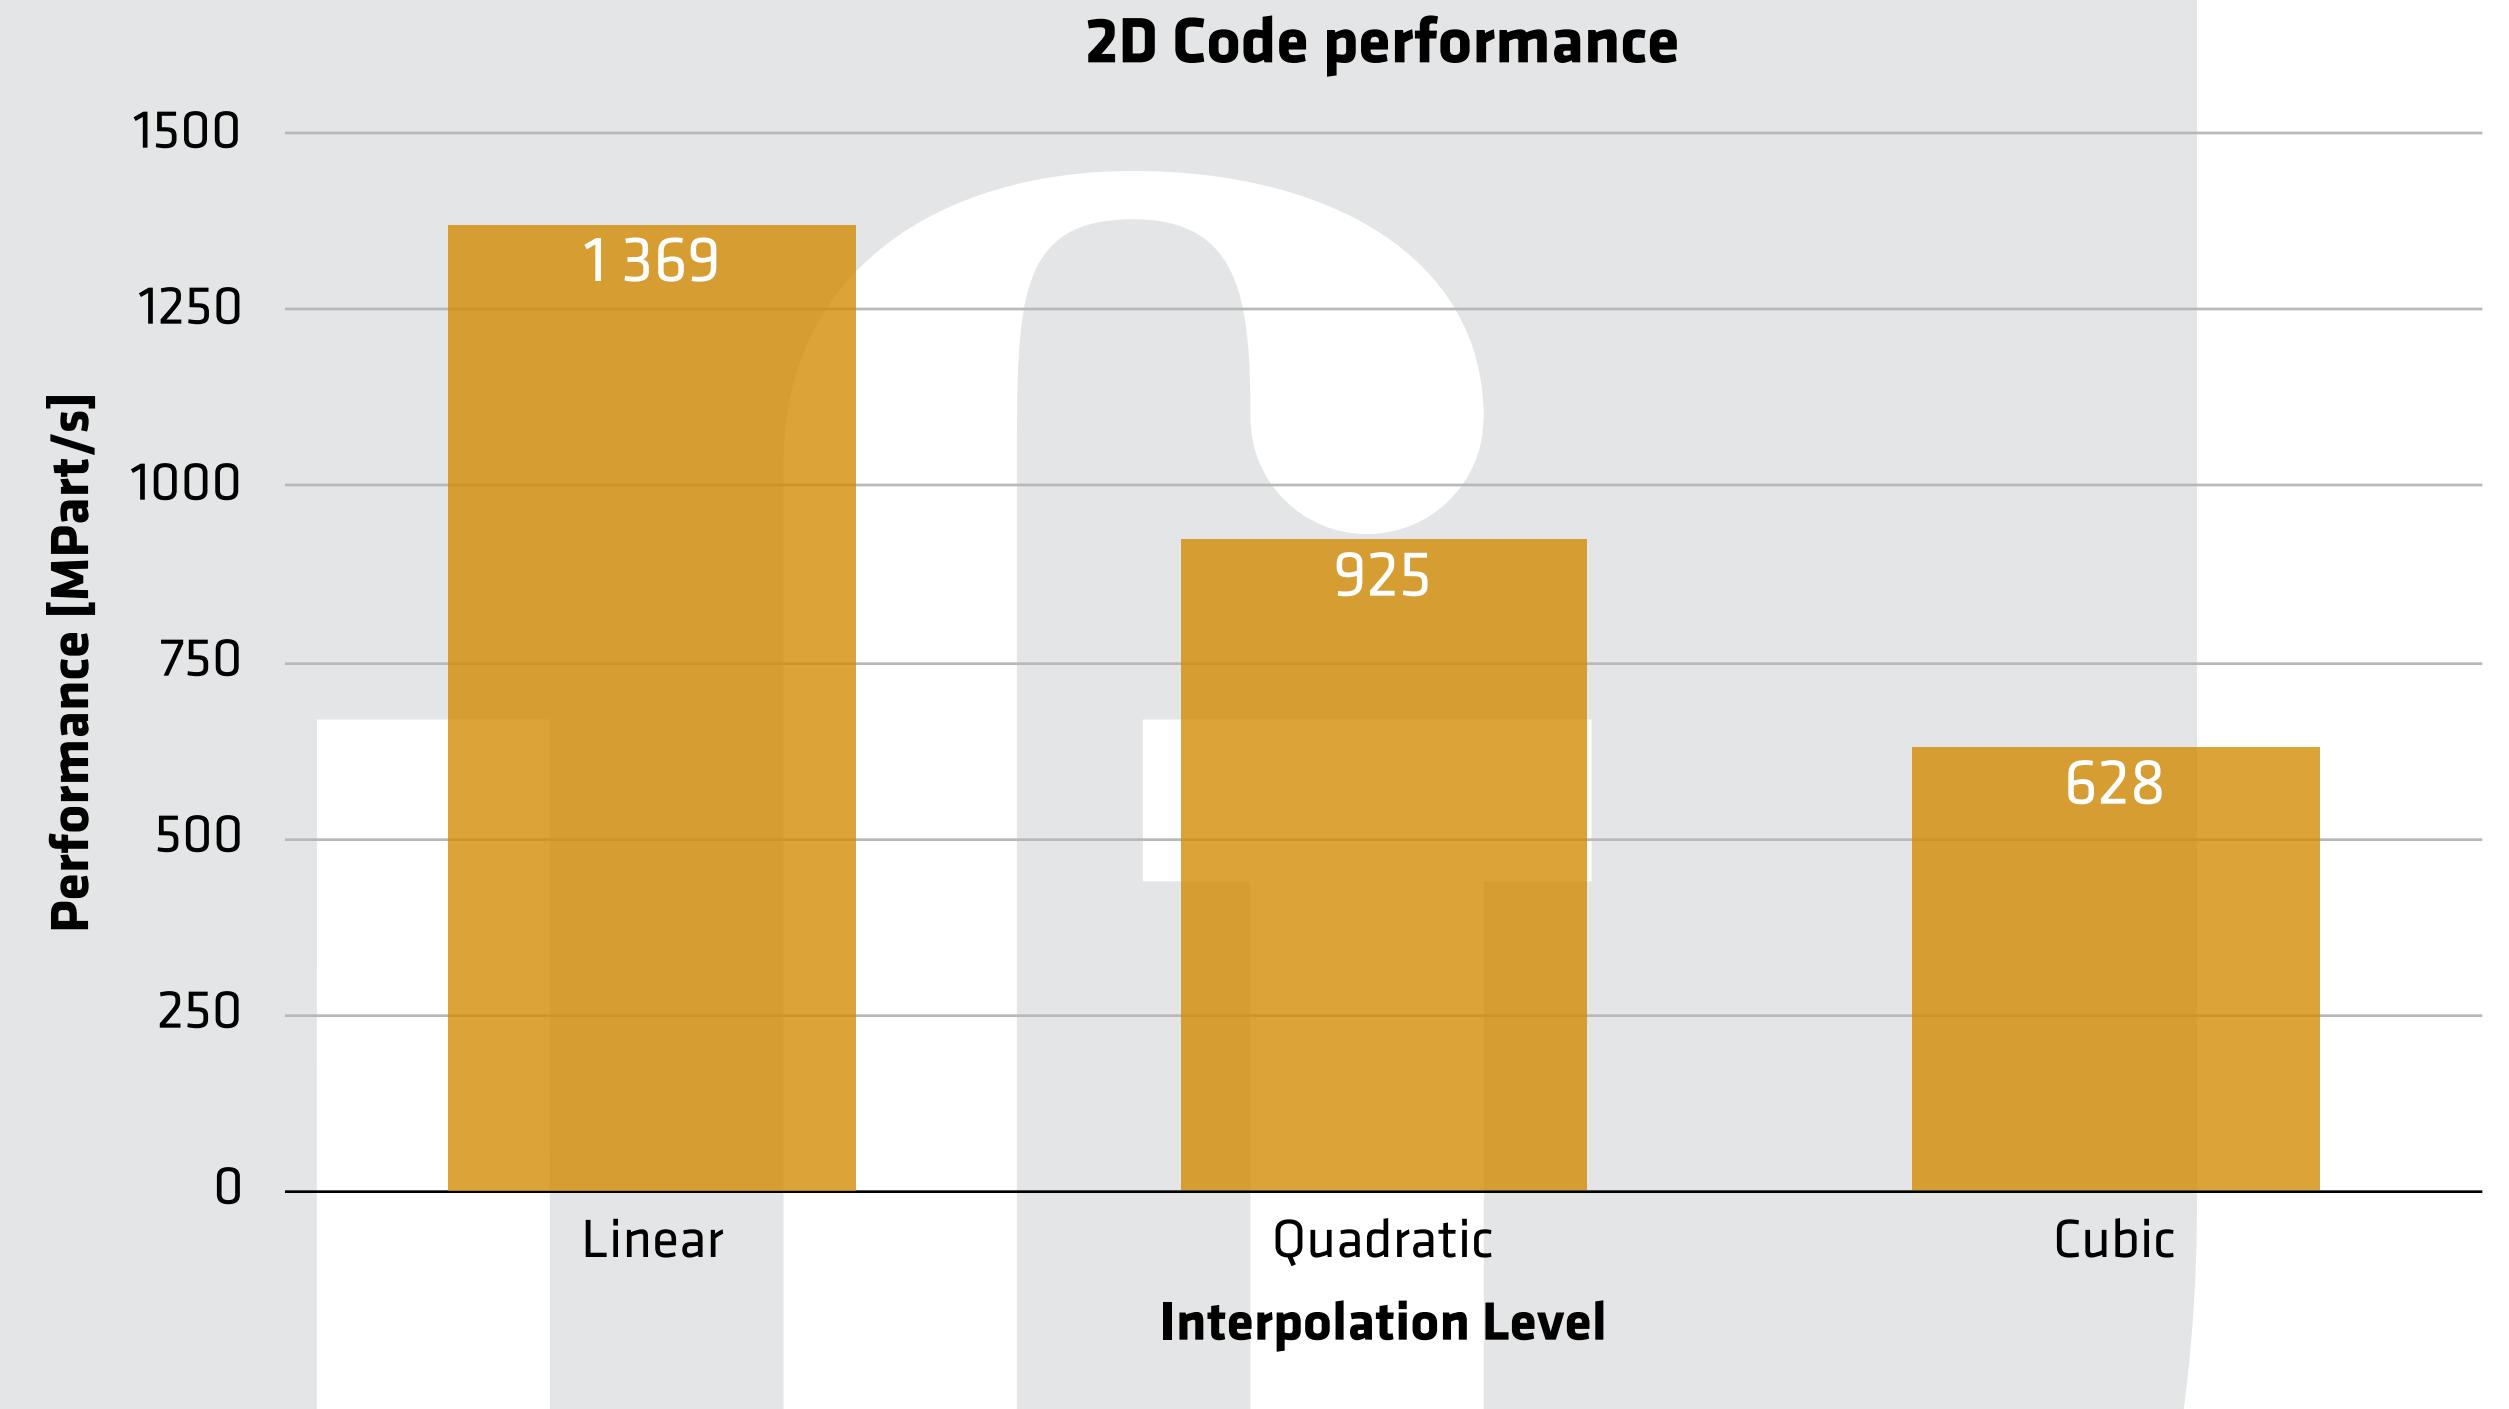
- 2D Problem
- $3400 \times 1600$  cells
- $340 \times 160 (c/\omega_p)^2$  box
- 25 particles / cell
- 136M particles total

## Configuration

- AVX-512 manual vectorization using intrinsics
- No OpenMP parallelism
- 272 MPI processes
  - 4 processes/core
- MCDRAM configuration
  - Flat / SNC4
  - Full problem in MCDRAM

## Performance

- Measured as number of particle pushes per second
- Measurement includes all the simulation loop, including the field solver and communications



### Test system software

- CentOS 7.3
- Intel(R) Parallel Studio XE 2017 Update 3 for Linux
- Intel(R) MPI Library for Linux\* OS, Version 2017 Update 2

### Compiler Flags

```
% mpiifort -xmic-avx512 -O3 -no-prec-div -fp-model fast=2 -fma -align array64byte \
-finline-functions
% mpiicc -xmic-avx512 -O3 -no-prec-div -fp-model fast=2 -ansi-alias -fma -align \
-finline-functions
```

# Processes per core / Hyperthreading

## KNL cores support 4 threads per core

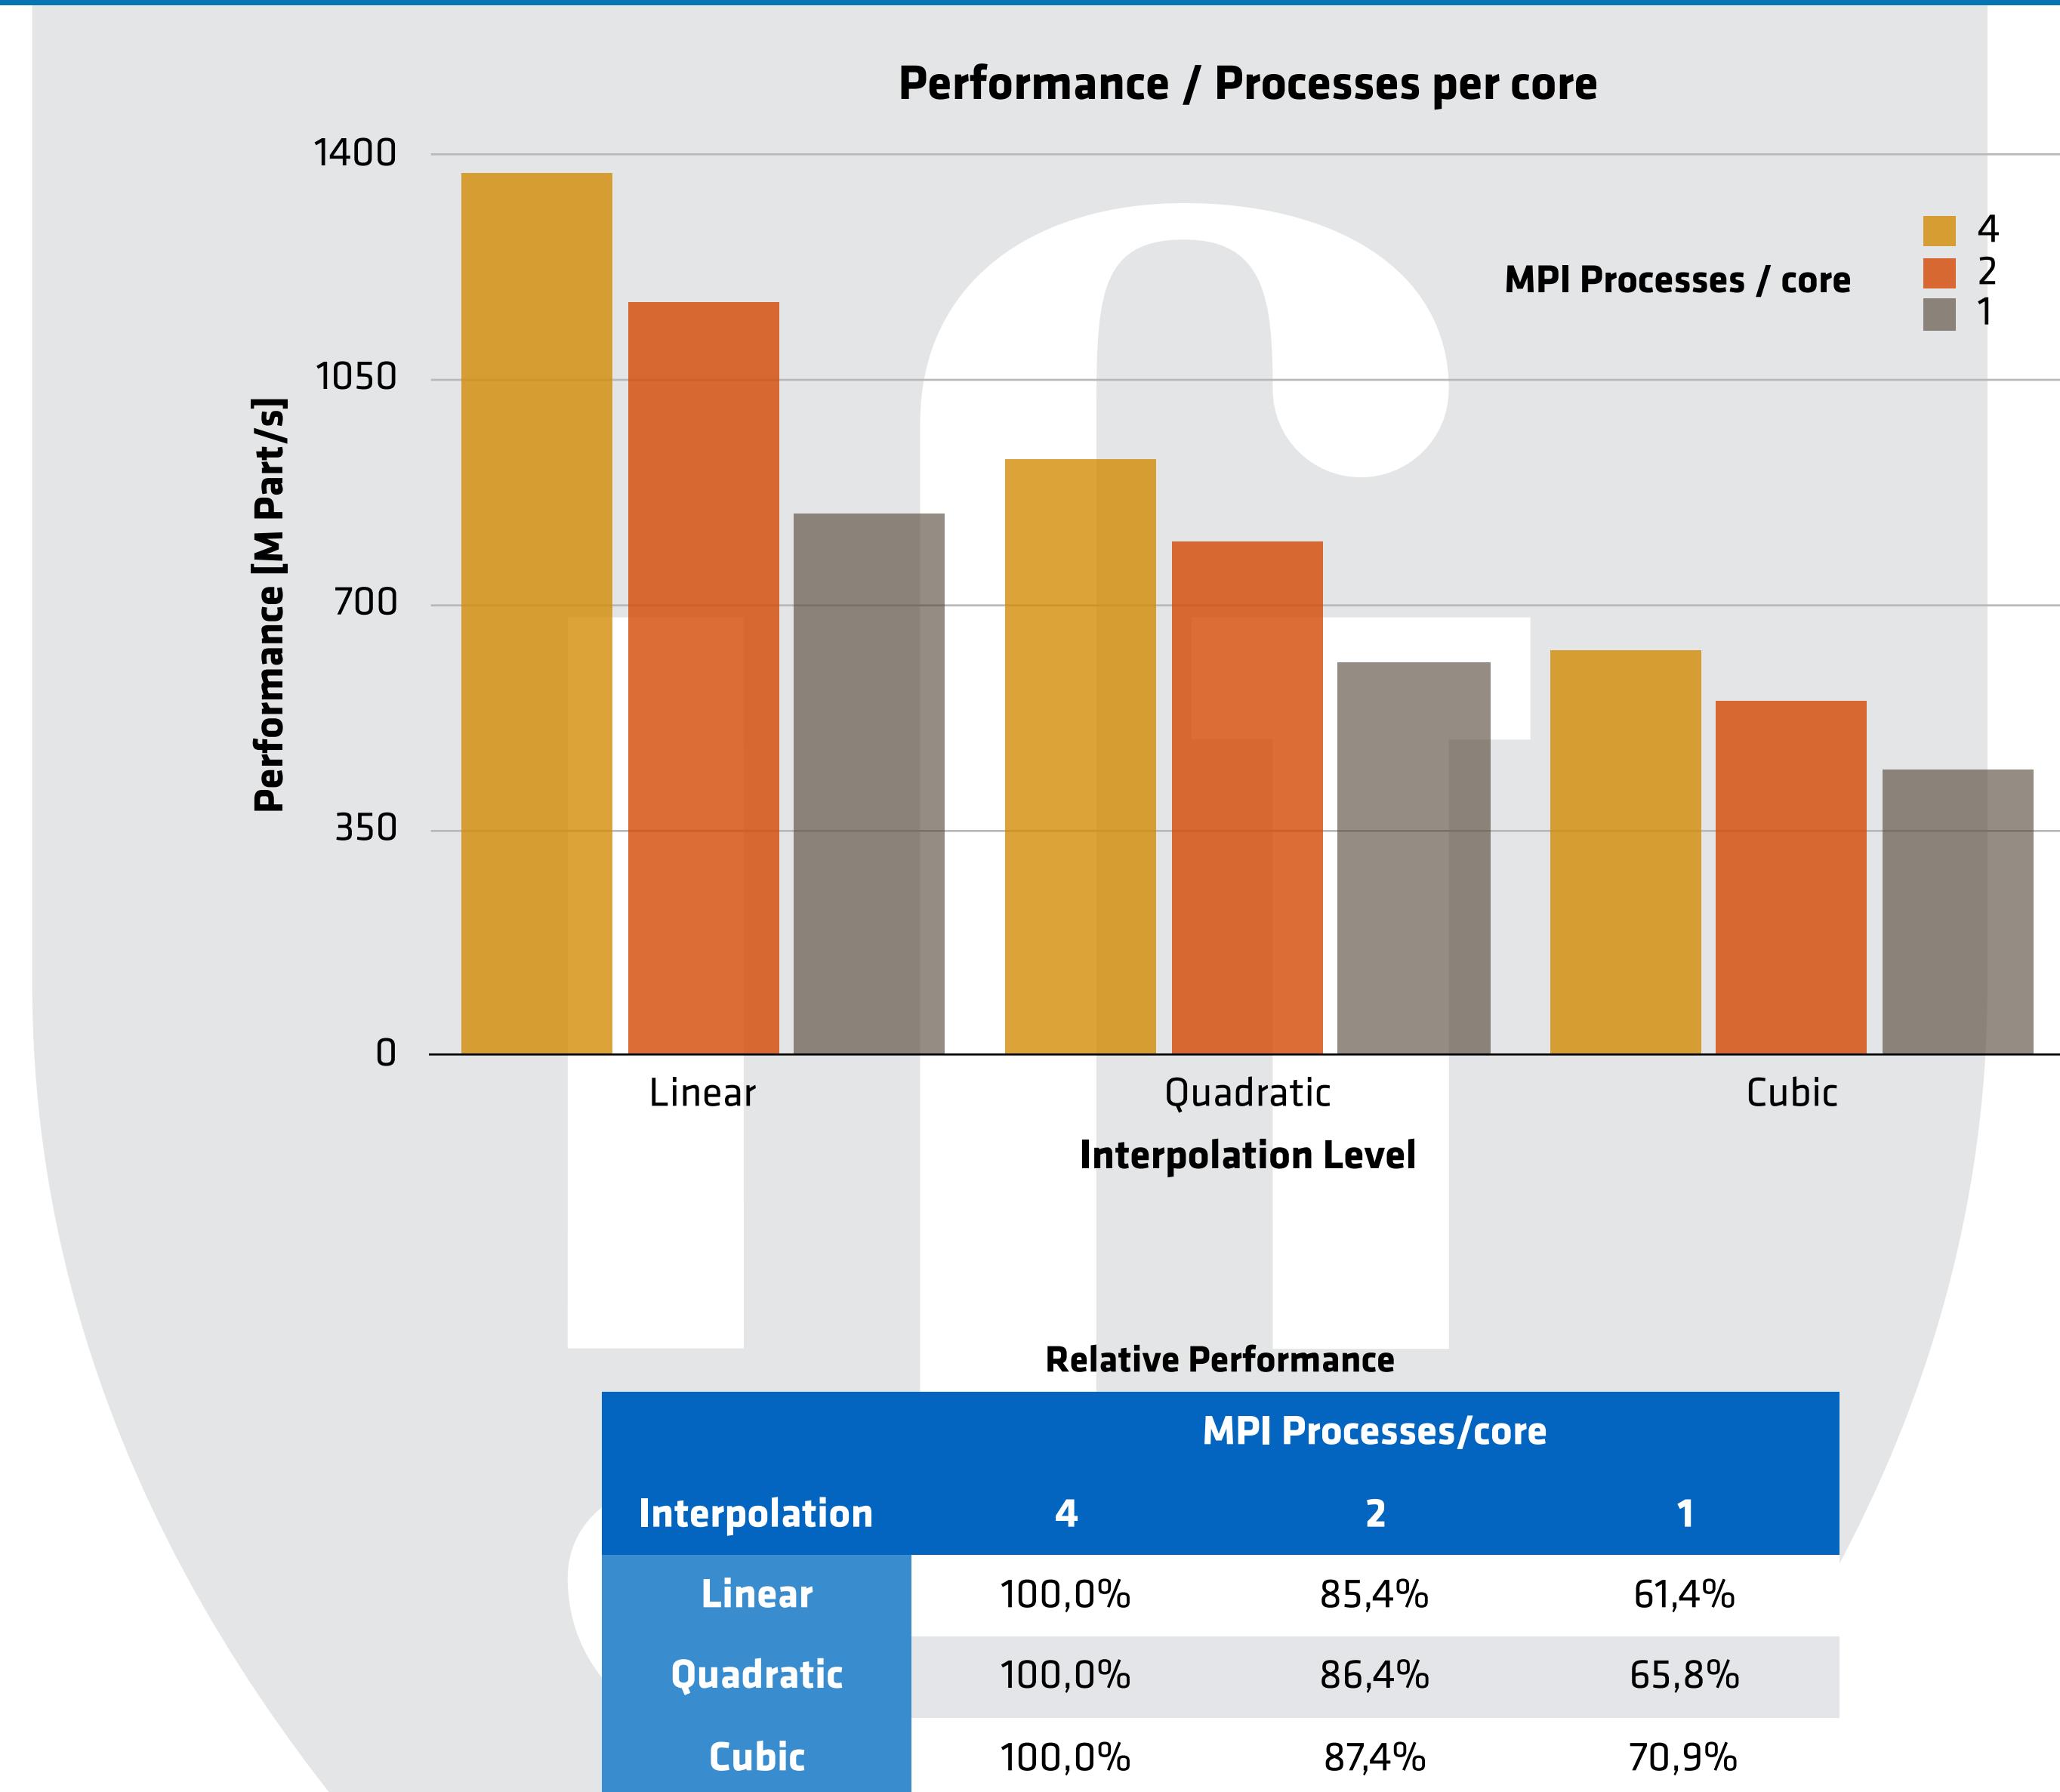
- Hide memory latency
- Saturate VPUs

## Test configuration

- Everything the same as reference configuration except
- 68 - 272 MPI processes
  - 4 processes/core

## Impact on Performance

- 38,6% drop on performance for linear interpolation from 4 to 1 processes per core
- Slightly smaller drop for cubic interpolation because of increased computational intensity
- Use of 4 threads per core recommended



# Shared memory (OpenMP) parallelism on the KNL

## OSIRIS supports hybrid MPI/OpenMP parallelism

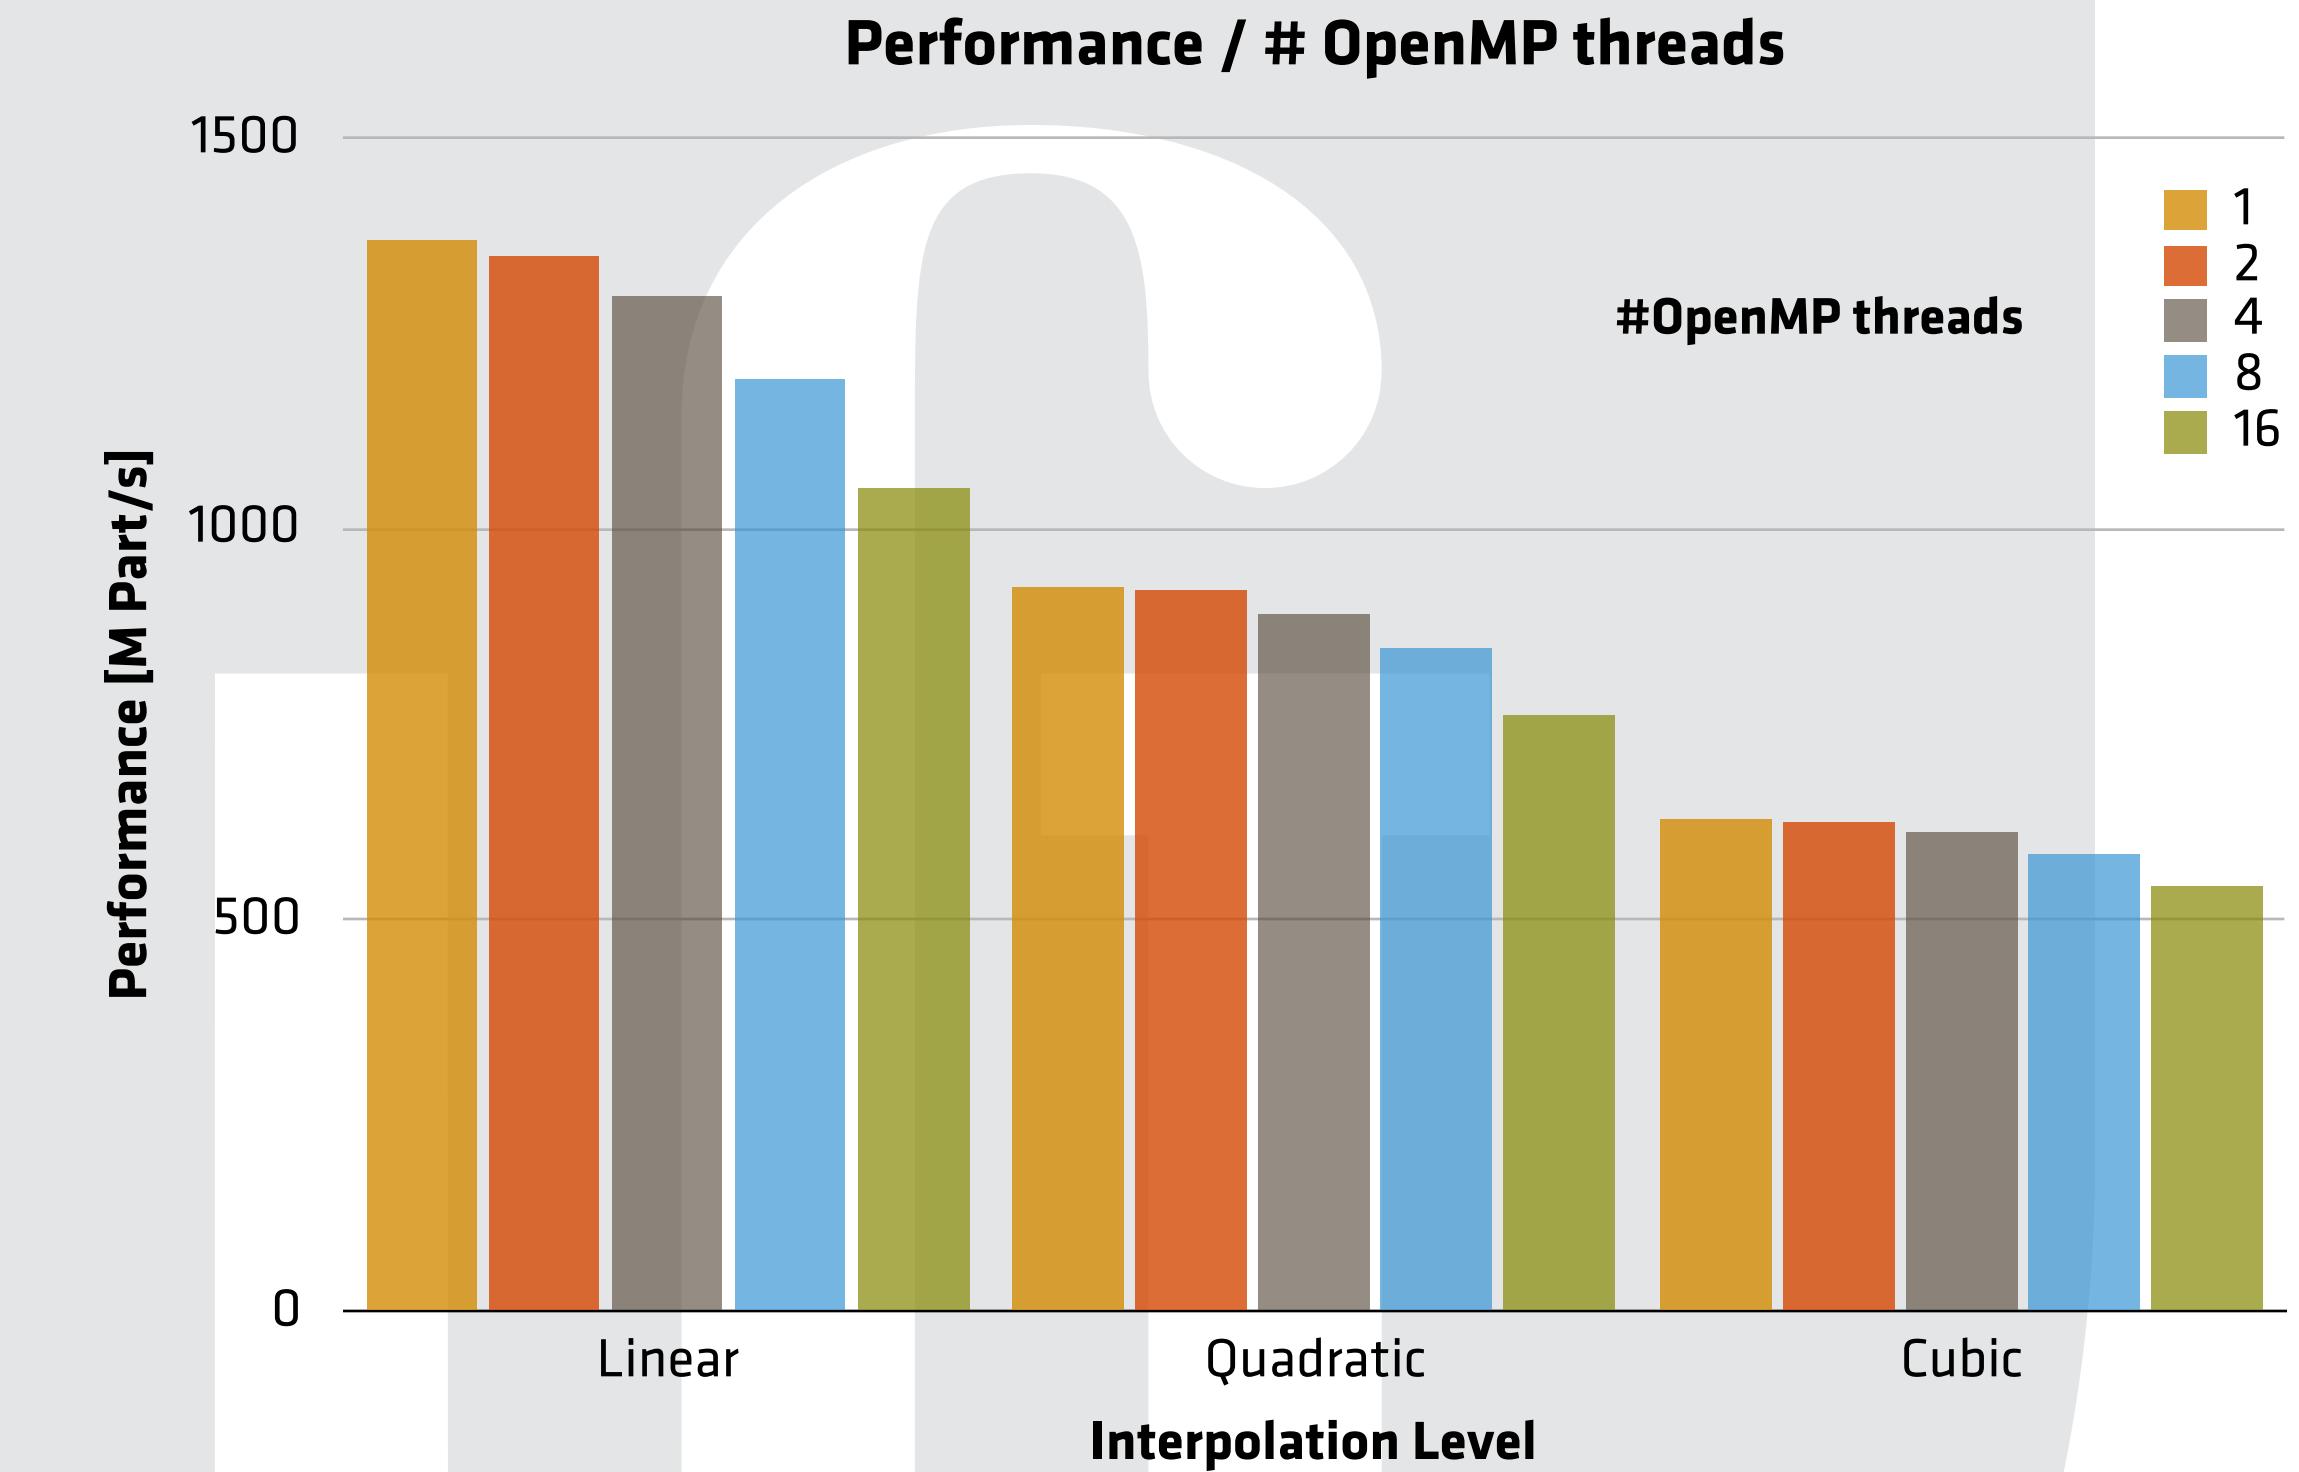
- Improve load imbalance
- Allow scalability for higher core counts
  - Minimal size per core for MPI parallelism
- OpenMP algorithm not as efficient as MPI for uniform plasmas

## Test configuration

- Everything the same as reference configuration except...
- Parallel partition  $n_p$  MPI processes  $\times$   $n_t$  OpenMP threads
- $n_t = 1, 2, 4, 8$  and  $16$  threads
- $n_p \times n_t = 272$  for all tests
  - 4 threads per core always

## Impact on Performance

- As expected, performance drops with increasing number of threads
- Negligible drop in performance ( $<\sim 5\%$ ) for up to 4 threads
- Even at 16 threads performance drop is acceptable, especially for higher interpolation levels



# Threads	Interpolation Level		
	Linear	Quadratic	Cubic
1	100,0%	100,0%	100,0%
2	98,4%	99,5%	99,4%
4	94,7%	96,5%	97,3%
8	87,0%	91,4%	93,2%
16	76,7%	82,2%	86,6%

# Explicit (intrinsics) vs. Automatic (compiler) Vectorization

## KNL cores have 2x 512bit vector units

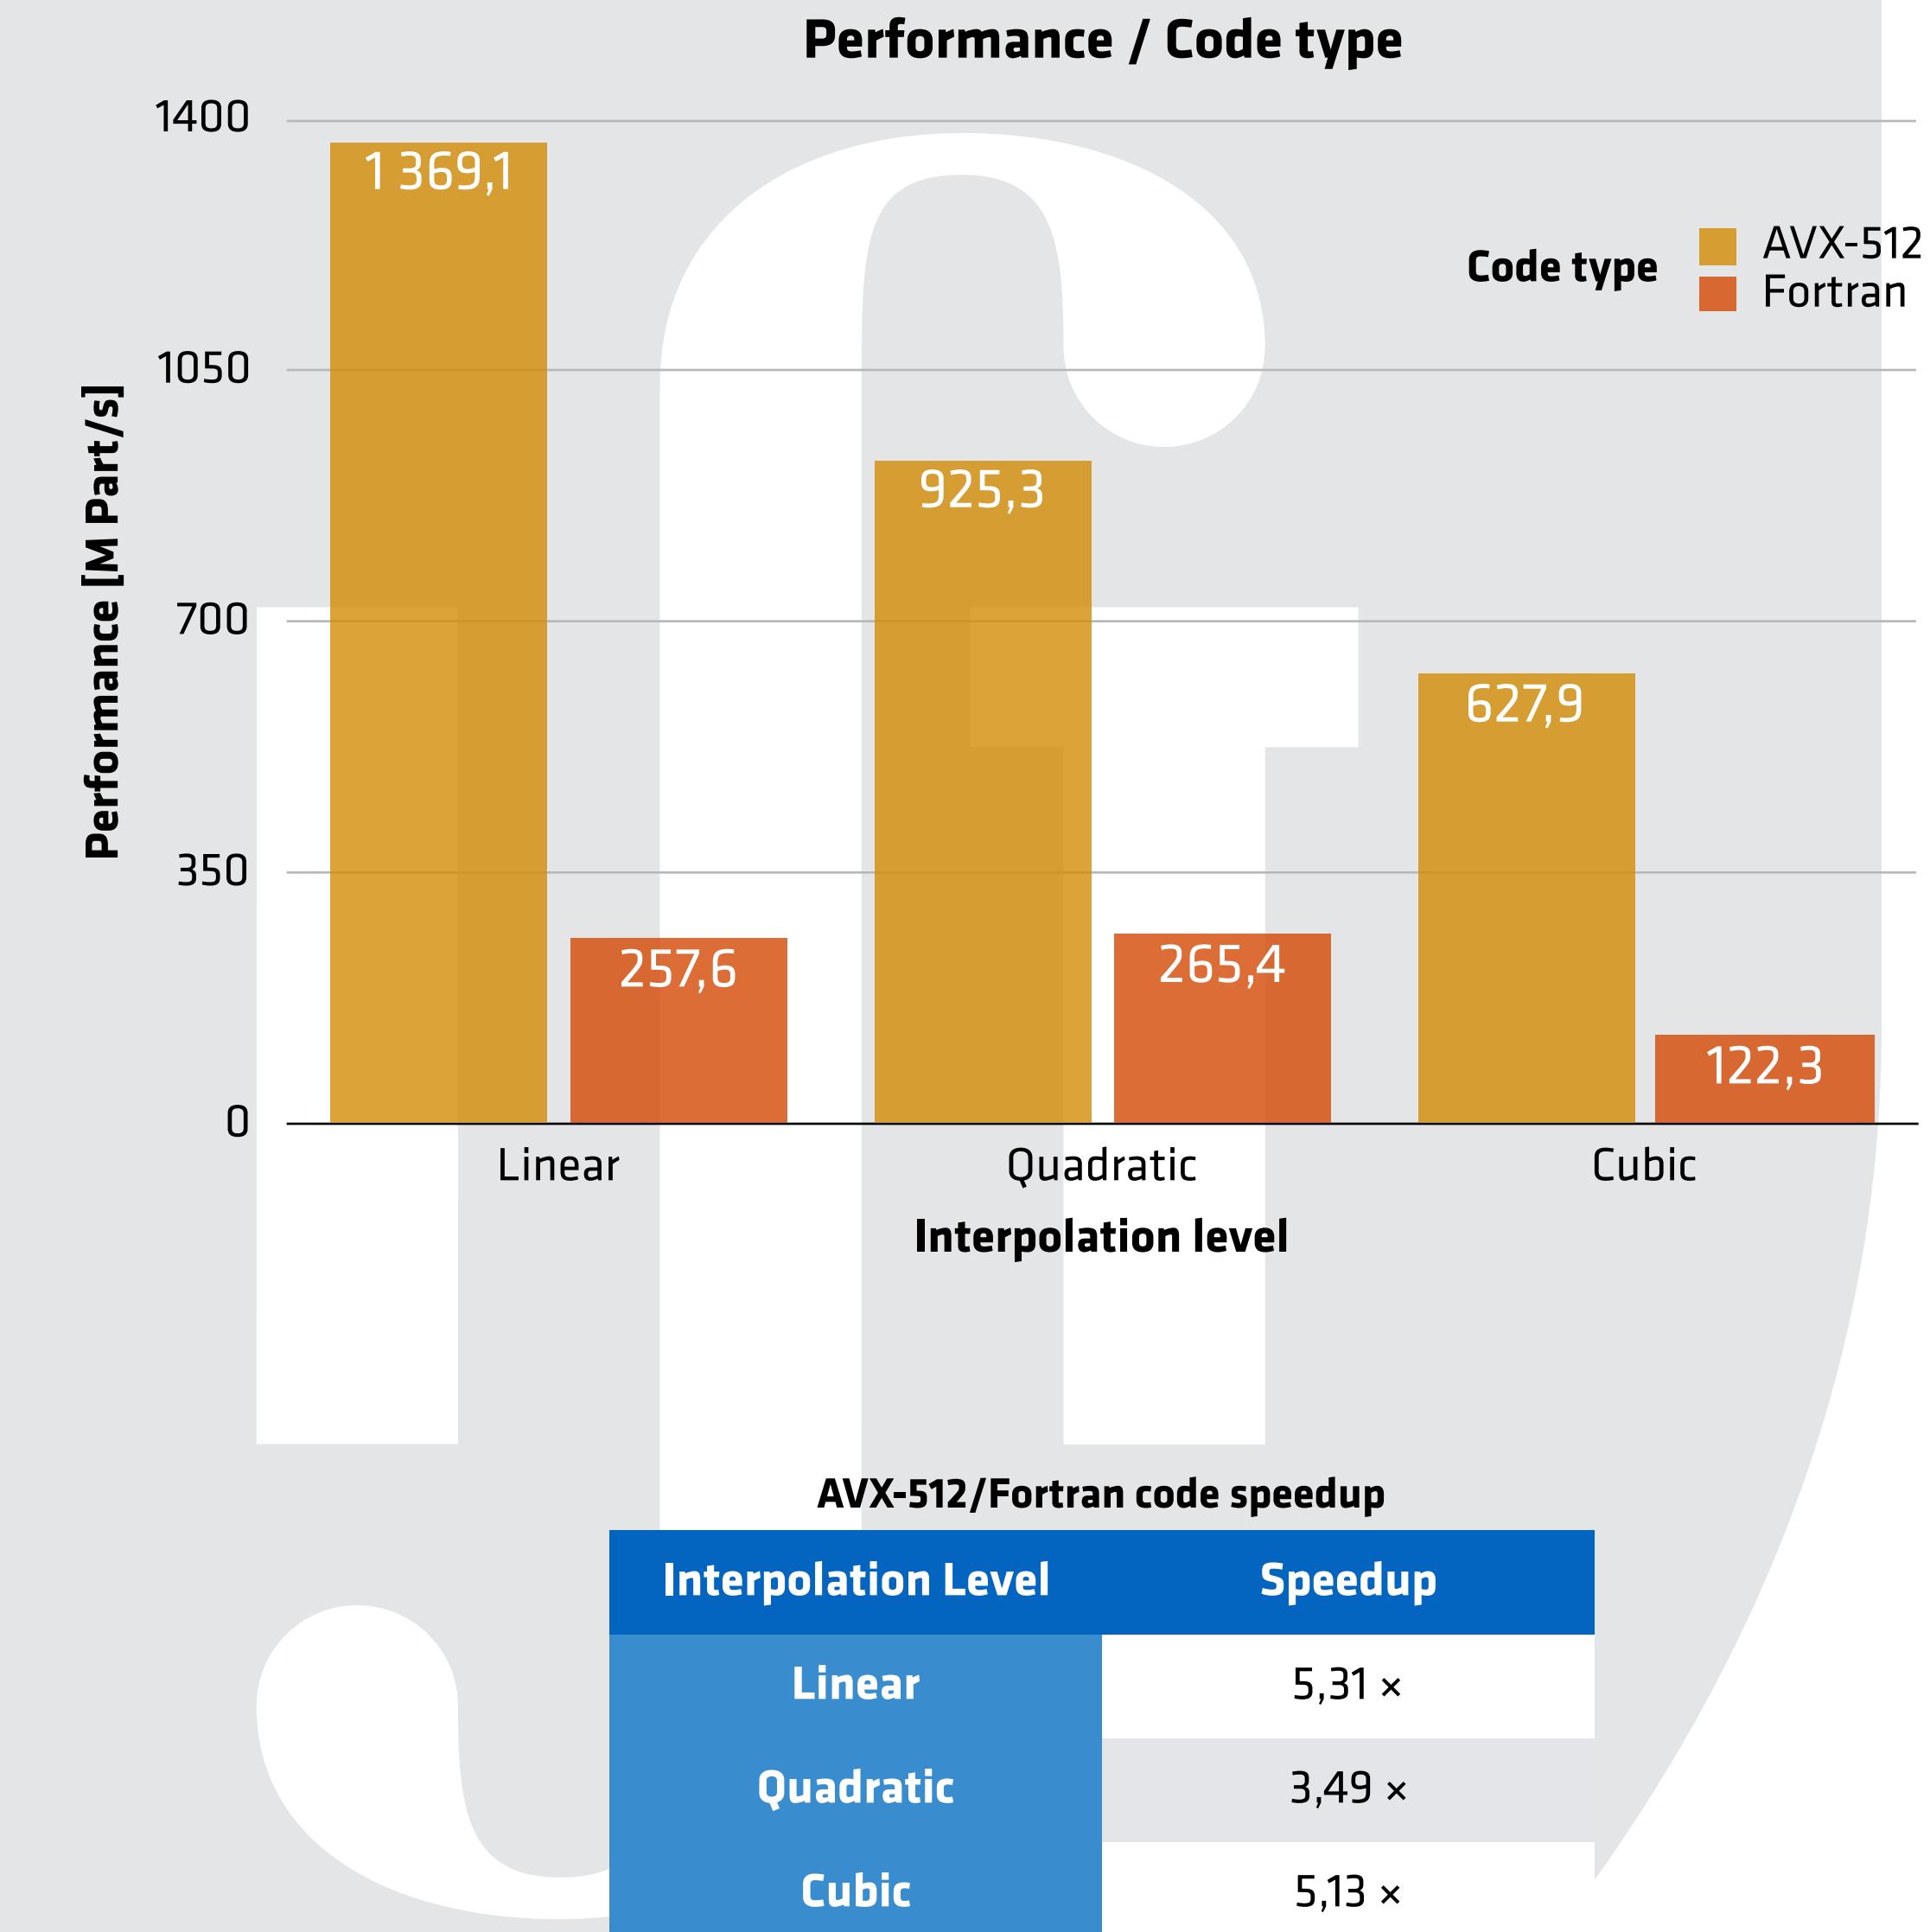
- Critical for performance
- Compiler can automatically vectorize the code to use these units
- Can be programmed explicitly using AVX-512 intrinsics

## Test configuration

- Everything the same as reference configuration except...
- Particle push is now done using the reference fortran pusher using automatic vectorization

## Impact on Performance

- Explicit vectorization using AVX-512 intrinsics achieves speedups from 3.5× to 5.3×
- The reference fortran code has not been modified (i.e. using pragmas to guide vectorization), may have room for improvement



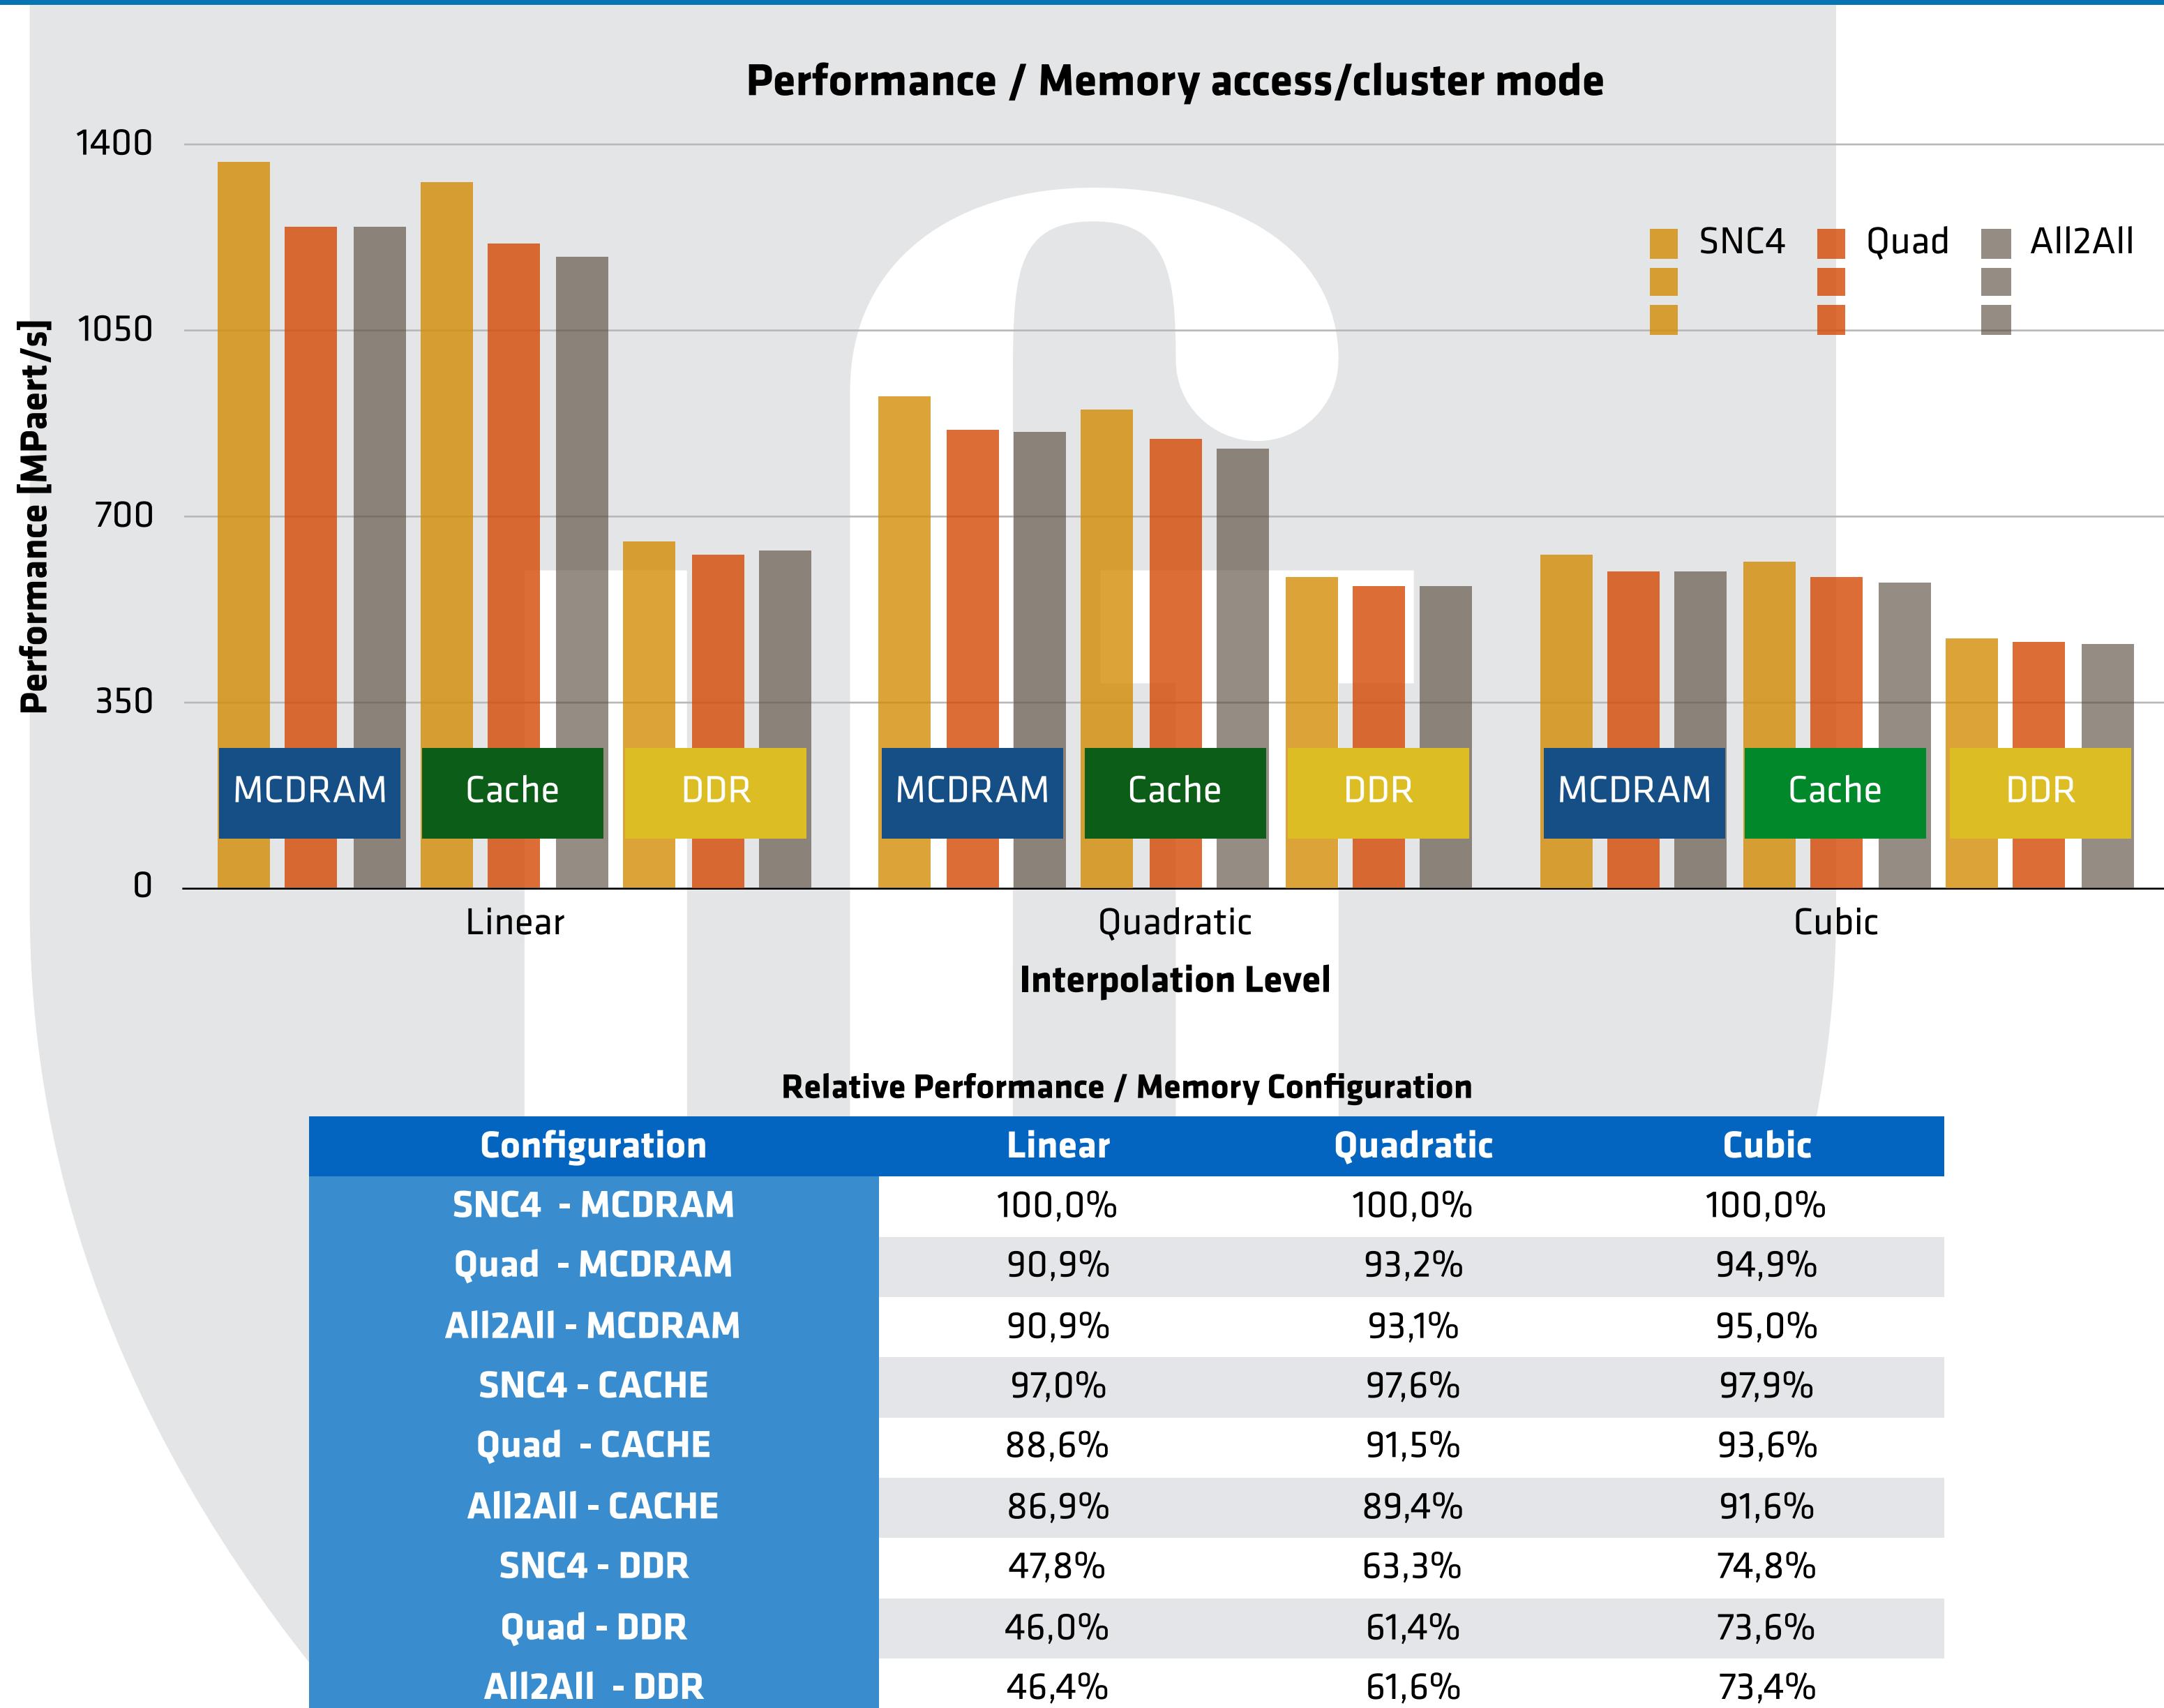
# Memory access/clustering mode

## KNL MCDRAM memory supports many configurations

- Access mode can be flat / cache
- Memory interconnect can be set to SNC4, Quadrant and All2All

## Impact on Performance

- Running the code completely in the MCDRAM yields best results
- Cache mode is extremely efficient, performance drop is below 2%
- Running the code strictly on DDR memory gives a 2.17 $\times$  slowdown in linear interpolation
  - On higher interpolation levels impact much lower, for cubic only a 1.36 $\times$  slowdown
- SNC4 faster than Quadrant and All2All, difference between Quadrant and All2All negligible
  - Again impact on lower interpolation levels more significant



# Knights Landing vs. Knights Corner

## KNL is the 2nd generation Intel MIC architecture

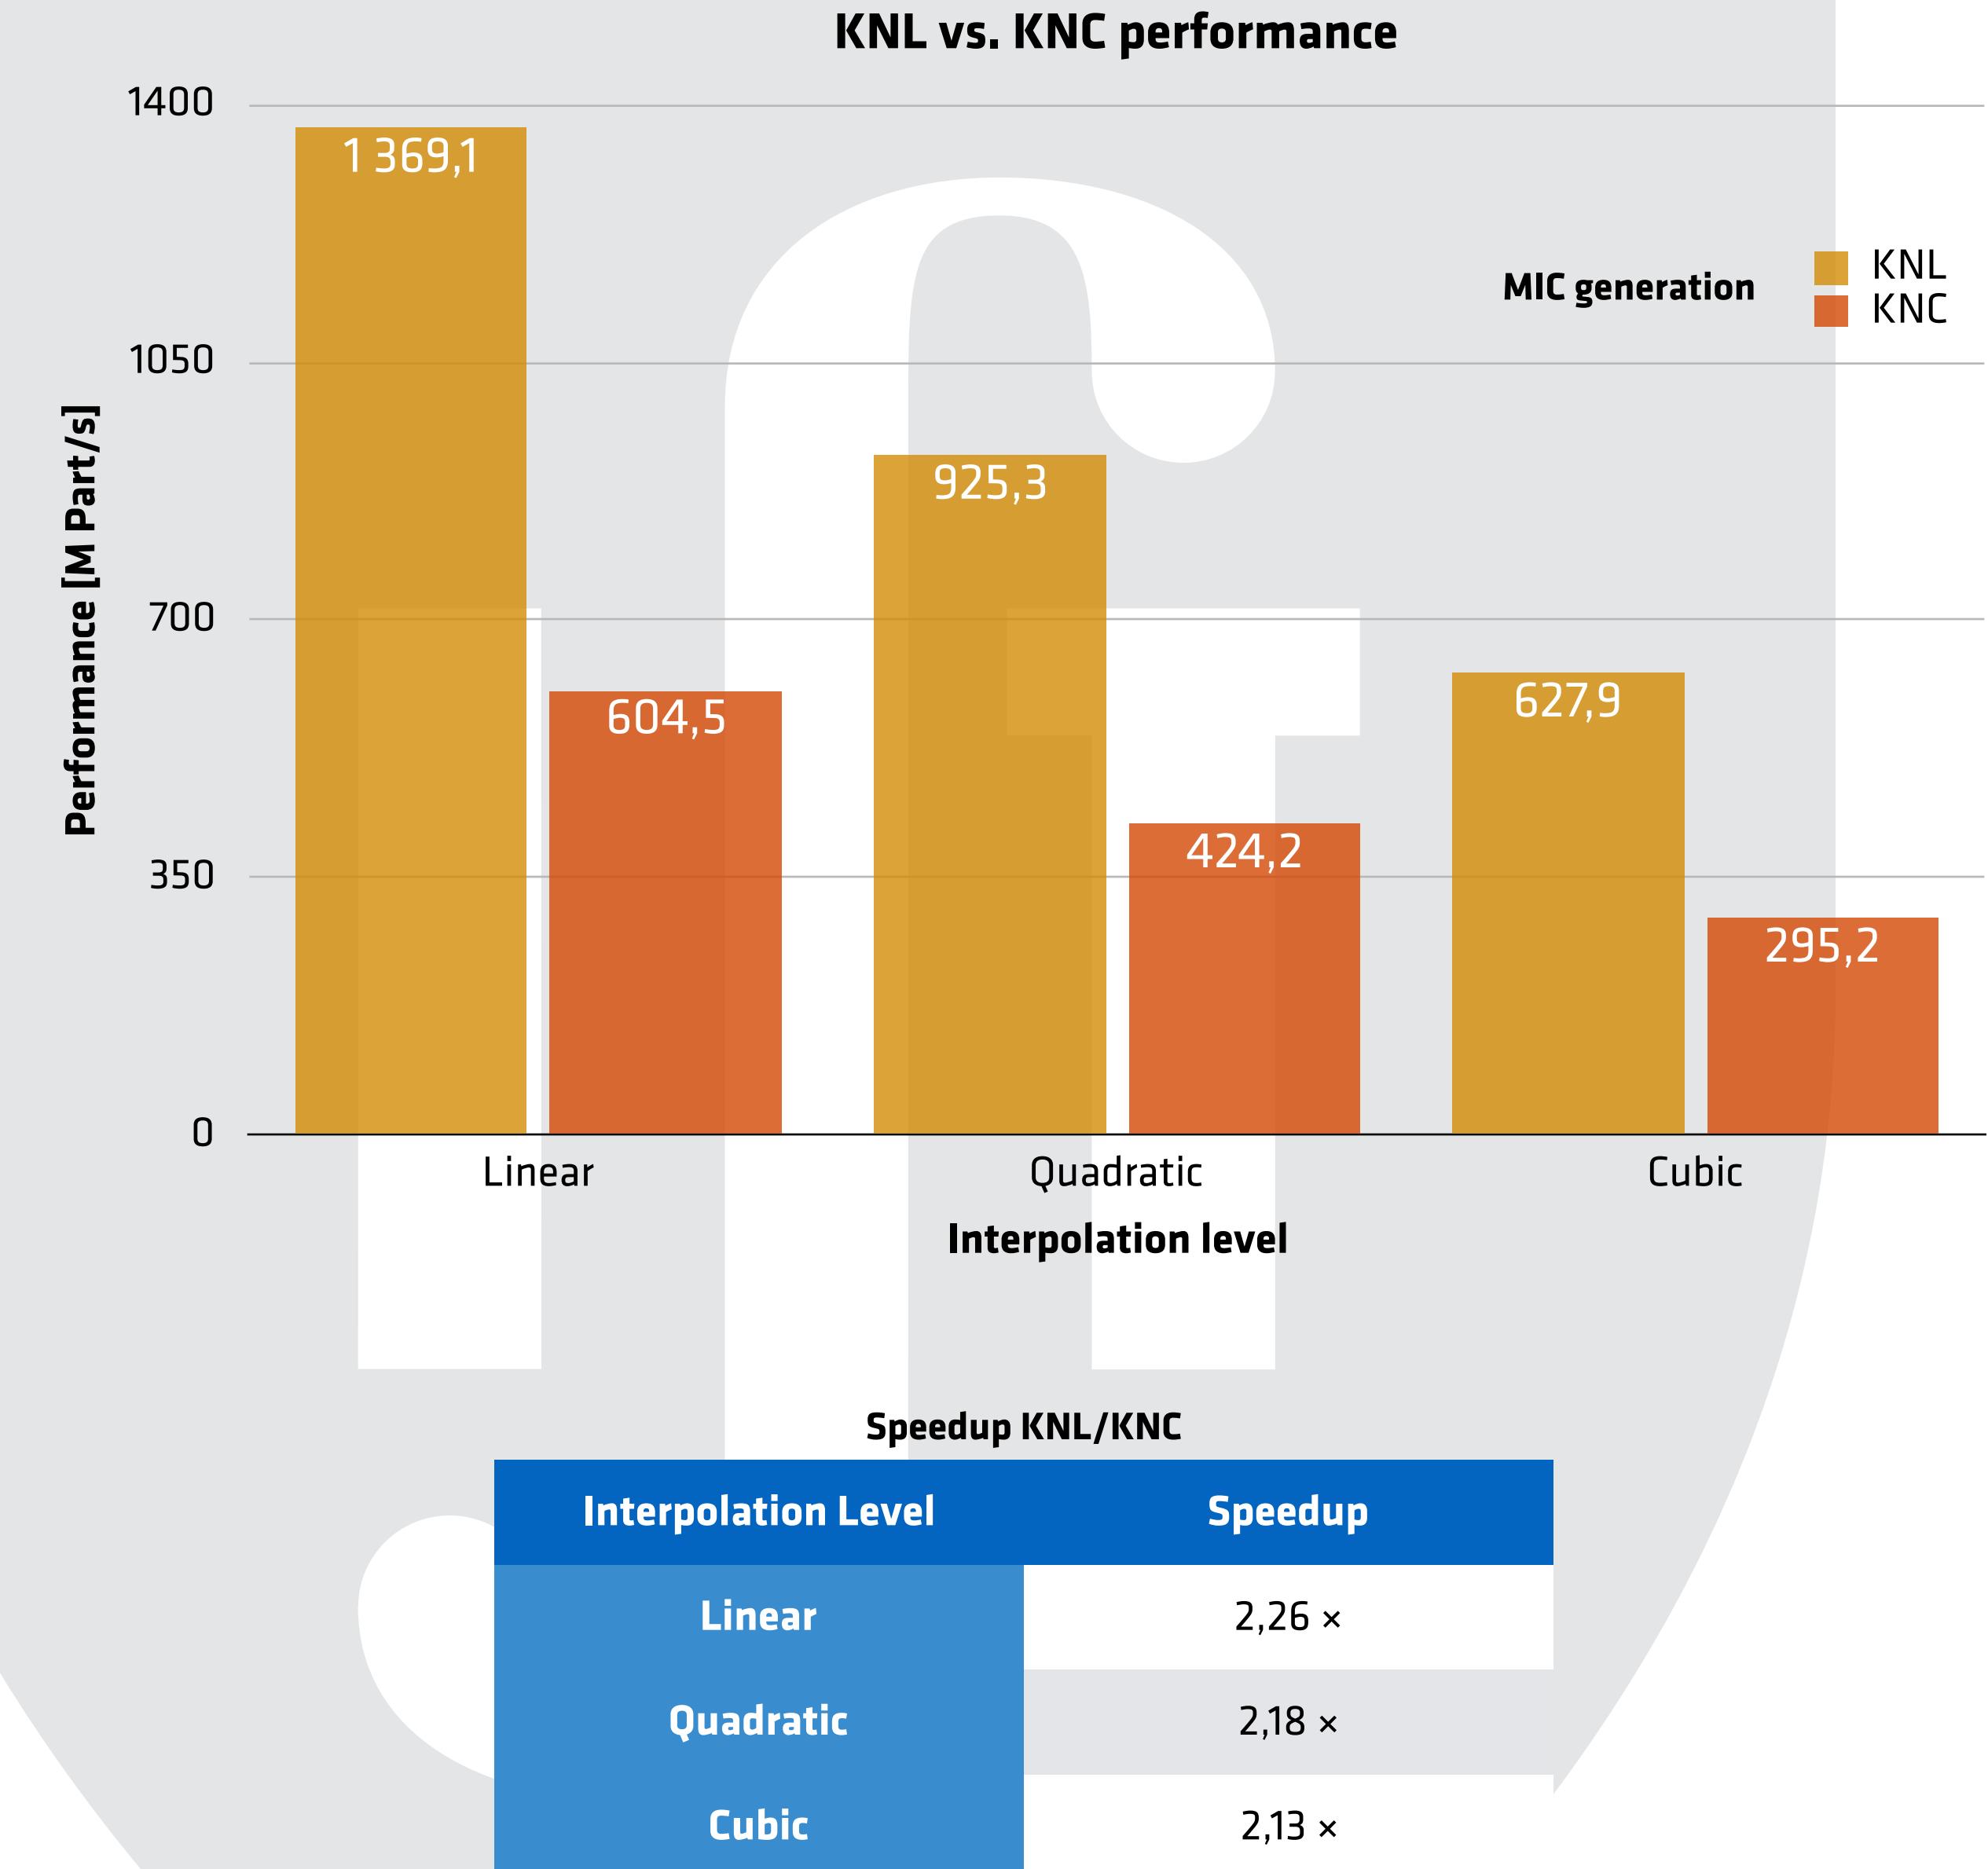
- First generation was the Knights Corner architecture
- Available only as coprocessor boards
- KNC board configuration
  - 60 x86\_64 cores @ 1.053 GHz
  - 1x 512bit VPU/core
  - 8 GB GDDR5 RAM
  - Peak FP ~ 2 TFlop/s (SP)
- KNL main differences
  - More cores / higher clock speed
  - Twice the VPU units / core
  - 16 GB MCDRAM
  - Peak FP ~ 6 TFlop/s (SP)

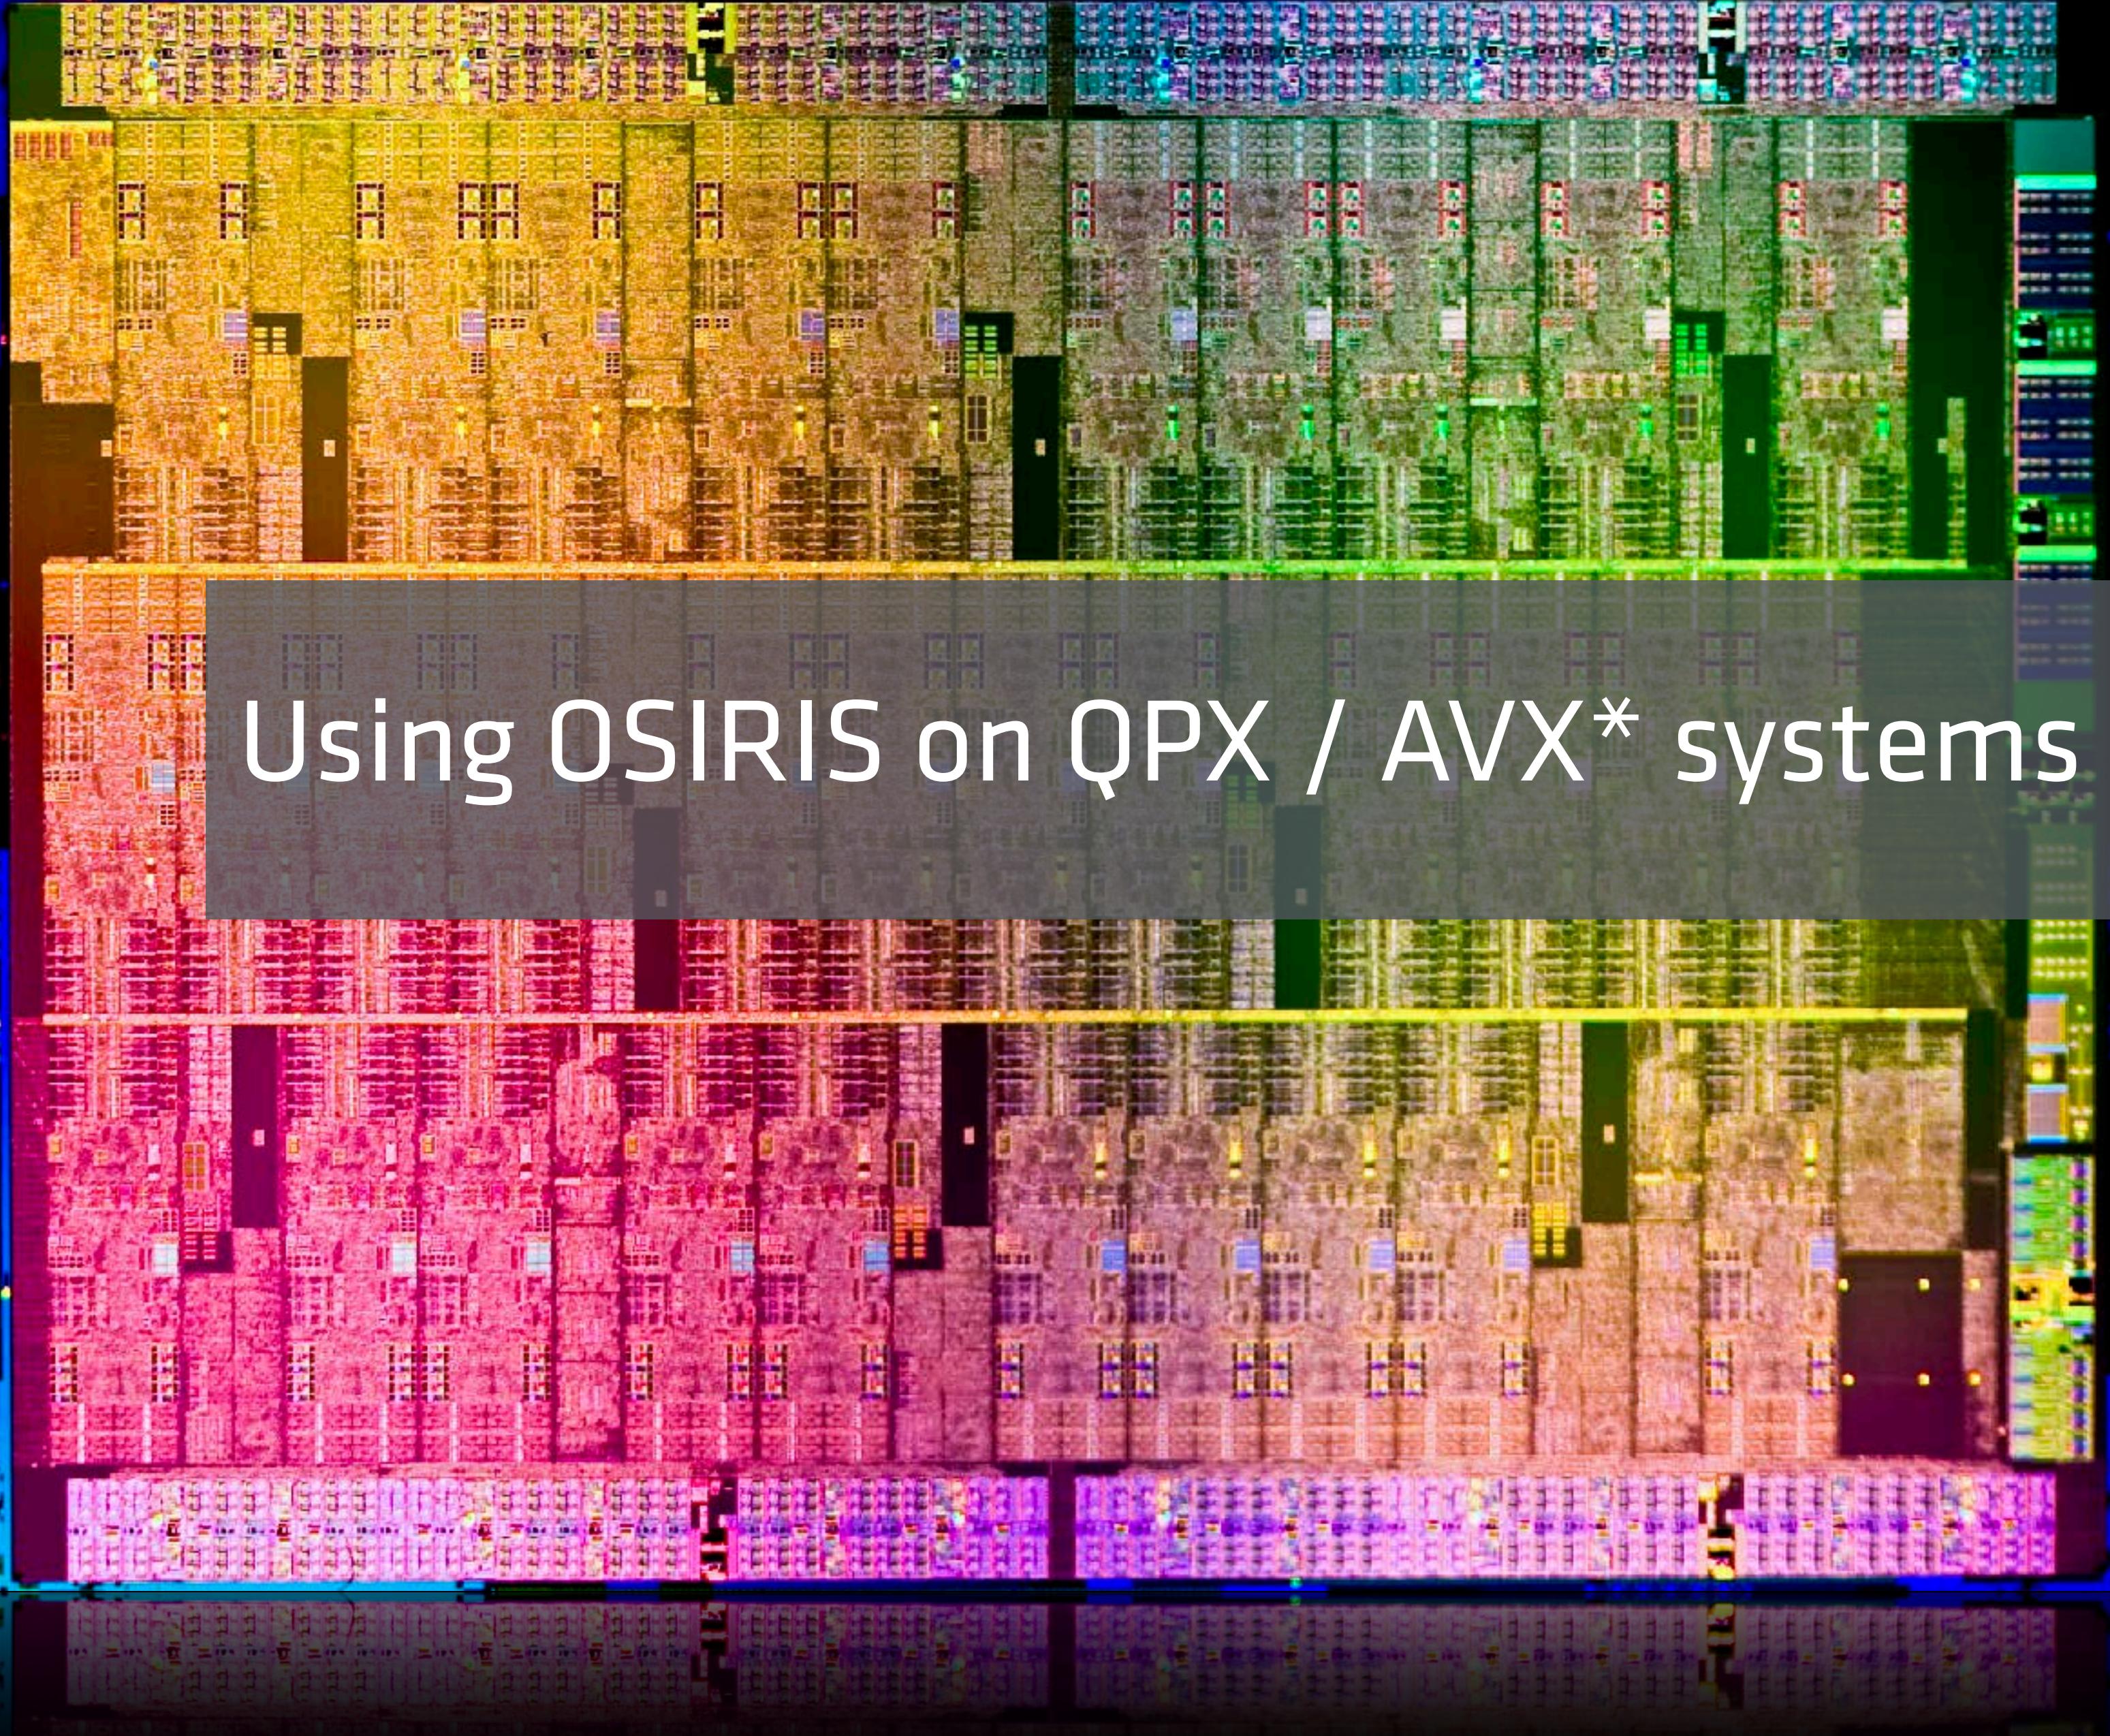
## Programming for KNL vs. KNC

- KNC intrinsics almost identical to AVX512 with a few exceptions, small changes required to vector code
- KNL has additional instructions for unaligned memory access
- Also additional AVX512 instructions (e.g. conflict detection), not explored yet.

## KNL vs. KNC performance

- Avg. speedup was 2.2×
- Floating point efficiency lower on KNL
- Room for improvement on KNL code





Using OSIRIS on QPX / AVX\* systems



Intel Xeon Phi 7290 die

# SIMD support on modern CPUS

	SIMD extension support					
Name	SSE2	AVX	AVX2	AVX512 <sup>1</sup>	KNC <sup>2</sup>	QPX
Architecture	x86	x86	x86	x86	x86	PowerPC A2
Width [bits]	128	256	256	512	512	256
Fused multiply-add operations	No	No	Yes	Yes	Yes	Yes
Vector Gather / Scatter instructions	No	No	Yes	Yes	Yes	No
Float vector width	4	8	8	16	16	n/a
Double vector width	2	4	4	8	8	4
Integer vector width	4	4+4	8	16	16	n/a
System	Jaguar	Supermuc	Trinity	Marenostrum 4	Tianhe-2	Sequoia

<sup>1</sup> AVX-512 extensions are available on both Intel Knights Landing and Xeon Platinum CPUs

<sup>2</sup> KNC systems are not binary compatible with other x86 systems

## Modern CPUS include SIMD vector unit(s)

- Vector registers of  $n_V$  values (int/float/double)
- Instructions act on vector registers
- Same operation on  $n_V$  values simultaneously
- Most also include support for fused multiply-add operations
  - $a \times b + c$  in a single op

## Programming SIMD units

- Let the compiler do it
  - Developer may give compiler hints to help automatic vectorization of code
- Use intrinsics
- Special operations that translate to vector operations directly

## OSIRIS includes support for all the main SIMD units

- Explicit vectorization of the particle advance/deposit routines
- Automatic vectorization of the field solver routines

# Using OSIRIS with vector (SIMD) hardware

## @ Compile time

- Make sure you have enabled the correct SIMD flag on the configure file

```
# SIMD
SIMD = AVX2
```

### • Available options are

- **QPX** - PowerPC A2 CPUs (BlueGene/Q systems)
- **AVX2** - x86 CPUs supporting AVX2 extensions, such as Intel Kaby Lake or AMD Ryzen
- **AVX512** - x86 CPUs supporting AVX512 extensions, such as Intel Knights Landing or Intel Xeon Platinum

### • Legacy options

- **SSE** - Really old x86 systems
- **AVX** - Not so old x86 systems that do not yet support AVX2 (e.g. my previous laptop)

### • SSE/AVX\* options are forward compatible

- e.g. AVX code will run on AVX512 cpus

## @ Run time

- OSIRIS will always compile the standard Fortran advance/ deposit also
- Users can choose between them using the **push\_type** parameter in the species section

```
species
{
    name = "positrons" ,
    num_par_x(1:2) = 4, 4,
    rqm=+1.0,
    ! push_type = "standard",
    push_type = "simd",
}
```

### • When SIMD support was chosen at the compile stage “simd” will be the default

- Users can choose the standard pusher using the “standard” option
- *Can also be used to choose the radiation cooling pusher (“radcool”)*

# Production Runs - Knights Landing

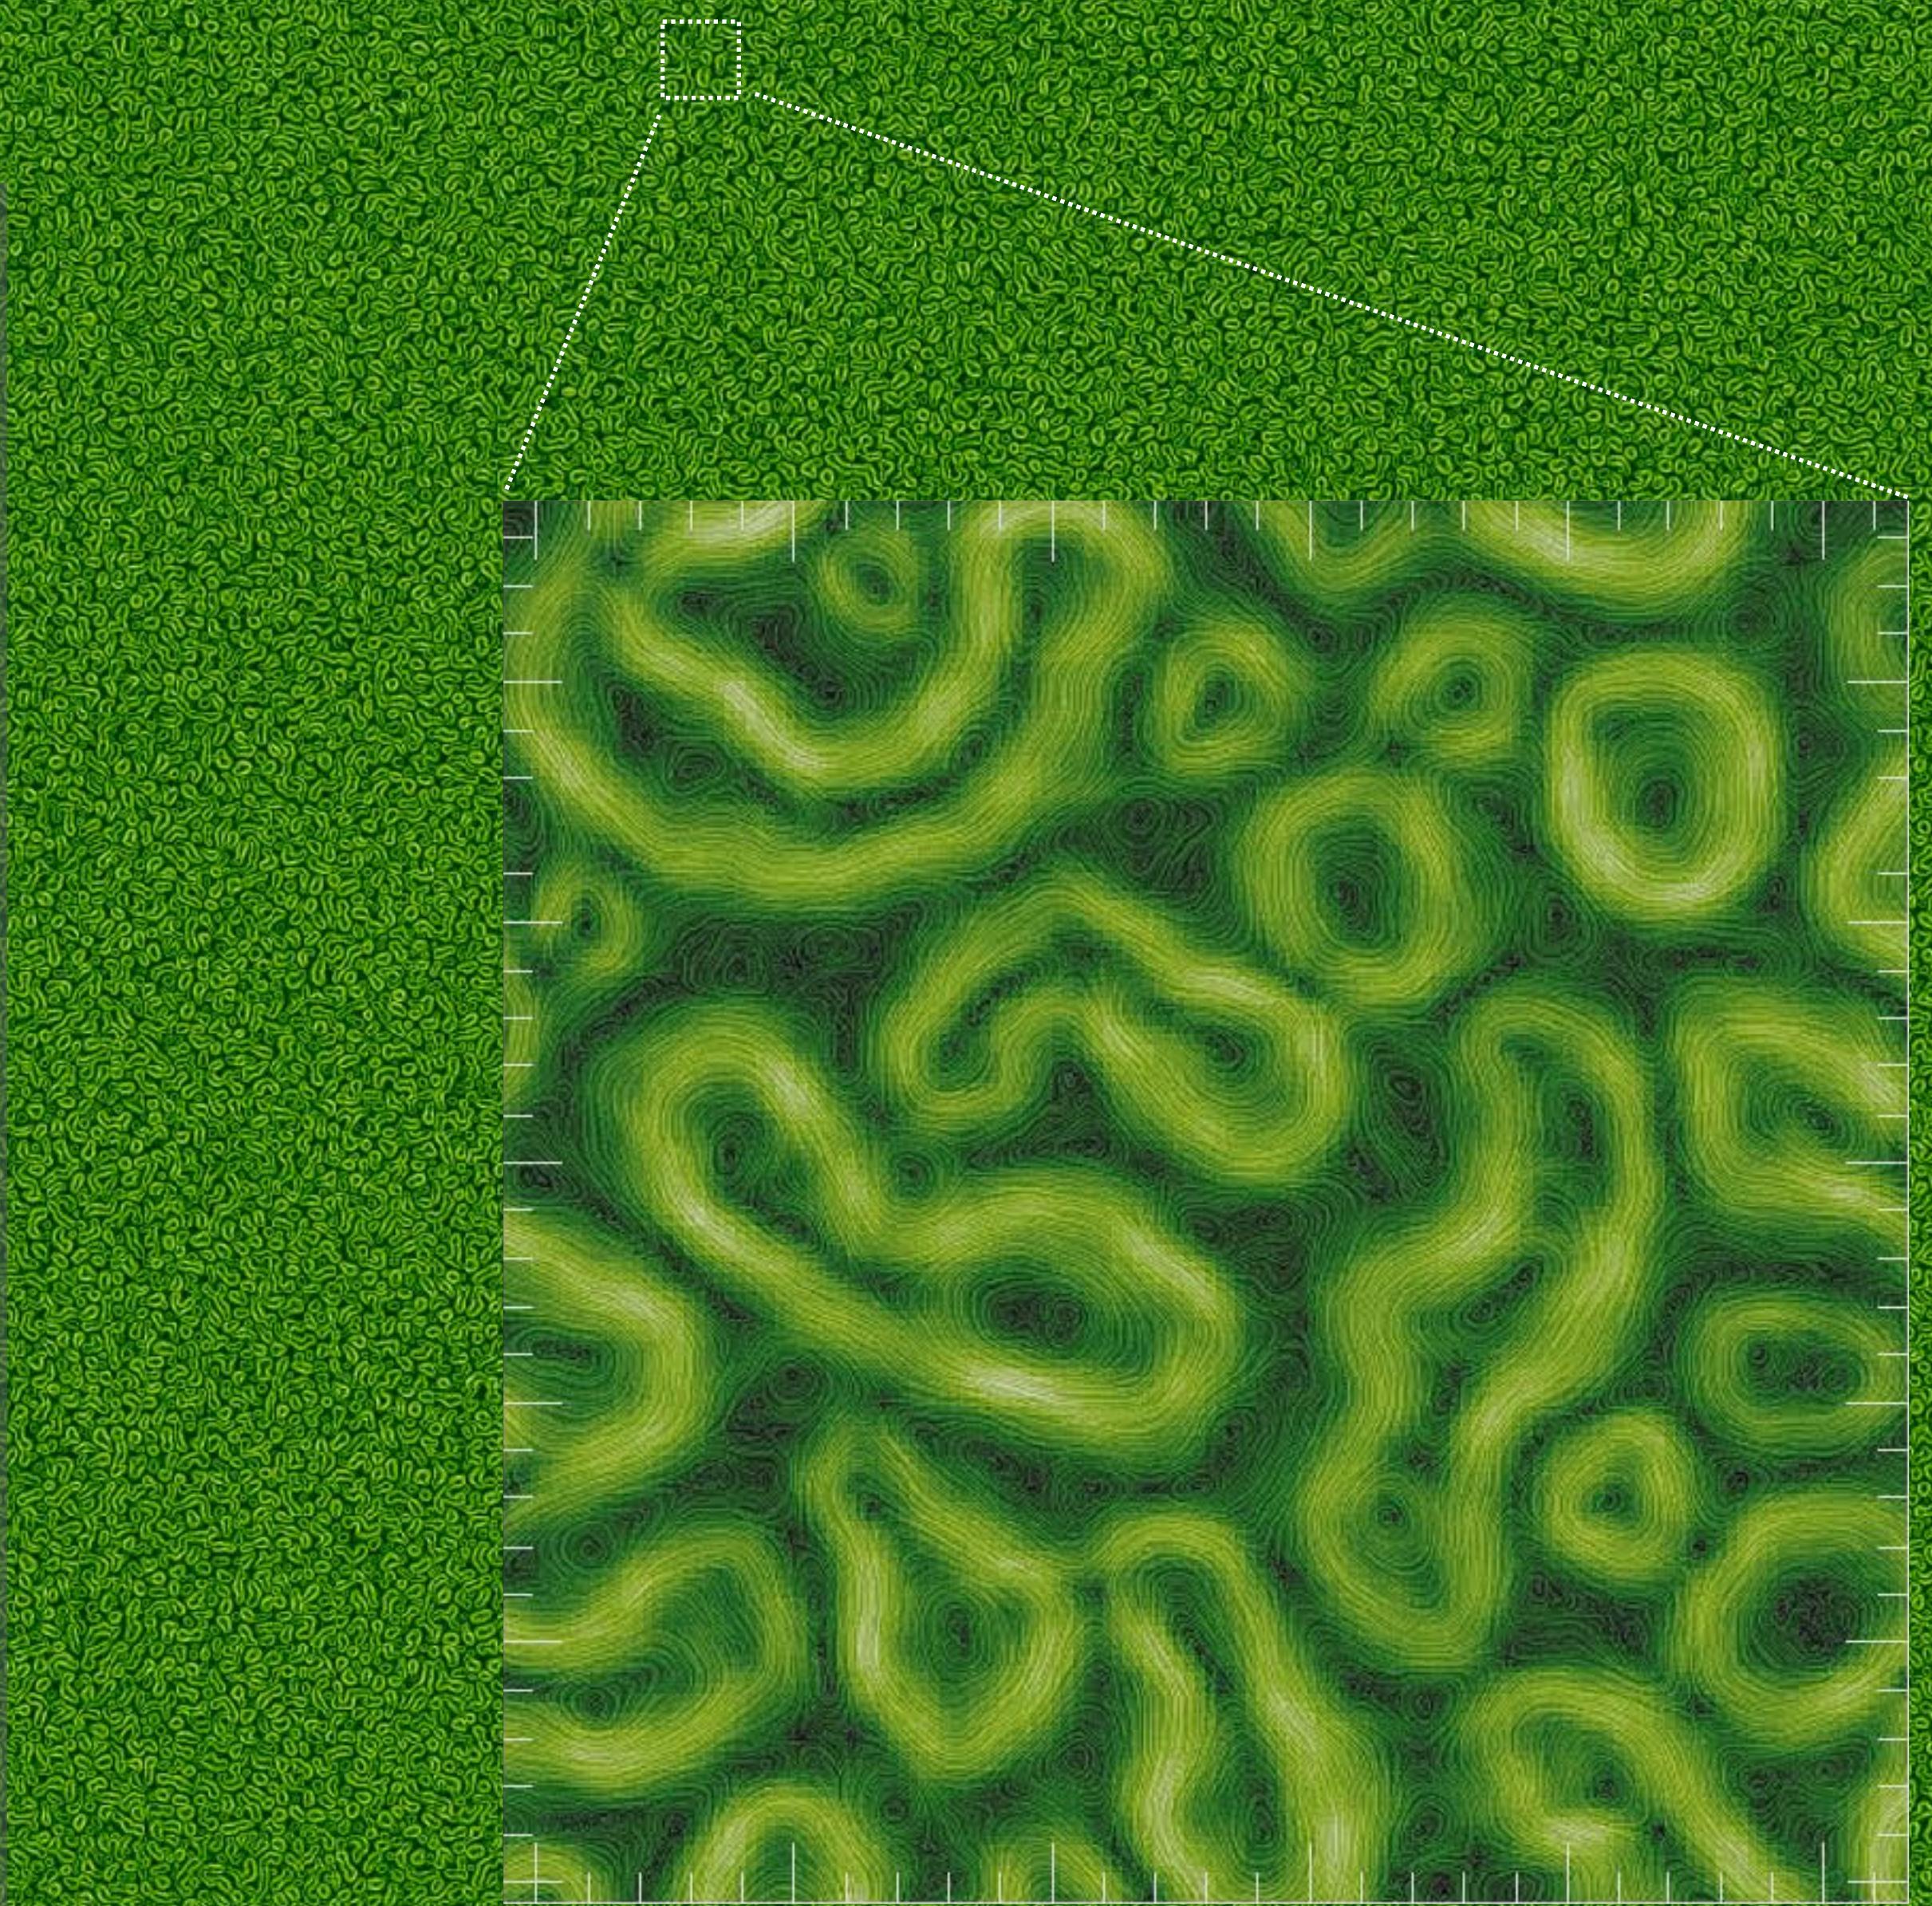
- **Simulation setup**

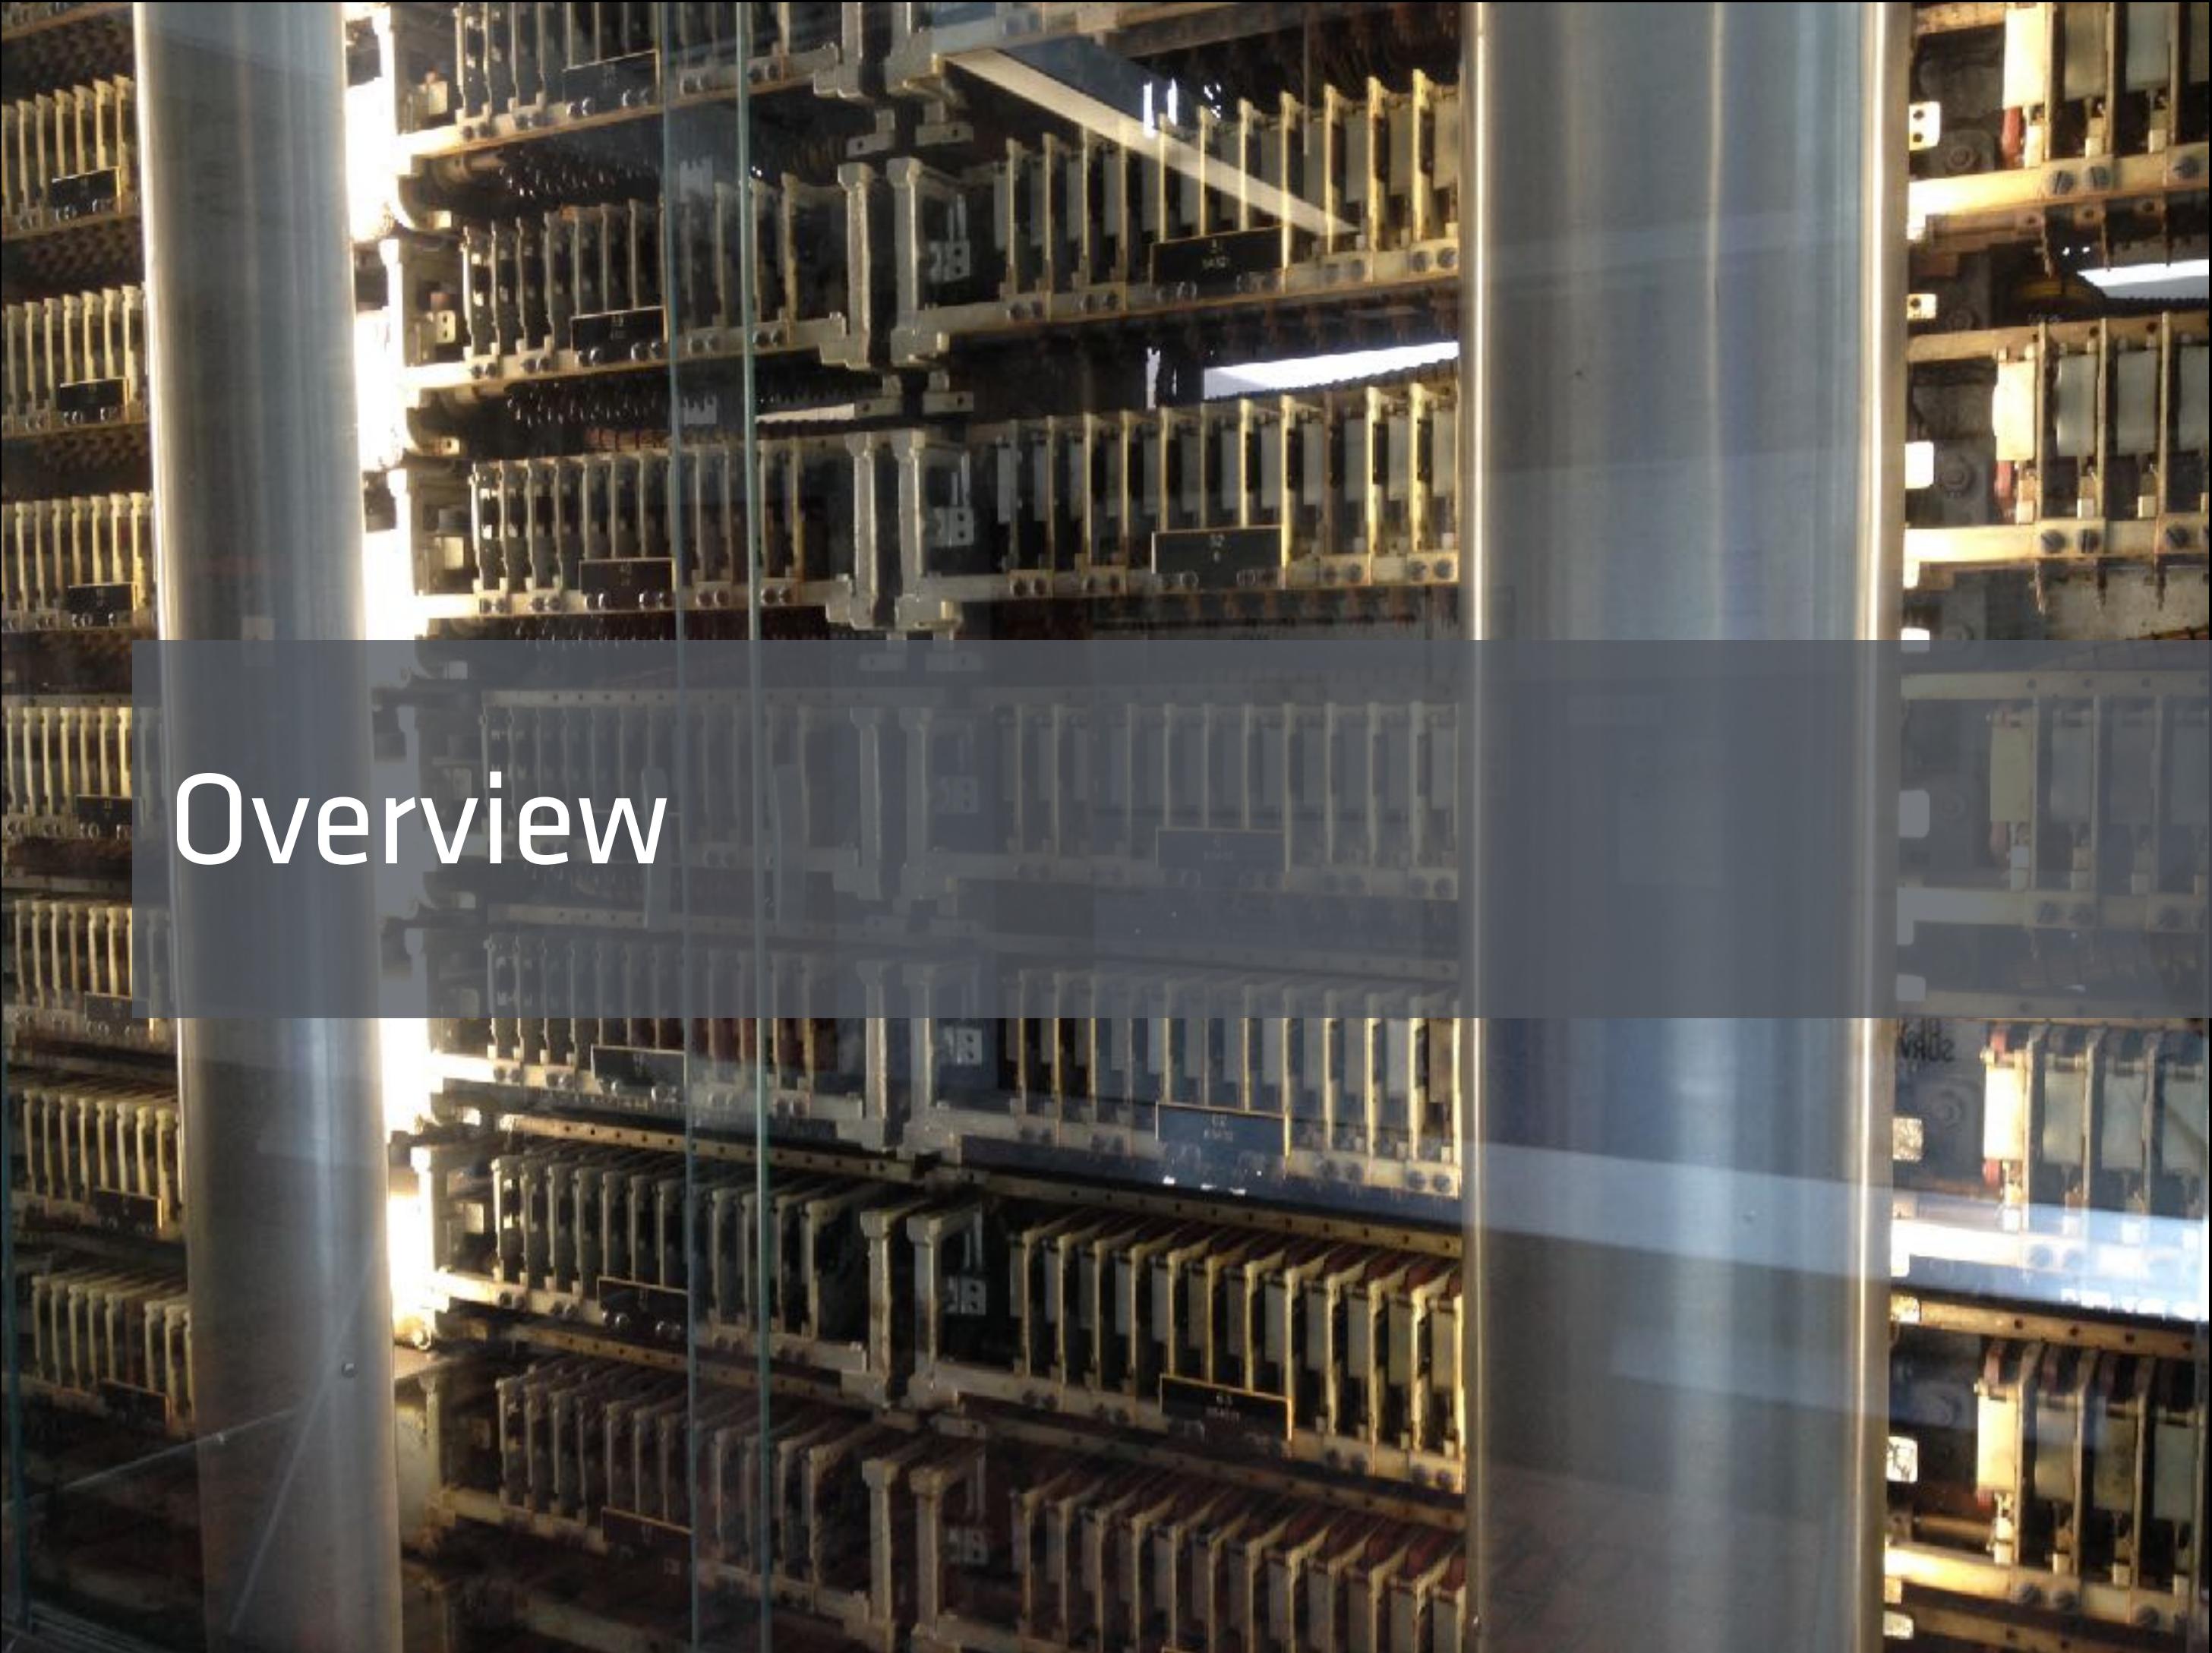
- Collision of two neutral electron/positron clouds
- Physics dominated by the Weibel instability
- 2D simulation in the perpendicular plane

- **Parameters**

- $8192 \times 8192$  cells
- 4 species,  $4.2 \times 10^9$  particles total
- $\gamma V_{fl} = \pm 0.6 c$
- $\gamma V_{th} = 0.1 c$
- SNC4 / Cache
- 50 GB ~ 50% system memory (DDR) used

- **Very good performance on a single KNL node**





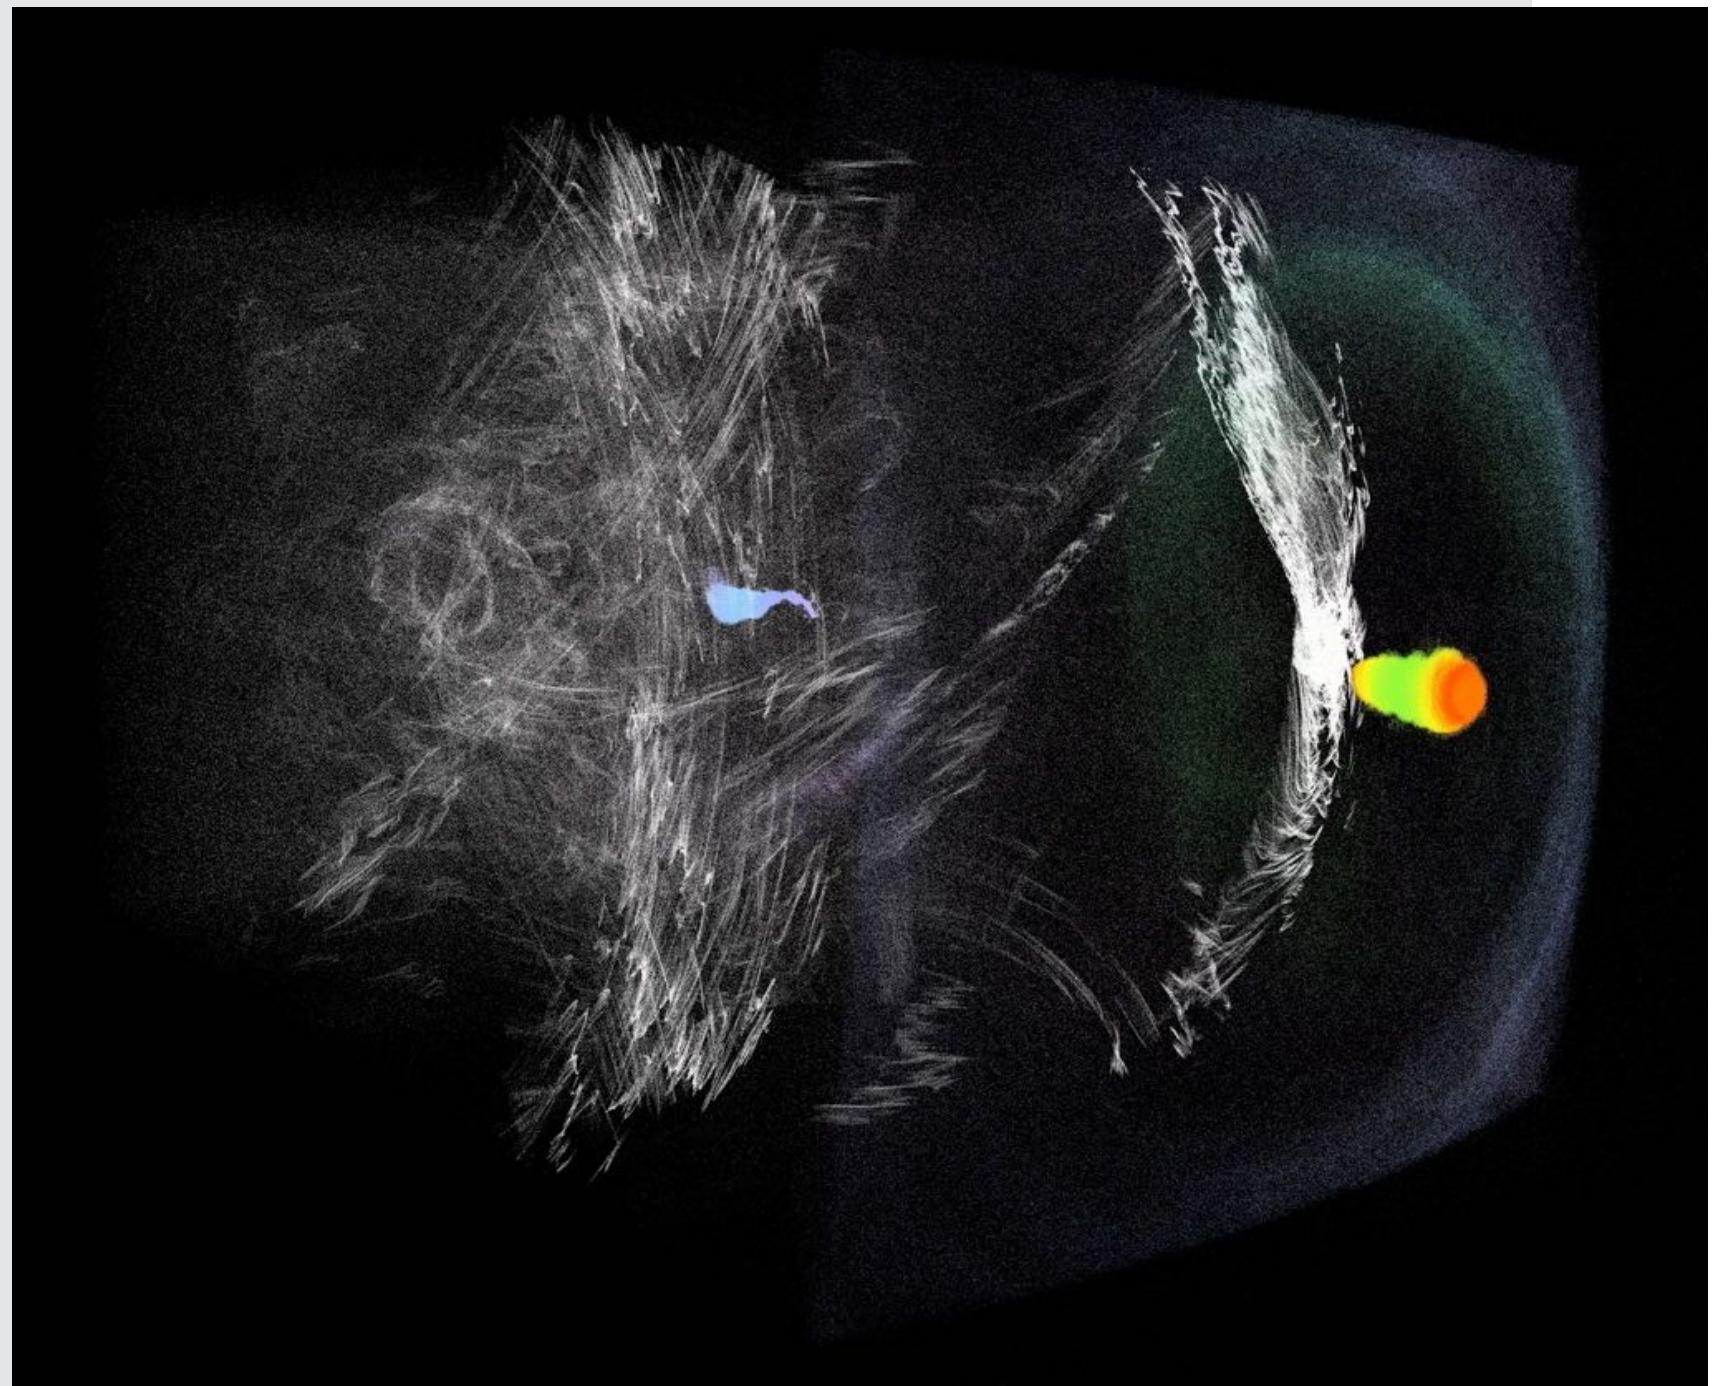
# Overview

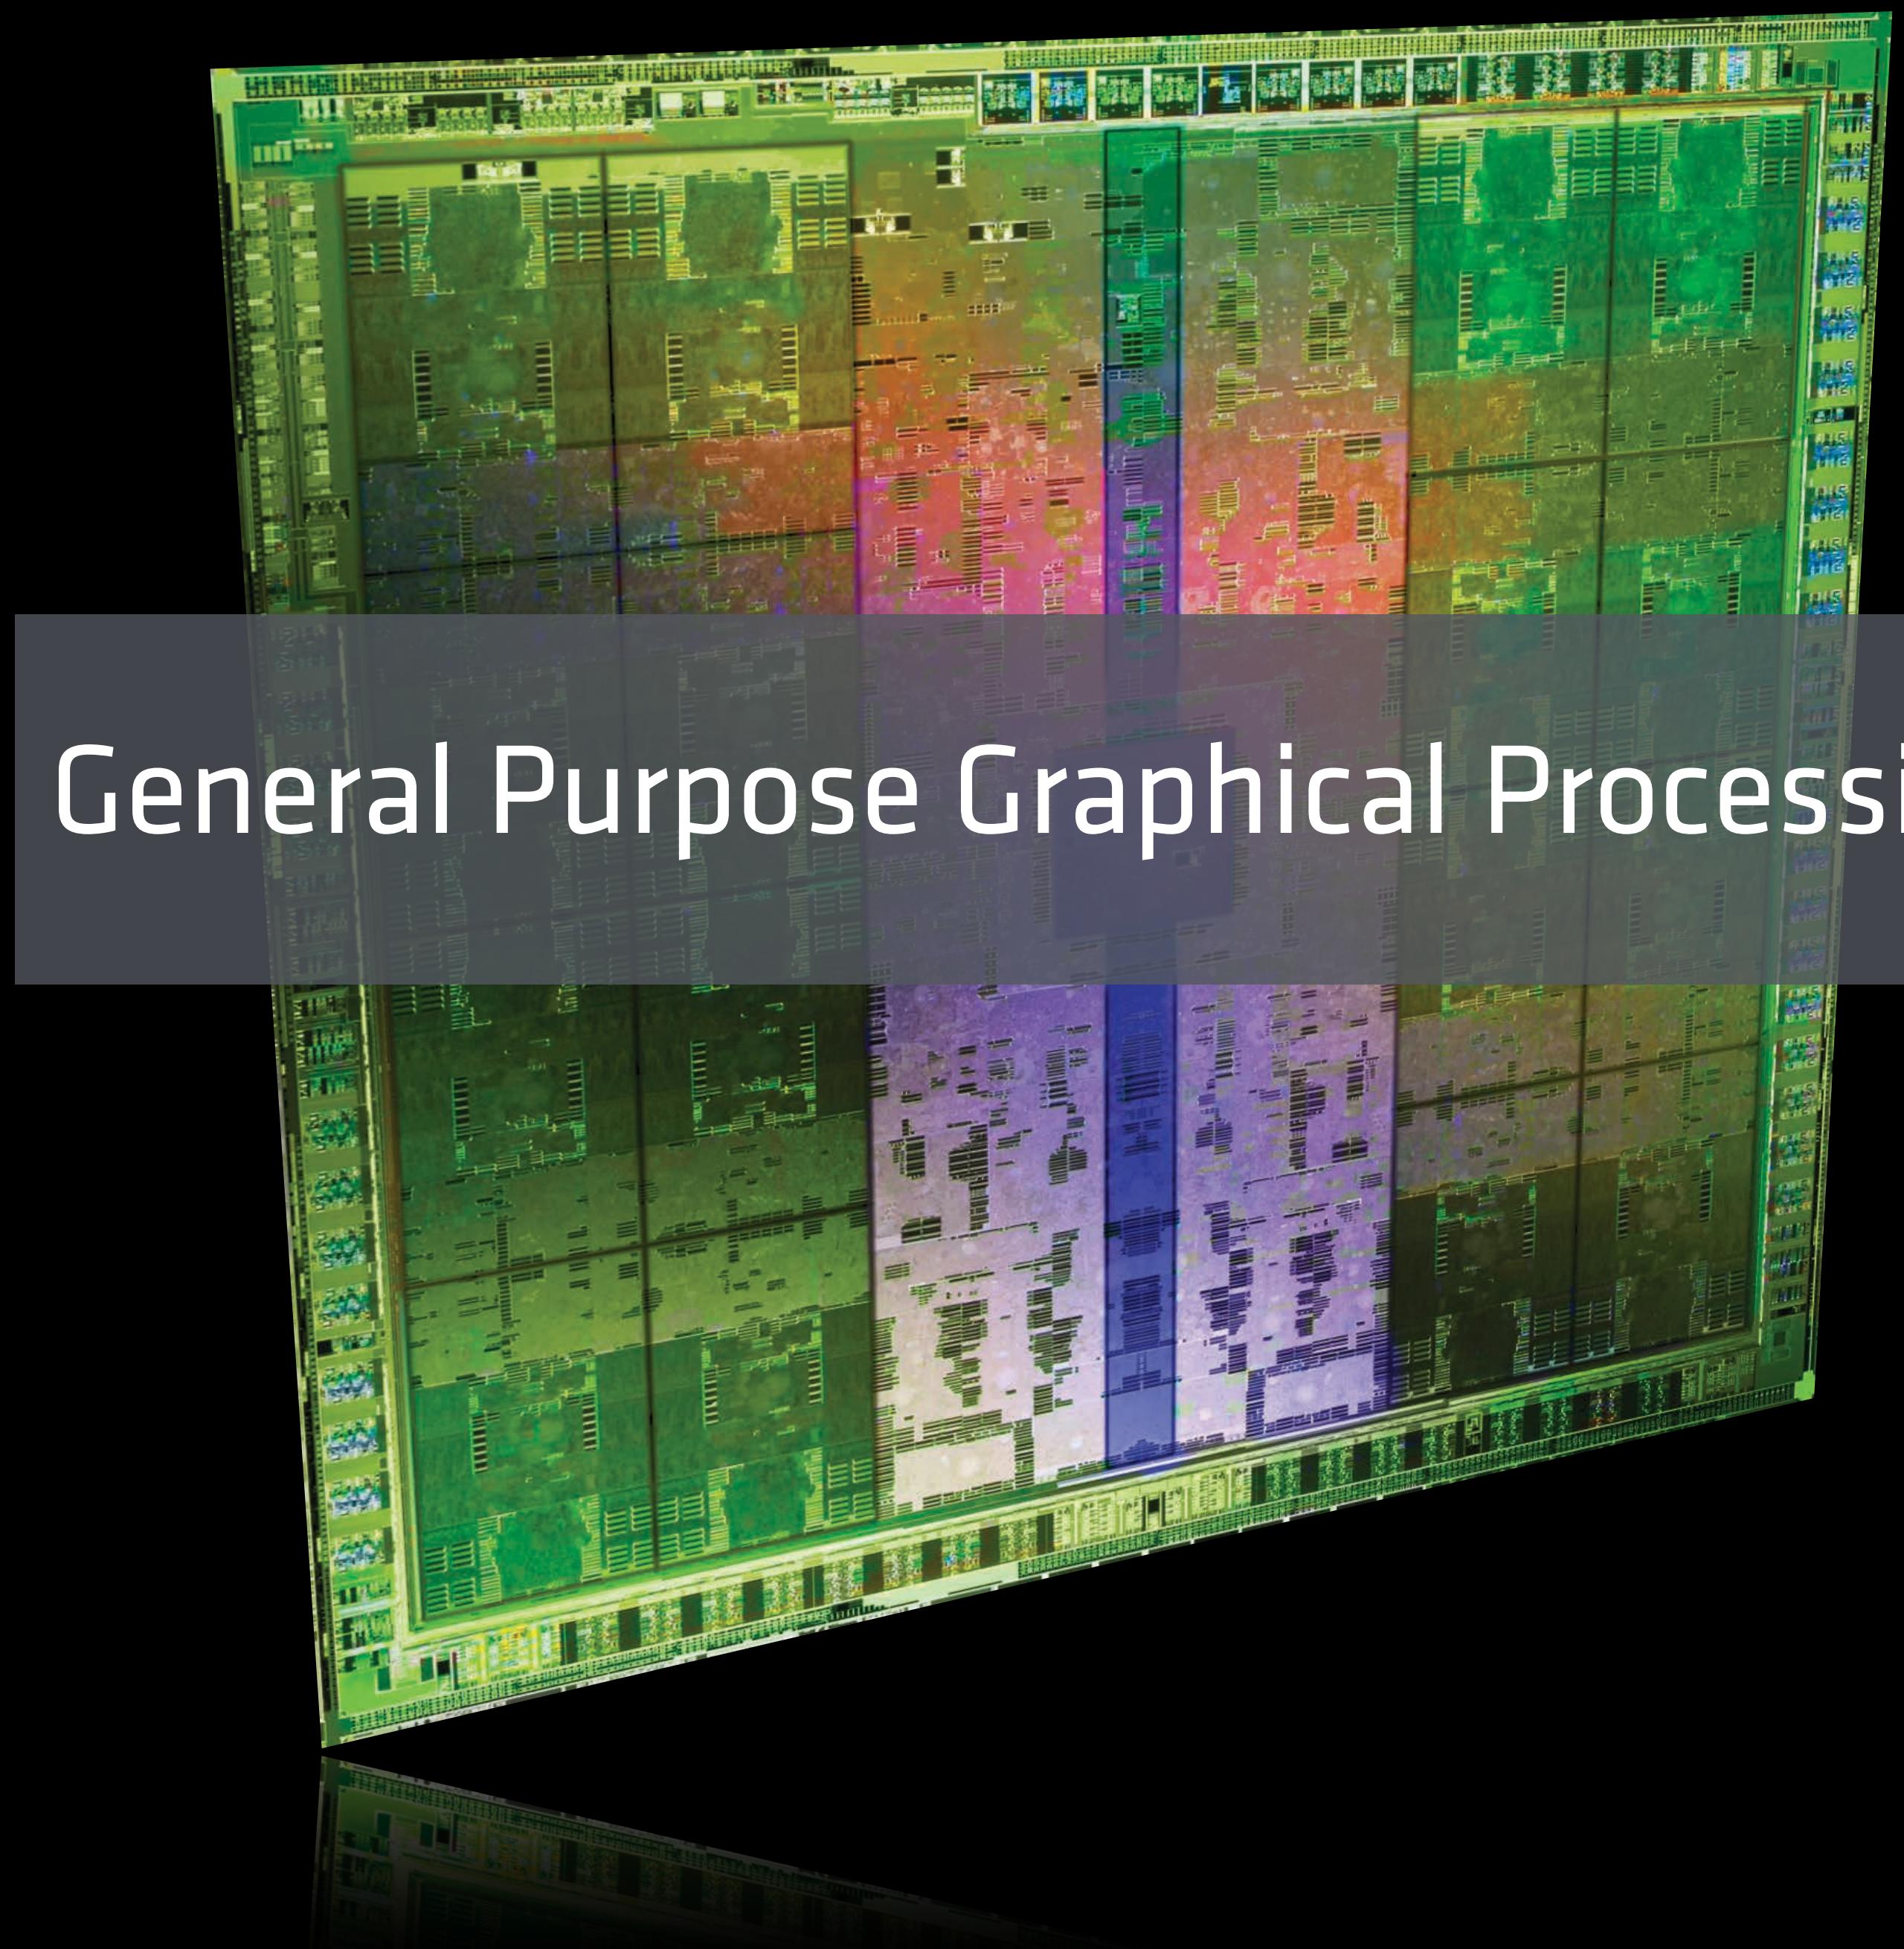
**Harvard Mark I - 1944**

Rear view of Computing Section

# Overview

- **Outstanding progress in computational power since 1950s**
  - Present systems can reach performances of 0.1 EFlop/s
  - Energy cost for calculations has gone down by 14 orders of magnitude
  - Continuous evolution of architectures and computing paradigms
- **OSIRIS is ready to take advantage of the most advanced computing resources in the world**
  - Only 1 system in the top-10 is not supported
  - 8 systems in the top-10 are fully supported, with computing kernels developed specifically to take advantage of advanced hardware features
- **The Intel Knights Landing Architecture is a promising architecture for PIC simulations**
  - Existing codes and algorithms can be efficiently ported to this architecture
  - Good support for MPI and OpenMP parallelism
  - Explicit vectorization critical for good performance
  - New MCDRAM memory has high impact on performance at lower interpolation orders





# General Purpose Graphical Processing Units

NVIDIA Fermi K20x die

